

Android Optimization: Myth and Reality

Jim Huang (黃敬群) <jserv@0xlab.org>

Developer, 0xlab

<http://0xlab.org/>

July 8, 2011

Rights to copy

© Copyright 2011 **0xlab**

<http://0xlab.org/>

contact@0xlab.org



Attribution – ShareAlike 3.0

You are free

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Corrections, suggestions, contributions and translations are welcome!

Latest update: July 10, 2011

Under the following conditions

- **BY:** **Attribution.** You must give the original author credit.
- **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.
- For any reuse or distribution, you must make clear to others the license terms of this work.
- Any of these conditions can be waived if you get permission from the copyright holder.

Your fair use and other rights are in no way affected by the above.

License text: <http://creativecommons.org/licenses/by-sa/3.0/legalcode>



Agenda

(1) Android Meets

Hardware Revolution

(2) Unproven "Optimization"

(3) Android Evolution

(4) Mythbusters



Revolution

- (1) an attempt, by a large number of people, to change the government of a country, especially by violent action (革命)
- (2) a great change in conditions, ways of working, beliefs, etc. that affects large numbers of people (巨變)

Evolution

- (1) the gradual development of plants, animals, etc. over many years, from simple to more complicated forms (進化)
- (2) the gradual development of something (演變 ; 發展 ; 漸進)



My interpret of Android:
Hardware is Revolution;
Software is basically Evolution;
Android is Hardware-driven Software Revolution



Android Meets Hardware Revolution



Take ARM for Example

- 32-bit instructions, with extension to support 16-bit Thumb & Thumb-2 instructions.
- Single unified memory address space (i.e. all peripherals and I/O are accessed like normal memory, at certain specific memory locations).
- Relatively low power consumption.
- Run at wide range of clock cycle
 - (June 1, 2011) Qualcomm announced qual-core MSM8964/APQ8064
 - 2.5 GHz / 2MB L2 cache
 - sample in early 2012

Do you remember the specifications of your old laptop computers?



Classic ARM Processors (1)

- ARM7TDMI family
 - Based on ARMv4T architecture with 3-stage pipeline
 - supports the 16-bit Thumb instruction set
 - supports the JTAG Debugger
 - includes a fast Multiplier to support DSP algorithm
 - supports the In-Circuit Emulation interface
- ARM9TDMI family
 - Based on ARMv4T with Harvard cache architecture
 - 5-stage pipeline
 - ARM920T is based on ARM9TDMI with a memory management unit (MMU)



Classic ARM Processors (2)

- ARM9E family / Intel's Xscale
 - Based on ARMv5E architecture
 - Enhanced with DSP instructions
 - Hardware support of Java bytecodes execution
- ARM10 family
 - Based on ARMv5E with MMU
- **ARM11** family
 - Based on ARMv6 architecture
 - Supports SIMD instructions
 - Reduce context switching cost

The first shipped Android phone belongs to ARM11 family.

Android 1.5 (2009/04/27)

ARMv5te optimizations: libc (thumb, memcpy, memset, clz),
dalvik (fast interpreter), skia, surfaceflinger, pixelflinger, audioflinger



Classic ARM Processors (3)

Android 2.0 (2009q4)

ARMv7-a optimizations: libc (memcpy/neon), dalvik (JIT), skia

- **Cortex families**
 - Based on ARMv7 architecture
 - Supports the new Thumb-2 instruction set
 - NEON SIMD engine
 - **Cortex-A**: For complex OS based applications
 - Production: Cortex-A8 / Cortex-A9
 - Cortex-R: For real-time embedded applications
 - Cortex-M: For deeply embedded, microcontroller type cost sensitive applications
- Cortex-A15
 - Triple dispatch superscalar, out-of-order
 - High performance prefetcher
 - Up to 4M integrated L2 cache
 - 1TB of physical address space
 - Enterprise class performance, mobile power

The era of Android:

High quality
mobile computing





Security Support

ARM11

ARM9

TrustZone™

SIMD

VFPv2

Jazelle®

V5

V6

Thumb-EE

VFPv3

NEON™
Adv SIMD

Thumb®-2

V7A

Execution Environments:
Improved memory use

Improved Media and DSP (2xV6)

Improved code density, approx 30%

Improved Media (2xV5)



ARM Optimizations in Android

- Thumb2 optimizations in Dalvik VM
 - On real memory systems with latency the additional instructions in I-cache latency, it makes Thumb2 higher performance than ARM (logic: 1.23x)
- NEON/VFP optimization in skia & dalvik
 - 1.45x in dalvik; 2-5x microbench in skia
- Thumb2 + VFP in V8 JavaScript engine
- ARMv6 atomic/rev in libc & dalvik
- NEON + VFP in StageFright
- SMP in libc, libcutils, dalvik

Again, we were informed of the improvements from hardware and dedicated software.
But, is it enough?



Unproven “Optimization”



“Optimization” → 最佳化 (?)



「最佳化的迷思」

操作績效不理想或想再創佳績
於是就開始進行程式最佳化
在操盤軟體上東調西調
嚐試在變化萬千的參數搭配中找尋聖杯
果然出現創紀錄的績效時
就以為得到聖杯
拿到市場實地操作
卻發現落差很大

最佳化的對象是誰
程式還是策略
拿頭文字 D 來比喻吧
程式好比藤原拓海的車子
車子當然可以越改越好
但是藤原拓海才是致勝靈魂
對比程式 策略才是致勝關鍵

策略核心邏輯不是藤原拓海等級
參數再怎麼調
也調不出來一個藤原拓海

回到以策略為中心點
檢視策略在買賣點、停損停利點與資金
控管上
可以改善的空間
才能突破你自己

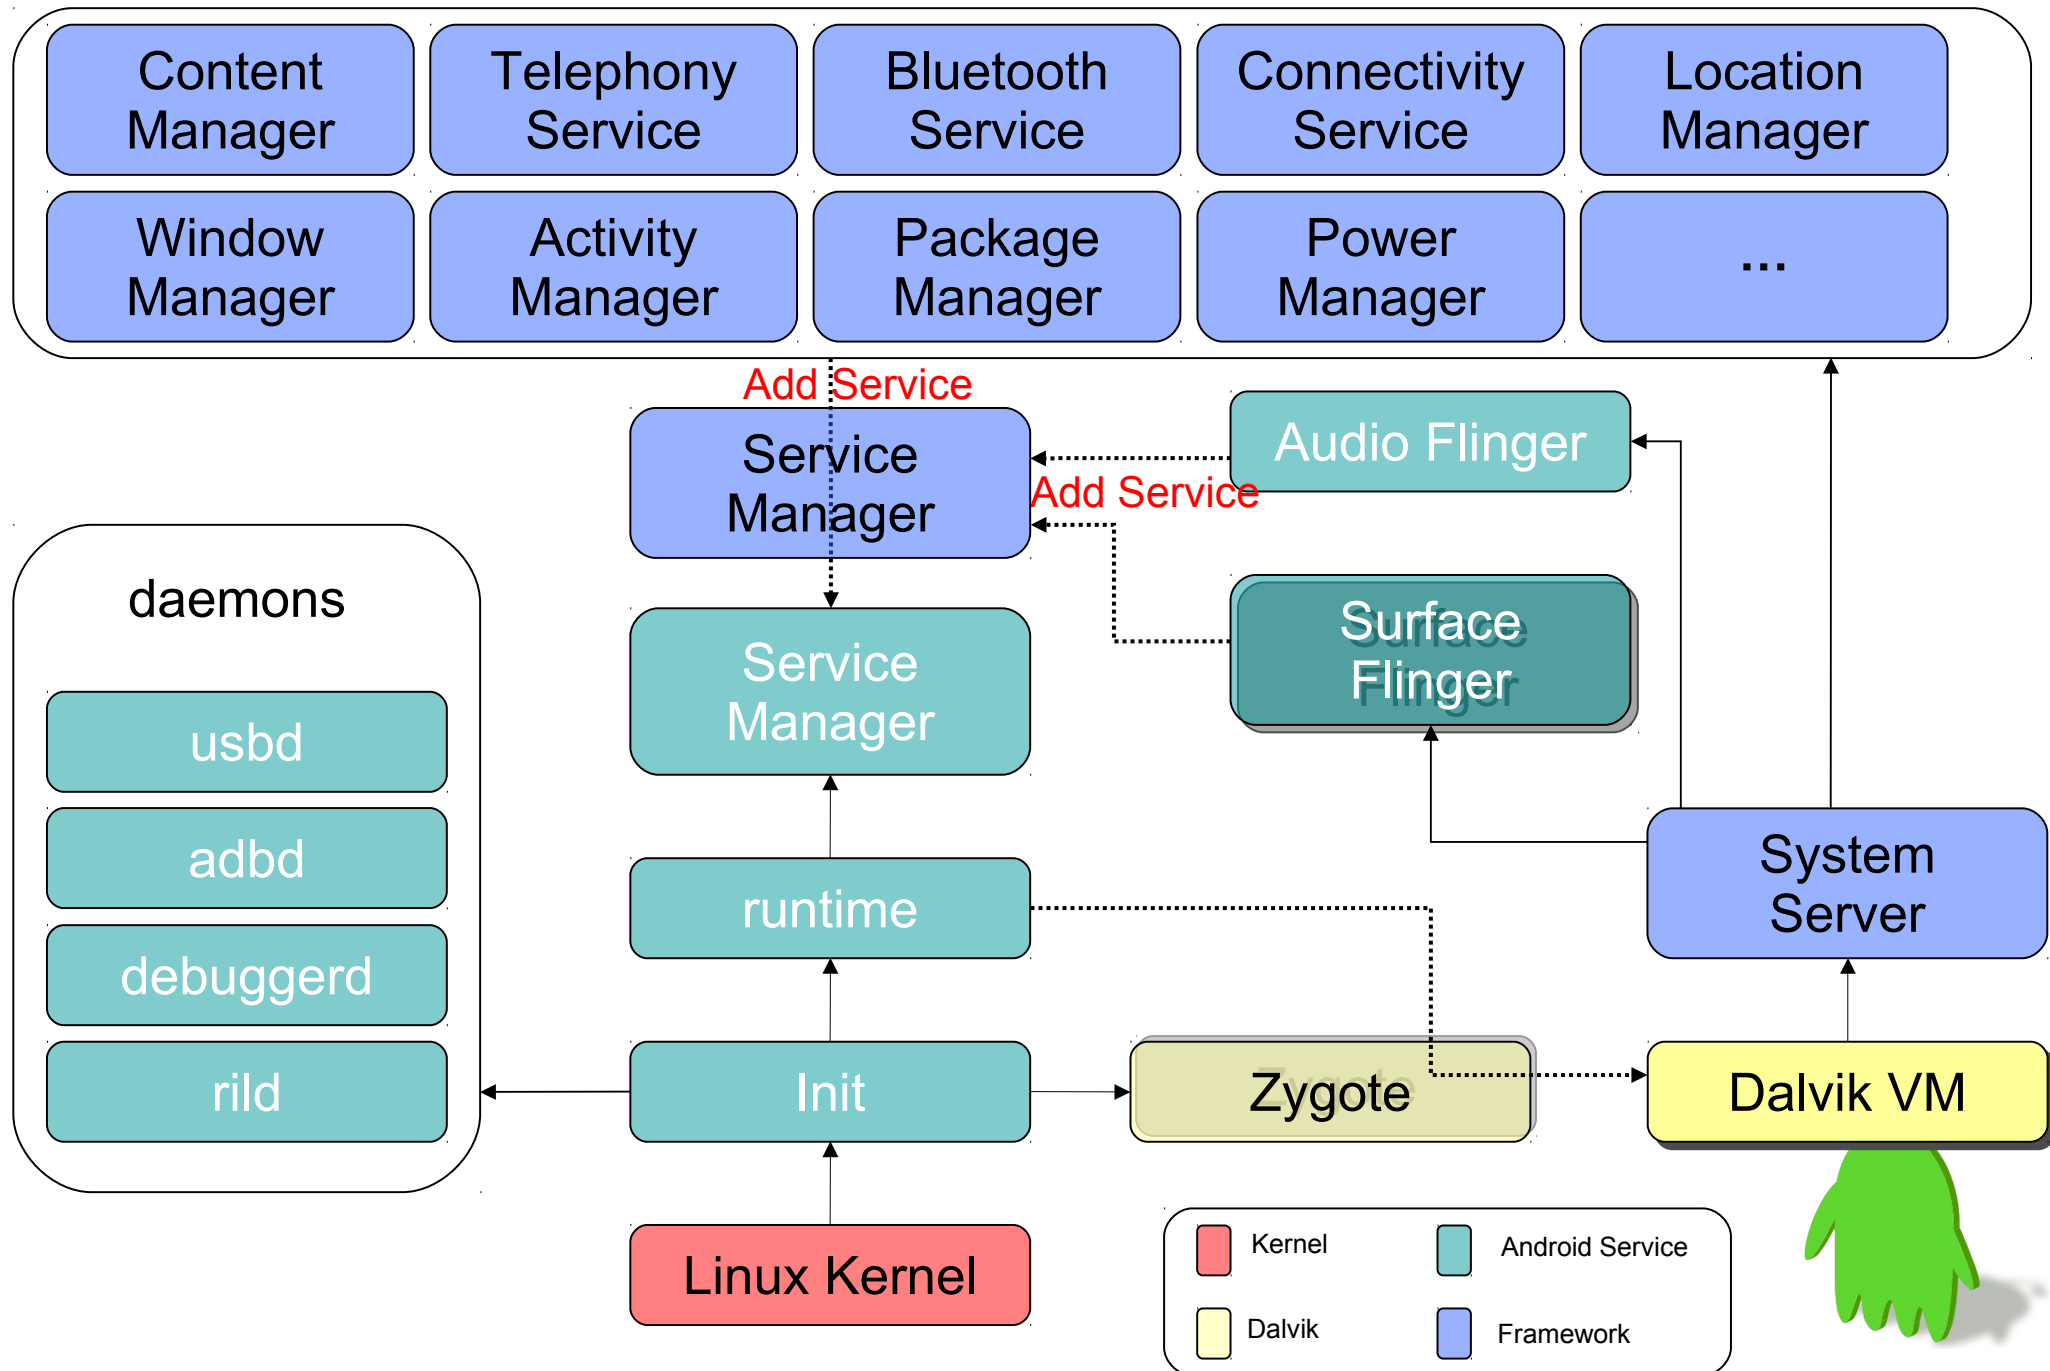


Key Concept: **Strategy**

Even runtime is (claimed as) fully optimized, we still have to look into details carefully.


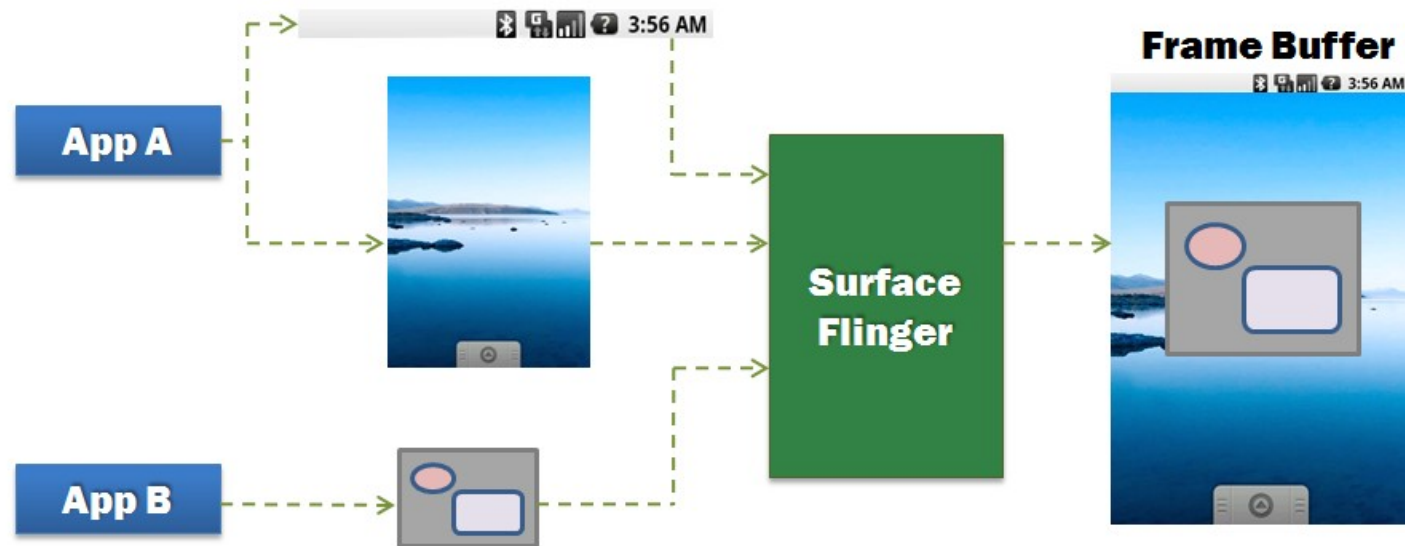



Android Services in Action



Android SurfaceFlinger

- Properties
 - Can combine 2D/3D surfaces and surfaces from multiple applications
 - Surfaces passed as buffers via Binder IPC calls
 - Can use OpenGL ES and 2D hardware accelerator for its compositions
 - Double-buffering using page-flip





Application Process

JAVA Application

Dalvik VM

Framework JNI

Framework Process

Multimedia Framework and its Components

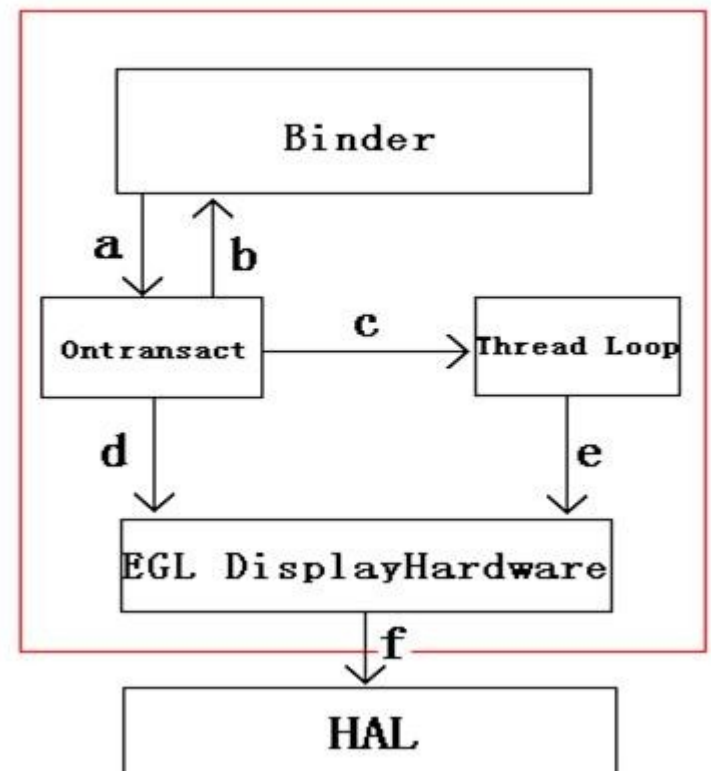
B
I
N
D
E
R

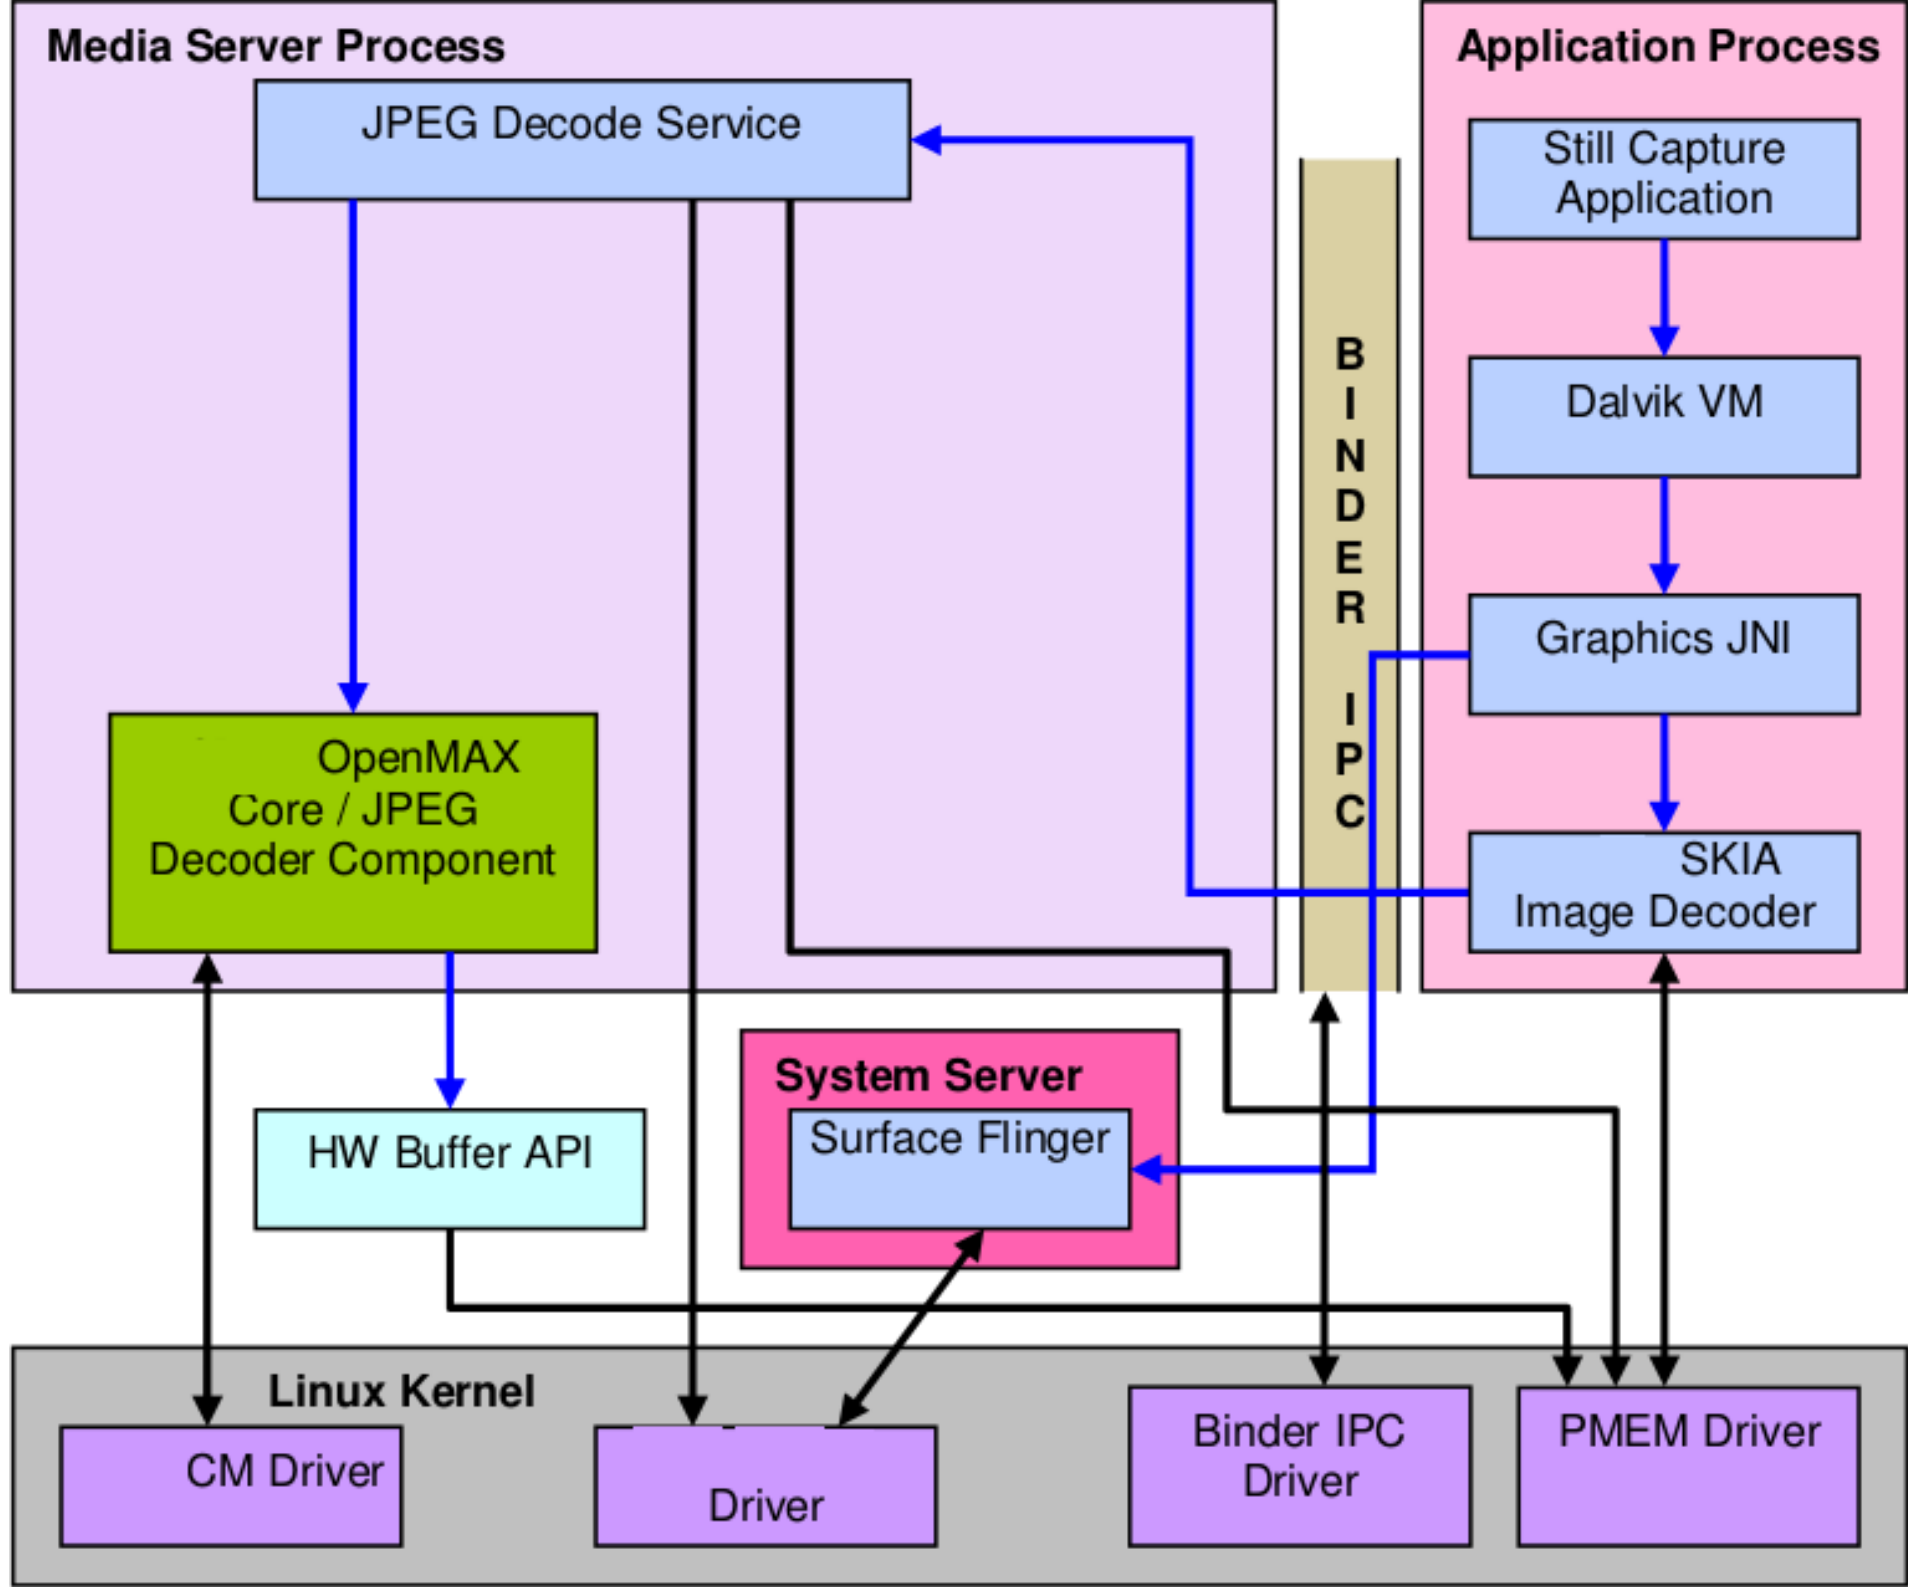
I
P
C

Linux Kernel

Binder IPC Driver

SurfaceFlinger





System Server Process

Surface Flinger Service

CopyBit HAL

EGL / OpenGL API

Linux Kernel

Driver

Driver

MALI Driver

Binder IPC Driver

Application Process

Still Capture Application

Dalvik VM


Display JNI

OpenGL / EGL JNI

EGL / OpenGL API

B
I
N
D
E
R



I
P
C




Optimizing SurfaceFlinger Techniques

- Measure/Profile the execution flow
- Scenario-driven system analysis
 - 2D/3D animation
 - Video Playback
 - Camera preview/recording
- Check list:
 - Color space converting
 - Eliminate the invalid cases (quality of HAL)
 - Allow zero copy among framework/HAL/Driver
 - Shared memory / IPC cost
 - Any misconfiguration?
 - Cache (write-through vs. write-back)
 - The last: essential routines






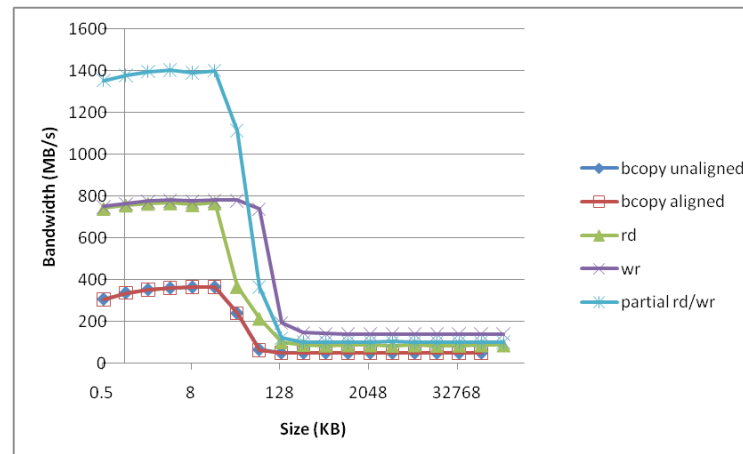

- Write bandwidth greatly affected if caches are configured as write-through
- Configure caches as write-back, with allocate-on -write



- Memory bandwidth

Single Instance

2 Instances




Possibly Premature optimizations

- **“Premature optimization is the root of all evil”**
 - Donald Knuth
- Bionic
 - Glibc-incompatible, No SysV IPC, partial Pthread, incomplete prelink
 - inactive/incorrect kernel header inclusion
 - May not re-use existing system utilities
- Assumed UI behaviors
 - Input event dispatching and handler
 - Strict / non-maintainable state machine (policy)
 - Depending on a certain set of peripherals
- Unclear HAL design and interface
 - Wifi, Bluetooth, GPS, ...



Android Evolution





Android is Hardware-driven Software Revolution

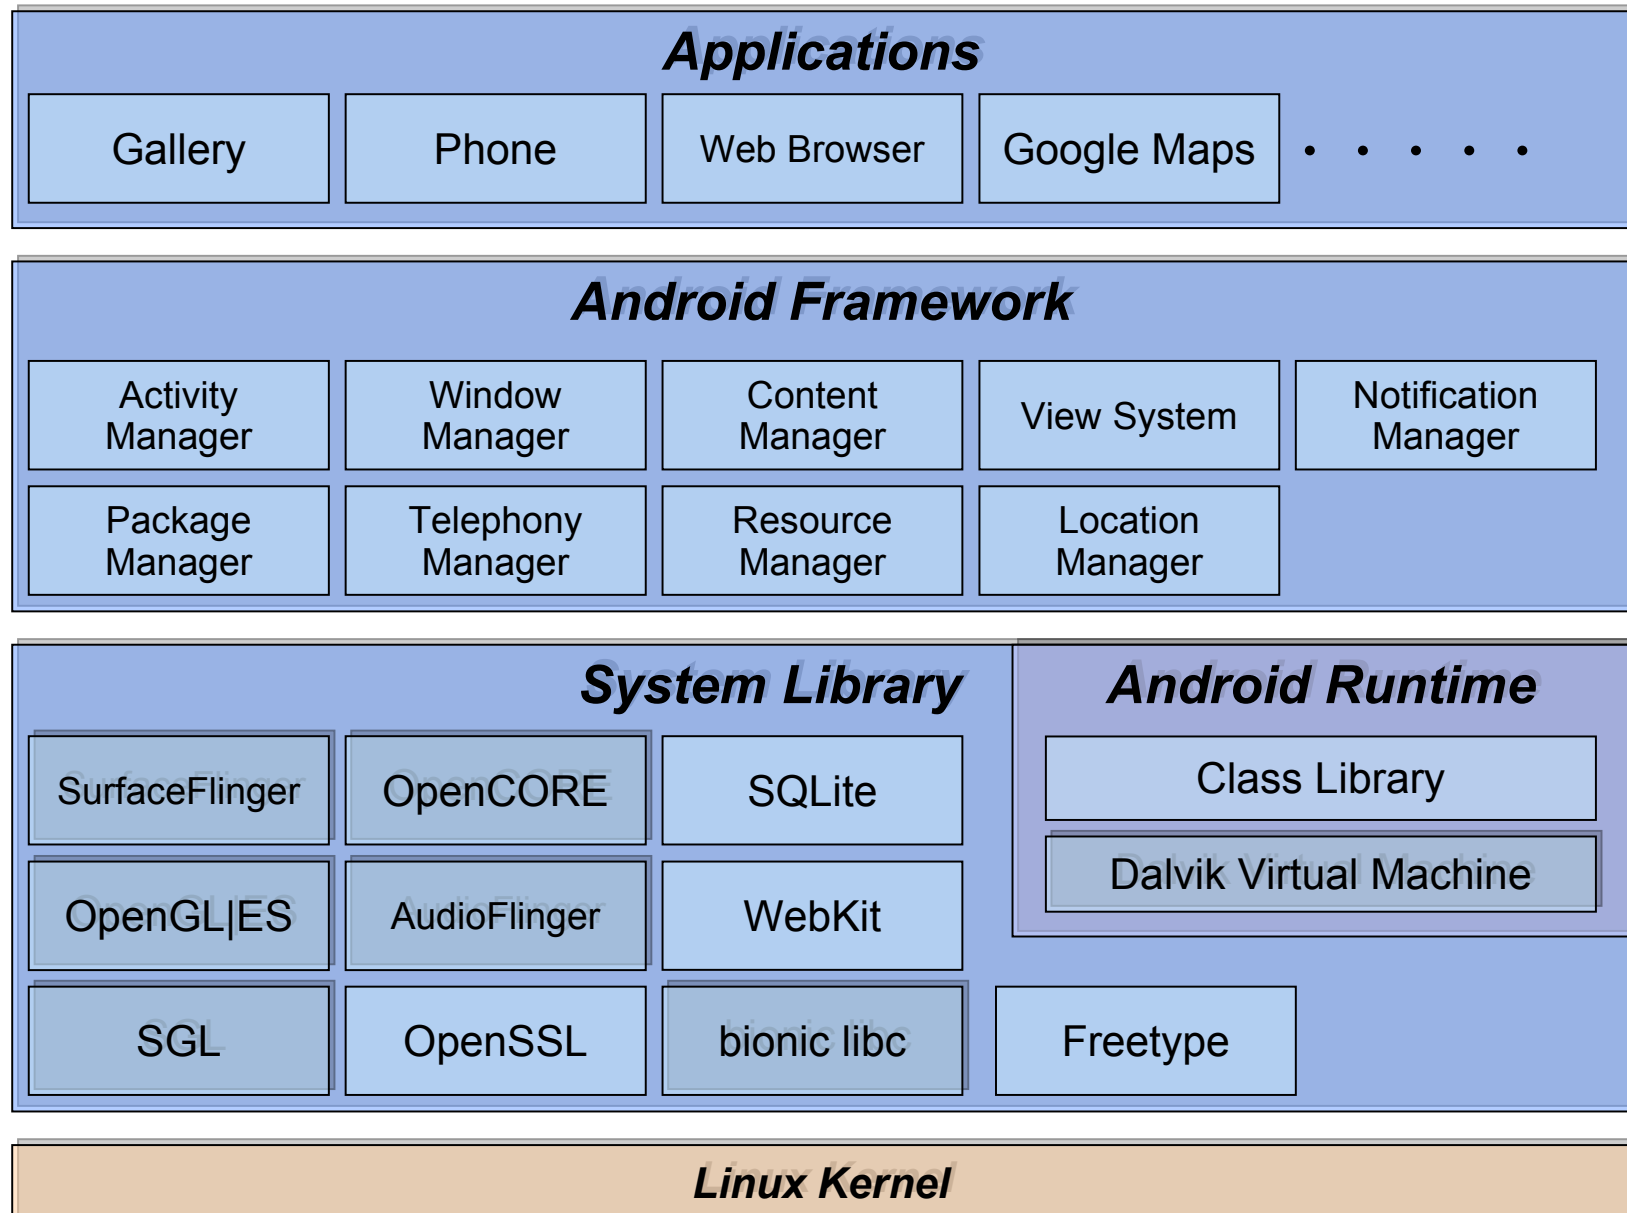


Android Software Evolution

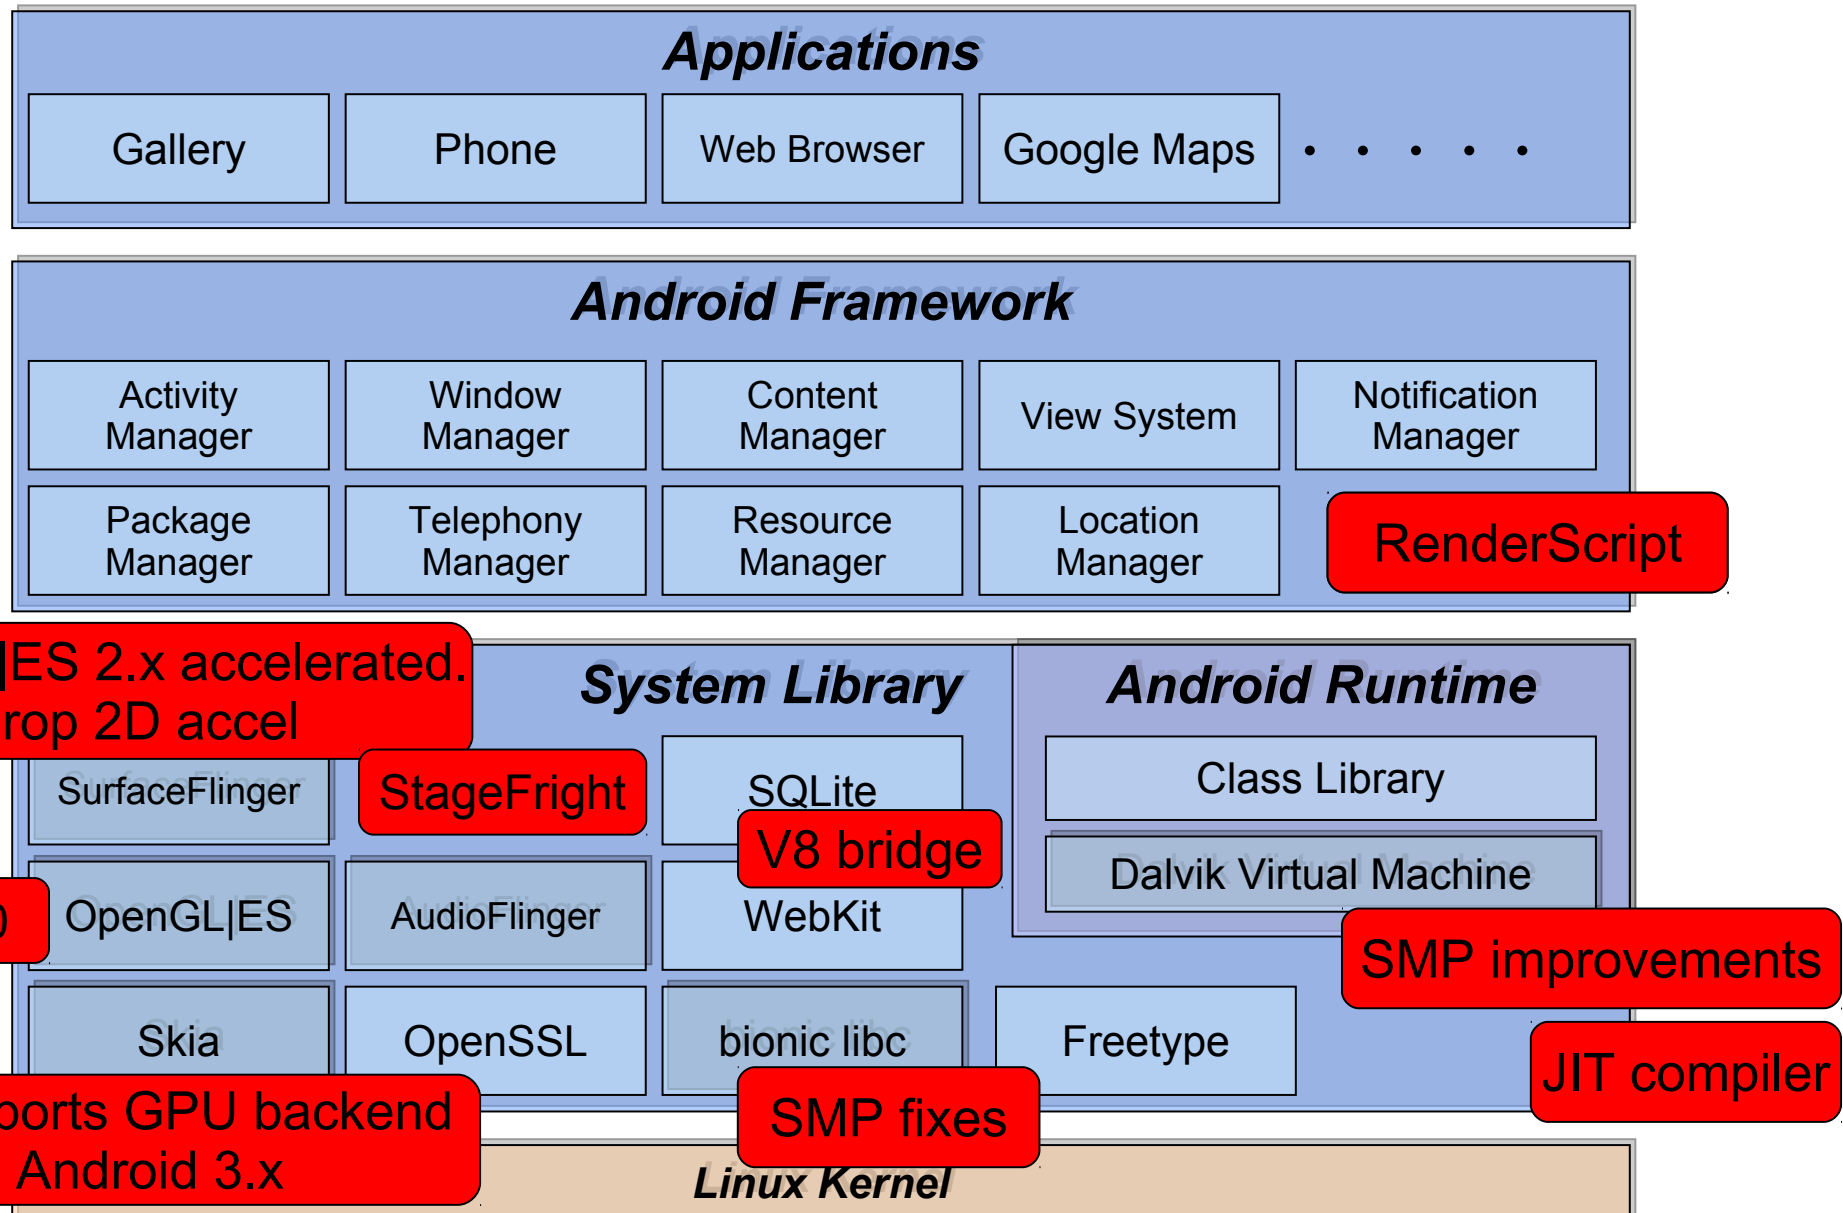
- Excluding the trickery, Android is a huge software project and publishes large parts of source repositories to AOSP (Android Open Source Project).
- “Hardware-driven software revolution” is the trickery. To fairly develop, it has to turn into "evolution".
 - SoC vendors built up the community(-like) hosted Android variants. → Qualcomm, TI, Samsung, ST-Ericsson, ...
- Quality-control is another key factor
 - Oxlab delivered Oxbench and related validation infrastructure.
- Solution to Android ecosystem is to continue to grow the AOSP community
 - need to unify to move forward
 - Reasons why Linaro exists



Functional View (1.5)



Functional View (2.3)



Frequently Asked Questions

- "What is the minimal hardware requirement for running Android?"
→ UX is the key.
- "May I reduce BOM cost in my Android product?"
→ cost down is always possible, but it depends on the overall performance.
- "How can I upgrade Android without pain?"
- "I need XXX feature from newer version of Android. How can I get it into the current and old codebase?"
→ Partially upgrading is feasible for system libraries, but always difficult for Android framework



Mythbusters



Think Difficult

- The Effects of Convergence
 - Mobile computing platform
 - Ecosystem
- Content providers lead the development of containers (devices)
- Innovation is integrated by open source software model. The cost of proprietary software gets much higher.
- To optimize or not to optimize, that is the question.
 - Merge Local optimizations != Optimized globally
 - Hardware revolution sharpens optimization strategies and product directions.
 - The role of your team





<http://0xlab.org>