

Implement Checkpointing for Android

(to speed up boot time and development process)

Kito Cheng (程皇嘉) <kito@0xlab.org>

Jim Huang (黃敬群) <jserv@0xlab.org>

Developer, 0xlab

Rights to copy

© Copyright 2012 **0xlab**

<http://0xlab.org/>

contact@0xlab.org



Attribution – ShareAlike 3.0

You are free

Corrections, suggestions, contributions and translations
are welcome!

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Latest update: Jun 12, 2012

Under the following conditions

- **BY:** **Attribution.** You must give the original author credit.
- **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.
- For any reuse or distribution, you must make clear to others the license terms of this work.
- Any of these conditions can be waived if you get permission from the copyright holder.

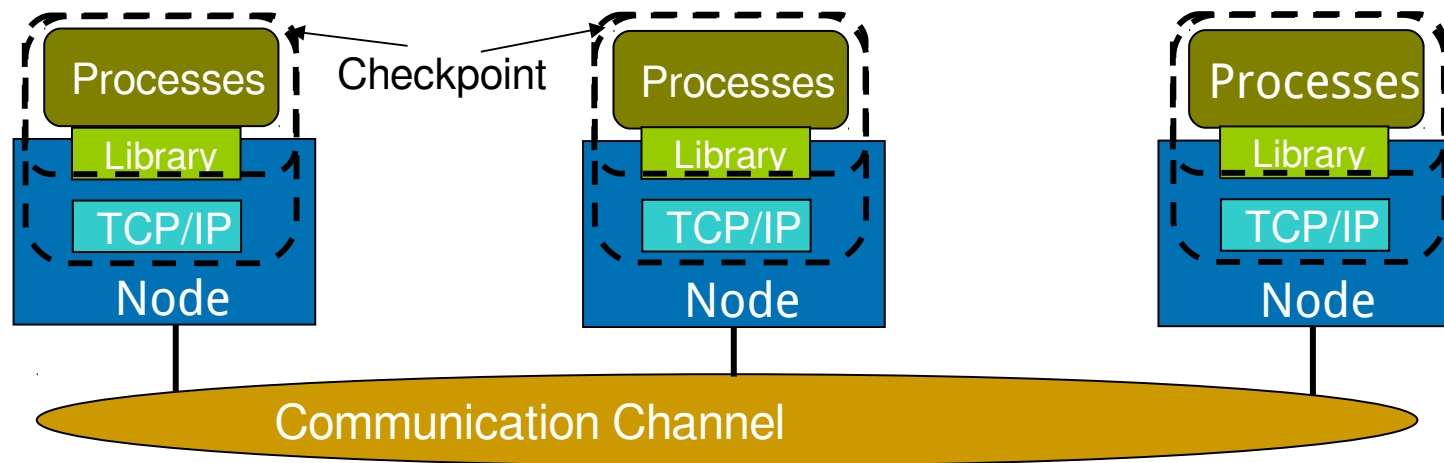
Your fair use and other rights are in no way affected by the above.

License text: <http://creativecommons.org/licenses/by-sa/3.0/legalcode>



Basic Idea: Process Migration

- Process Migration (in past applications)
 - Distributed Load Balancing
 - Efficient Resource Utilization
- Crash Recovery and Rollback Transaction
 - Useful to system admin

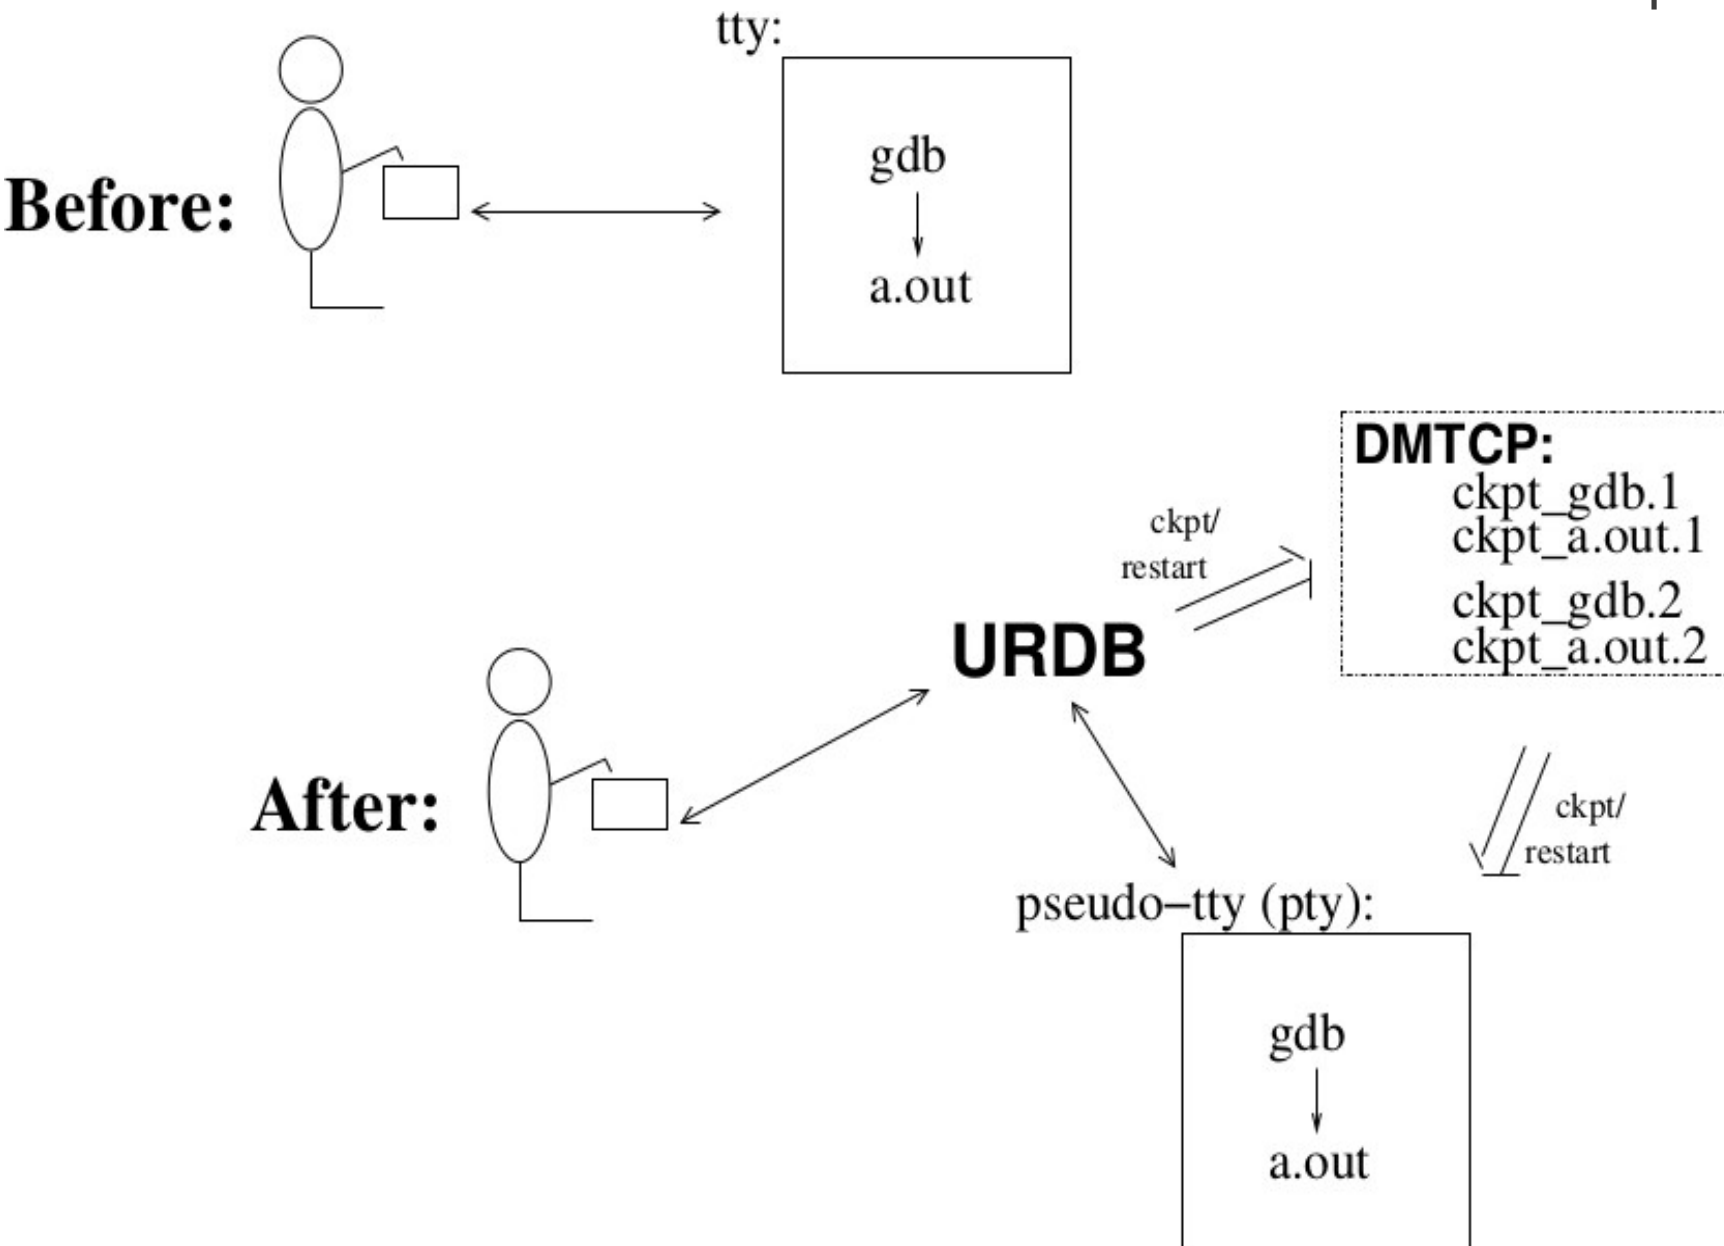


Checkpointing

- From Wikipedia:
 - ... is a technique for inserting fault tolerance into computing systems. It basically consists of storing a snapshot of the current application state, and later on, use it for restarting the execution in case of failure.
- Used in distributed shared memory systems
- Even used in reversible debugger
- Different from virtual machine level snapshot/resume mechanisms
 - Checkpointing emphasizes on process level.



Reversible Debugger using Checkpointing



Reasons to Implement Checkpointing for Android

- Resume to stored state for faster Android boot time
- Better product field trial experience due to regular checkpointing
- Deploy problematic states for engineering analysis and debugging transparently
- Q&A stress test purpose



Expectations of Checkpointing

- Application-transparent
 - supports applications without modifications or recompilation
- Supports a broad class of applications
 - Databases
 - parallel / MPI apps
 - desktop apps
- Comprehensive support for user-level state, kernel-level state, and distributed computation and communication state
- Supported on unmodified Linux kernel
 - checkpoint-restart should be integrated by addons



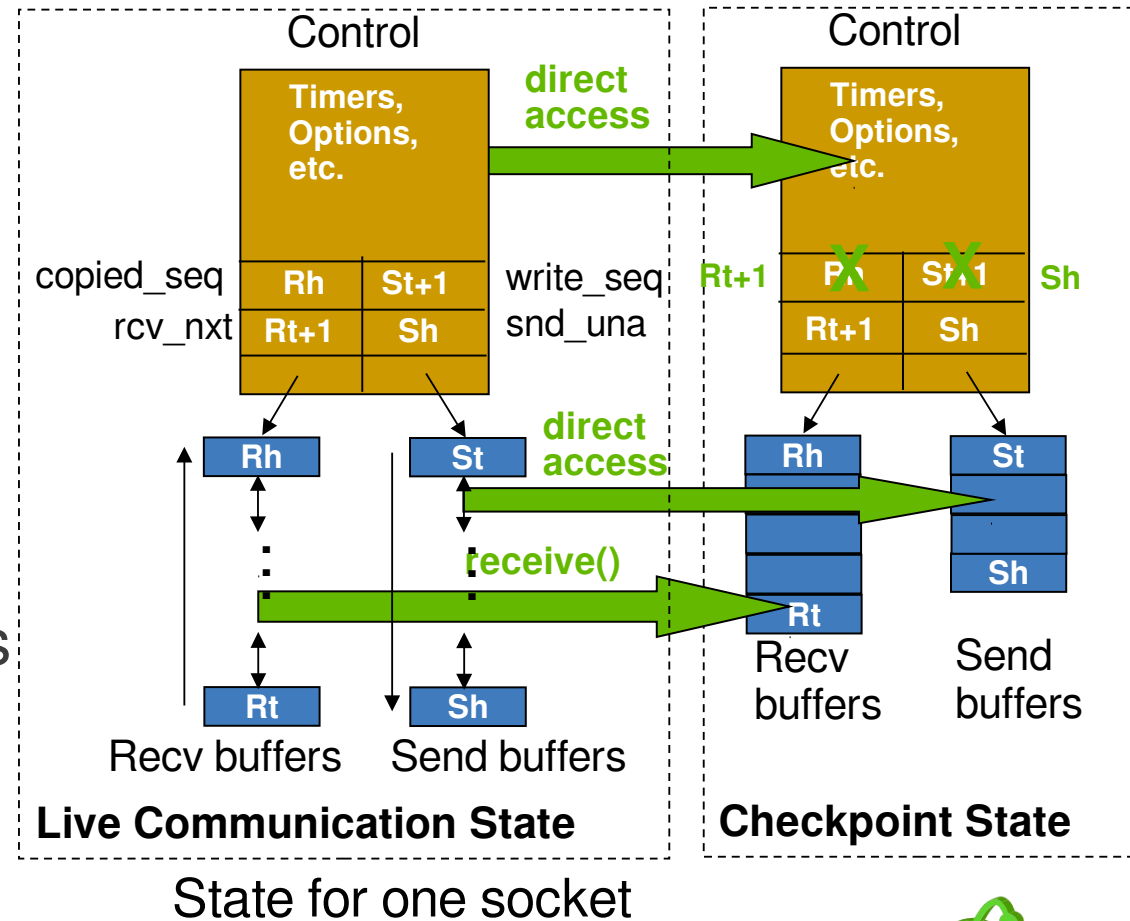
Challenges in checkpoint and restore

- Network stack will continue to execute even after application processes are stopped
- No system call interface to read or write control state
- No system call interface to read send socket buffers
- No system call interface to write receive socket buffers
- Consistency of control state and socket buffer state



Communication state checkpoint

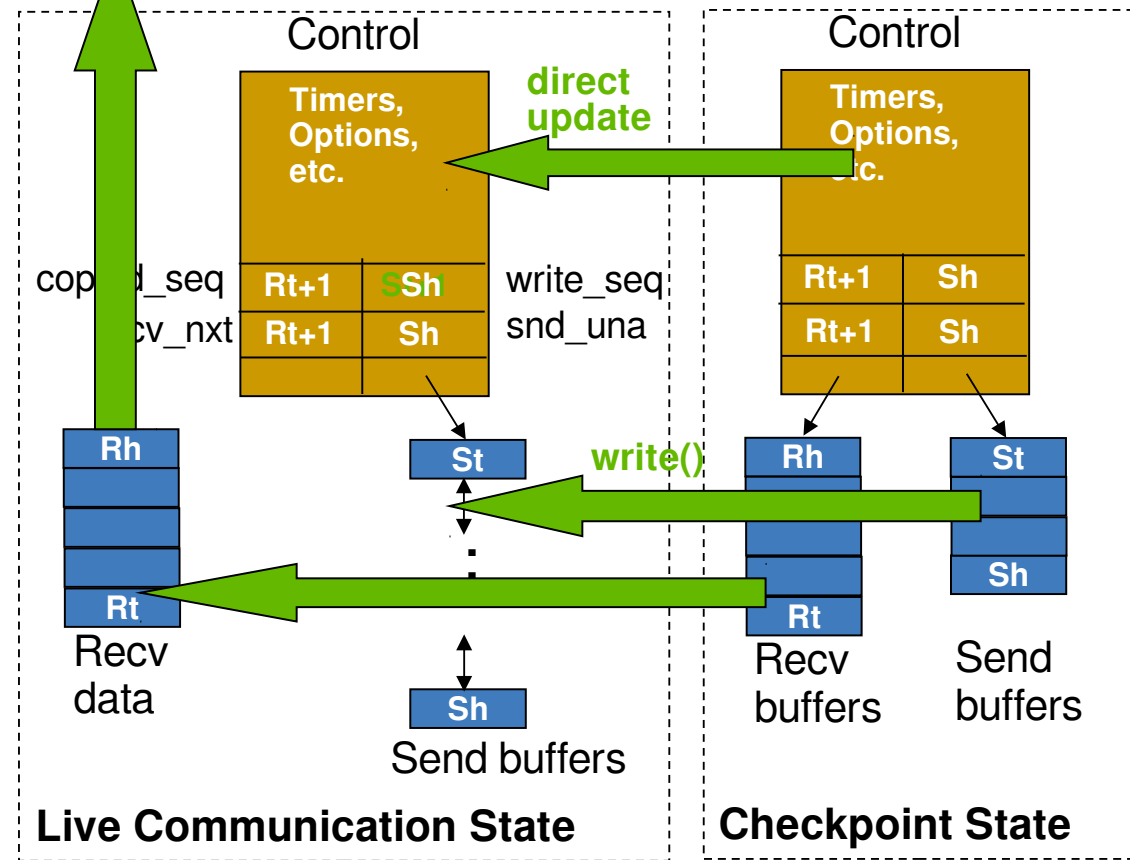
- Acquire network stack locks to freeze TCP processing
- Save receive buffers using socket receive system call in peek mode
- Save send buffers by walking kernel structures
- Copy control state from kernel structures
- Modify two sequence numbers in saved state to reflect empty socket buffers



Communication state restart

- Create a new socket
- Copy control state in checkpoint to socket structure
- Restore checkpointed send buffer data using the socket write call
- Deliver checkpointed receive buffer data to application on demand

To App by intercepted receive system call



Existing Checkpointing mechanisms

- CryoPID
 - <http://cryopid.berlios.de/>
- BLCR (Berkeley Lab Checkpoint/Restart)
 - <https://ftg.lbl.gov/projects/CheckpointRestart/>
- DMTCP
 - <http://dmtcp.sourceforge.net/>



Implementation Considerations

- Checkpointing can be implemented in
 - kernel modifications + helpers in userspace
 - pure userspace
- Introduce a virtualization layer groups processes into specific states with private virtual name space
 - Intercepts system calls to expose only virtual identifiers (e.g., vpid)
 - Preserves resource names and dependencies across migration
- Mechanism to checkpoint and restart states
 - User and kernel-level state
 - Primarily uses system call handlers
 - File system not saved or restored




- **D**istributed **M**ulti-**T**hreaded **C**heck**P**ointing.
- Works with Linux Kernel 2.6.9 and later.
- Supports sequential and multi-threaded computations across single/multiple hosts.
- Entirely in user space (no kernel modules or root privilege).
 - Transparent (no recompiling, no re-linking).
- Written in Northeastern University and MIT and under active development since 2006.
- License: GNU LGPL (allows freely using and linking)



Process Structure

CT = DMTCP checkpoint thread
T = User Thread



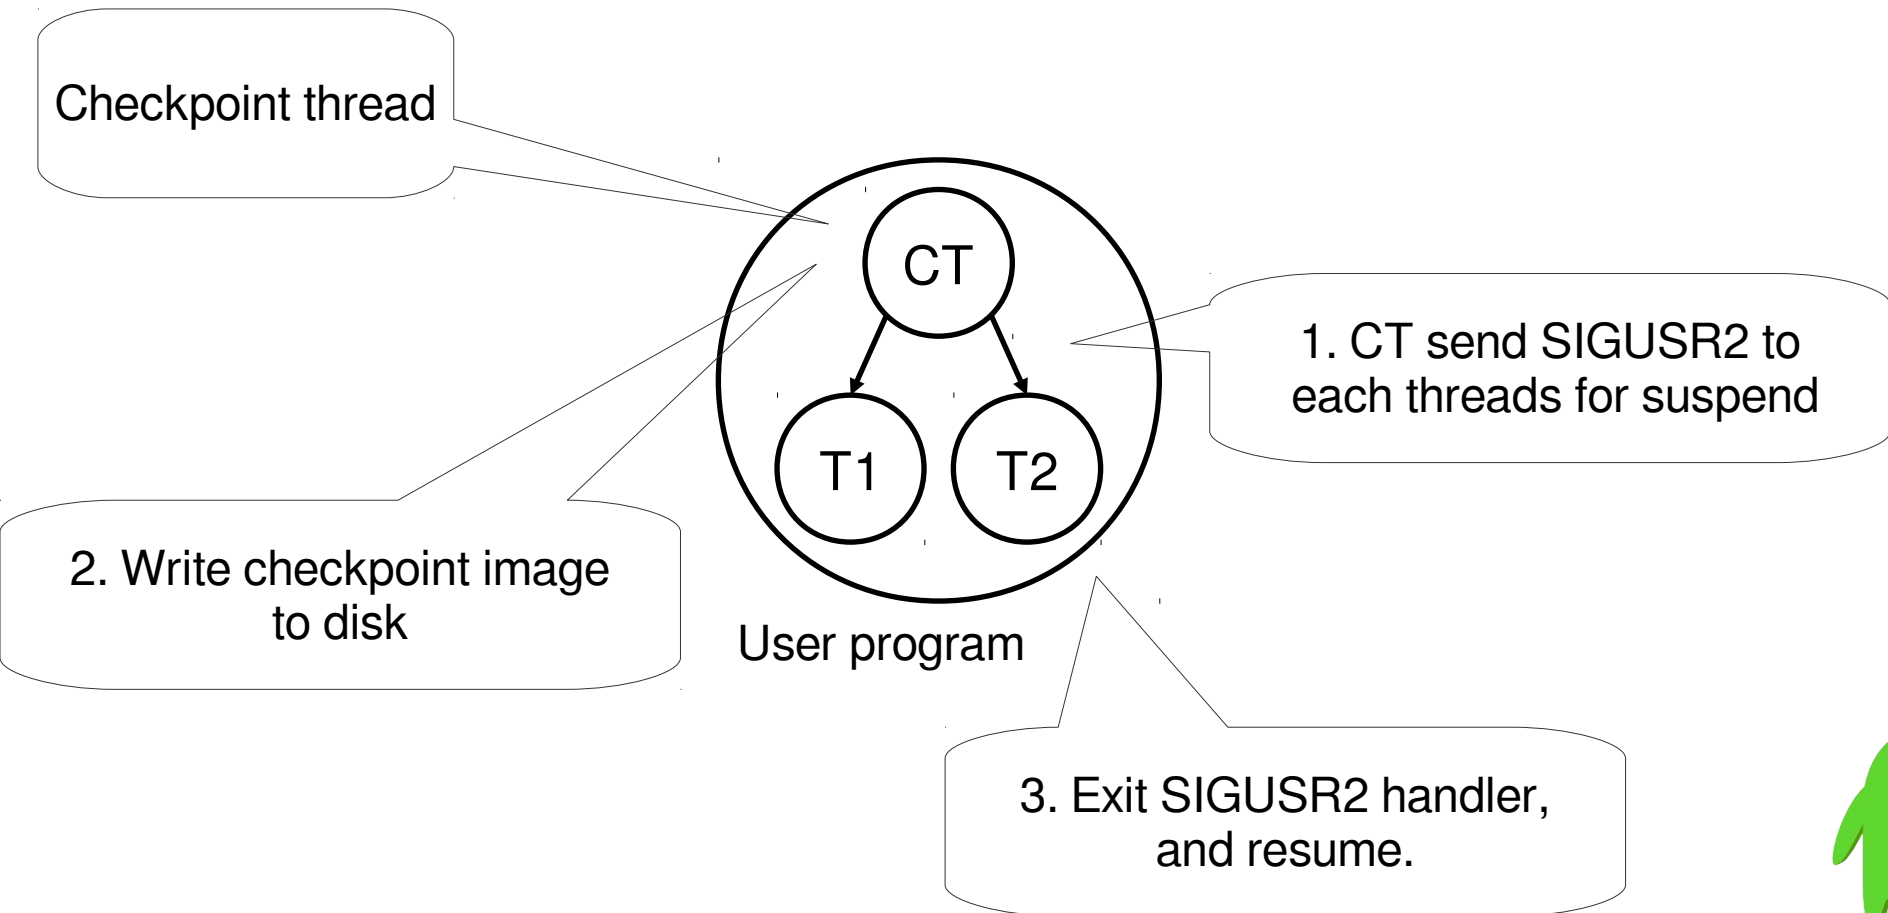
```
dmtcp_checkpoint <EXE> # starts coordinator  
dmtcp_command -c      # talks to coordinator  
dmtcp_restart ckpt_<EXE>-* .dmtcp
```

- Coordinator: a stateless synchronization server for the distributed checkpointing algorithm.
- Checkpoint/Restart performance related to size of memory, disk write speed, and synchronization.



How DMTCP works (1/4)

- MTCP : component for checkpoint single-process
- SIGUSR2: Used internally from checkpoint thread to user threads.



How DMTCP works (2/4)

- **LD_PRELOAD:** Transparently preloads checkpoint libraries ``dmtcphijack.so`` which installs libc wrappers and checkpointing code.
- **Wrappers:** Only on less heavily used calls to libc
 - `open, fork, exec, system, pipe, bind, listen, setsockopt, connect, accept, clone, close, ptsname, openlog, closelog, signal, sigaction, sigvec, sigblock, sigsetmask, sigprocmask, rt_sigprocmask, pthread_sigmask`
 - Overhead is negligible.



How DMTCP works (3/4)

- Additional wrappers when process id & thread id virtualization is enabled
 - getpid, getppid, gettid, tcgetpgrp, tcsetpgrp, getgrp, setpgrp, getsid, setsid, kill, tkill, tkill, wait, waitpid, waitid, wait3, wait4

```
...  
pid = getpid();  
...
```

User Program

```
int getpid(){  
...  
real_pid = funcs[_getpid]();  
return pid_table[real_pid];  
}
```

dmtcphijack.so

```
int getpid(){  
...  
}
```

libc.so



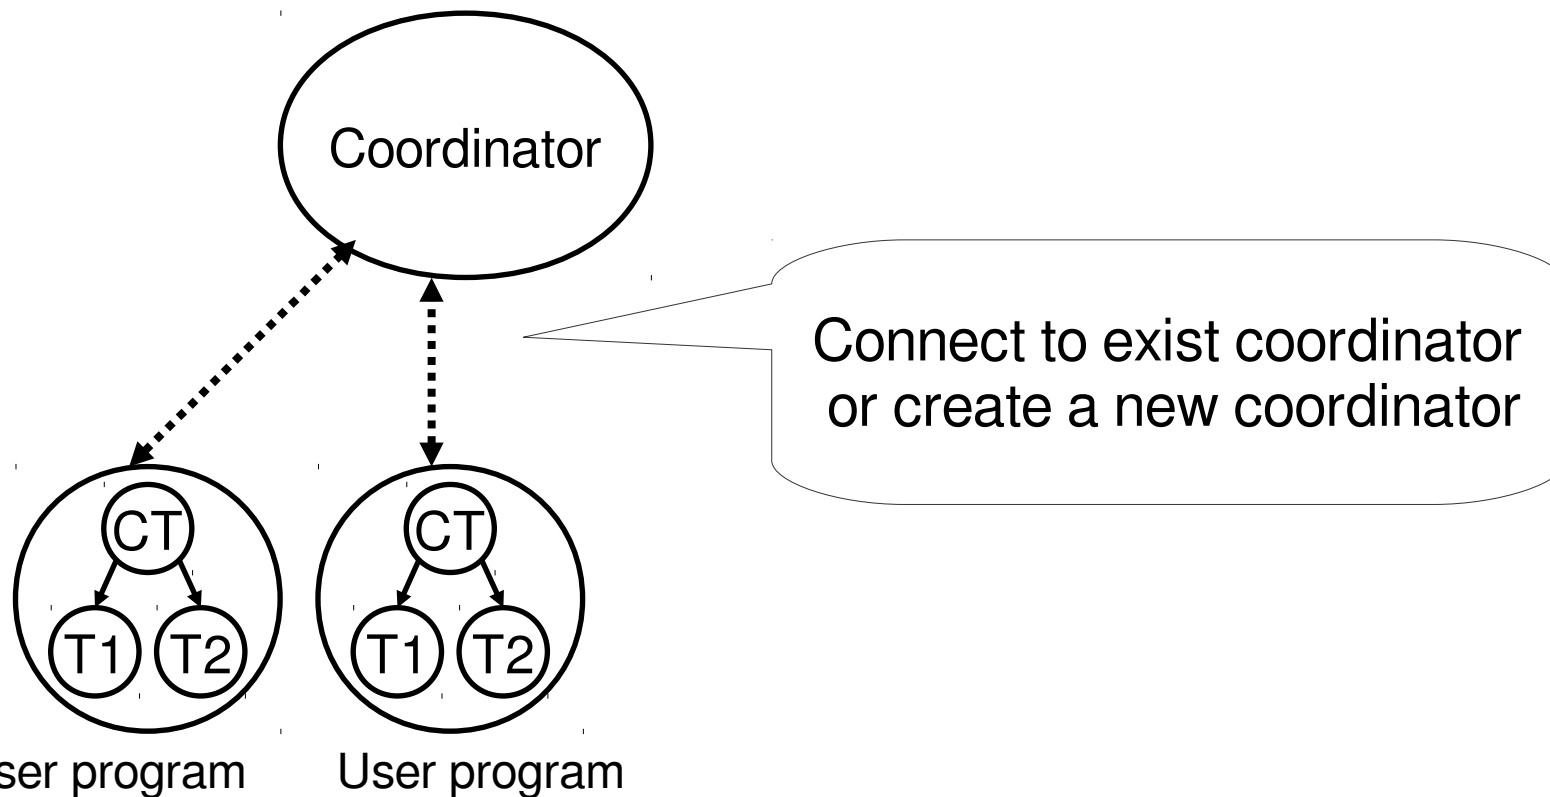
How DMTCP works (4/4)

- Checkpoint image compression on-the-fly (default).
- Currently only supports dynamically linking to libc.so. Support for static libc.a is feasible, but not implemented.



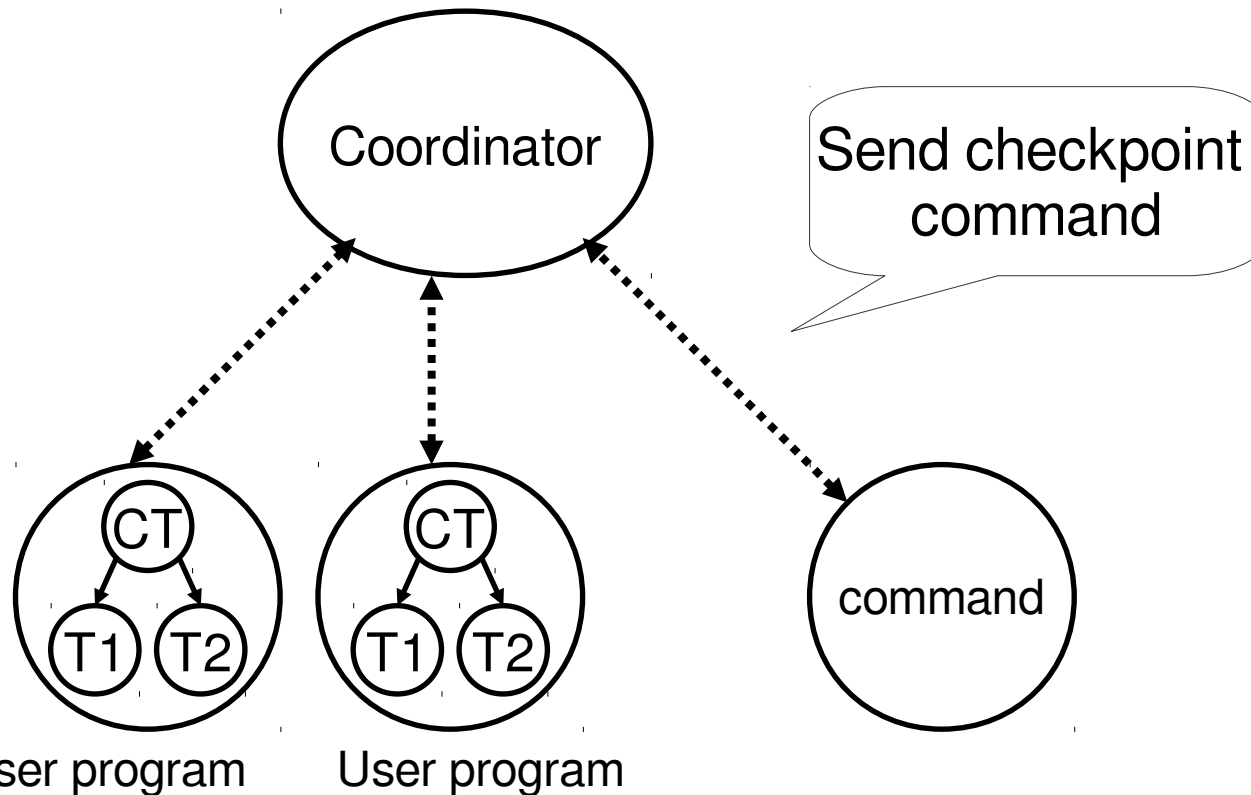
Checkpoint under DMTCP_(1/7)

- dmtcphijack.so and libmtcp.so present in executable's memory.
 - `dmtcp_checkpoint` <EXE>



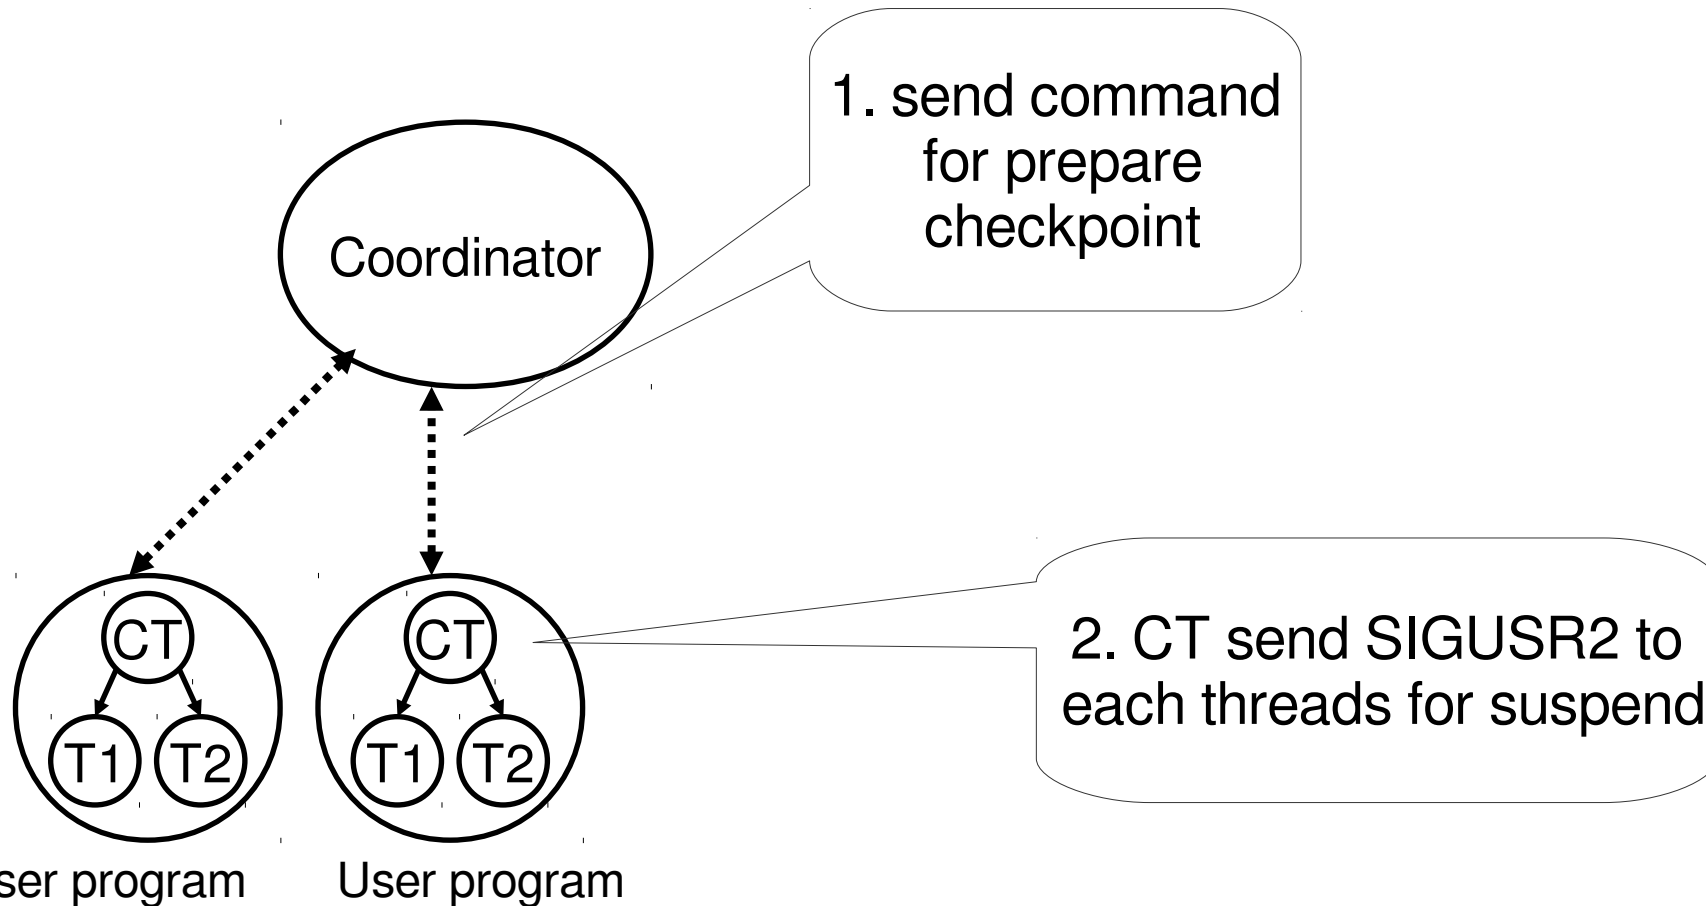
Checkpoint under DMTCP_(2/7)

- Ask coordinator process for checkpoint via `dmtcp_command`.
 - `dmtcp_command -c`
- DMTCP also provides API to send command or query status



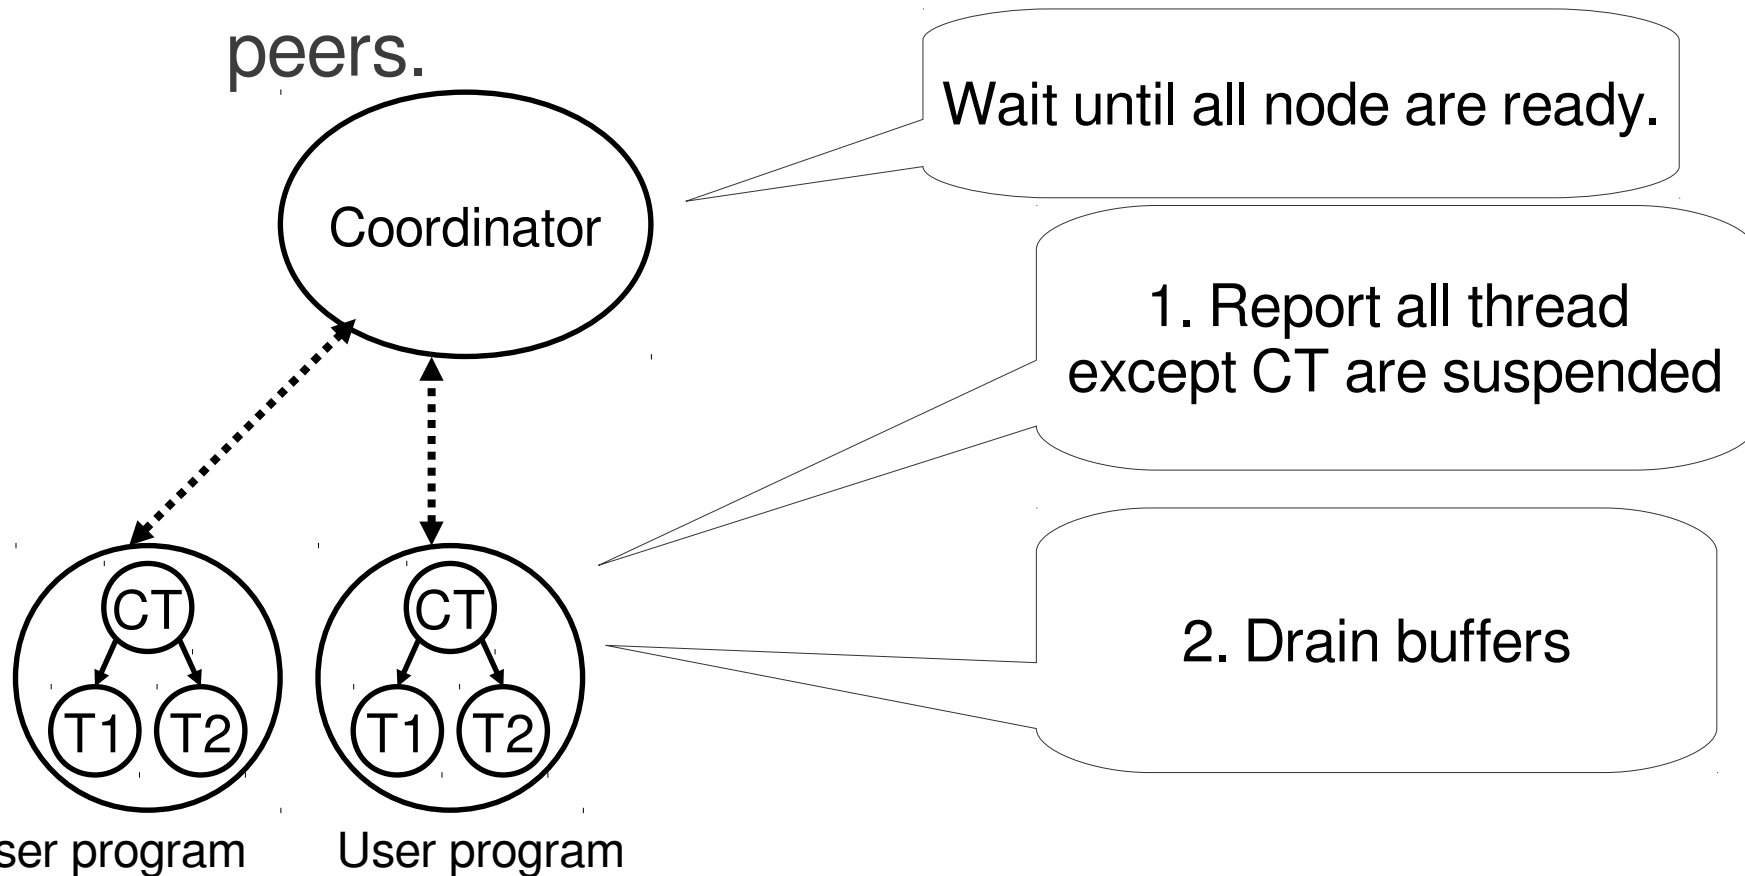
Checkpoint under DMTCP_(3/7)

- Suspend user threads with SIGUSR2.



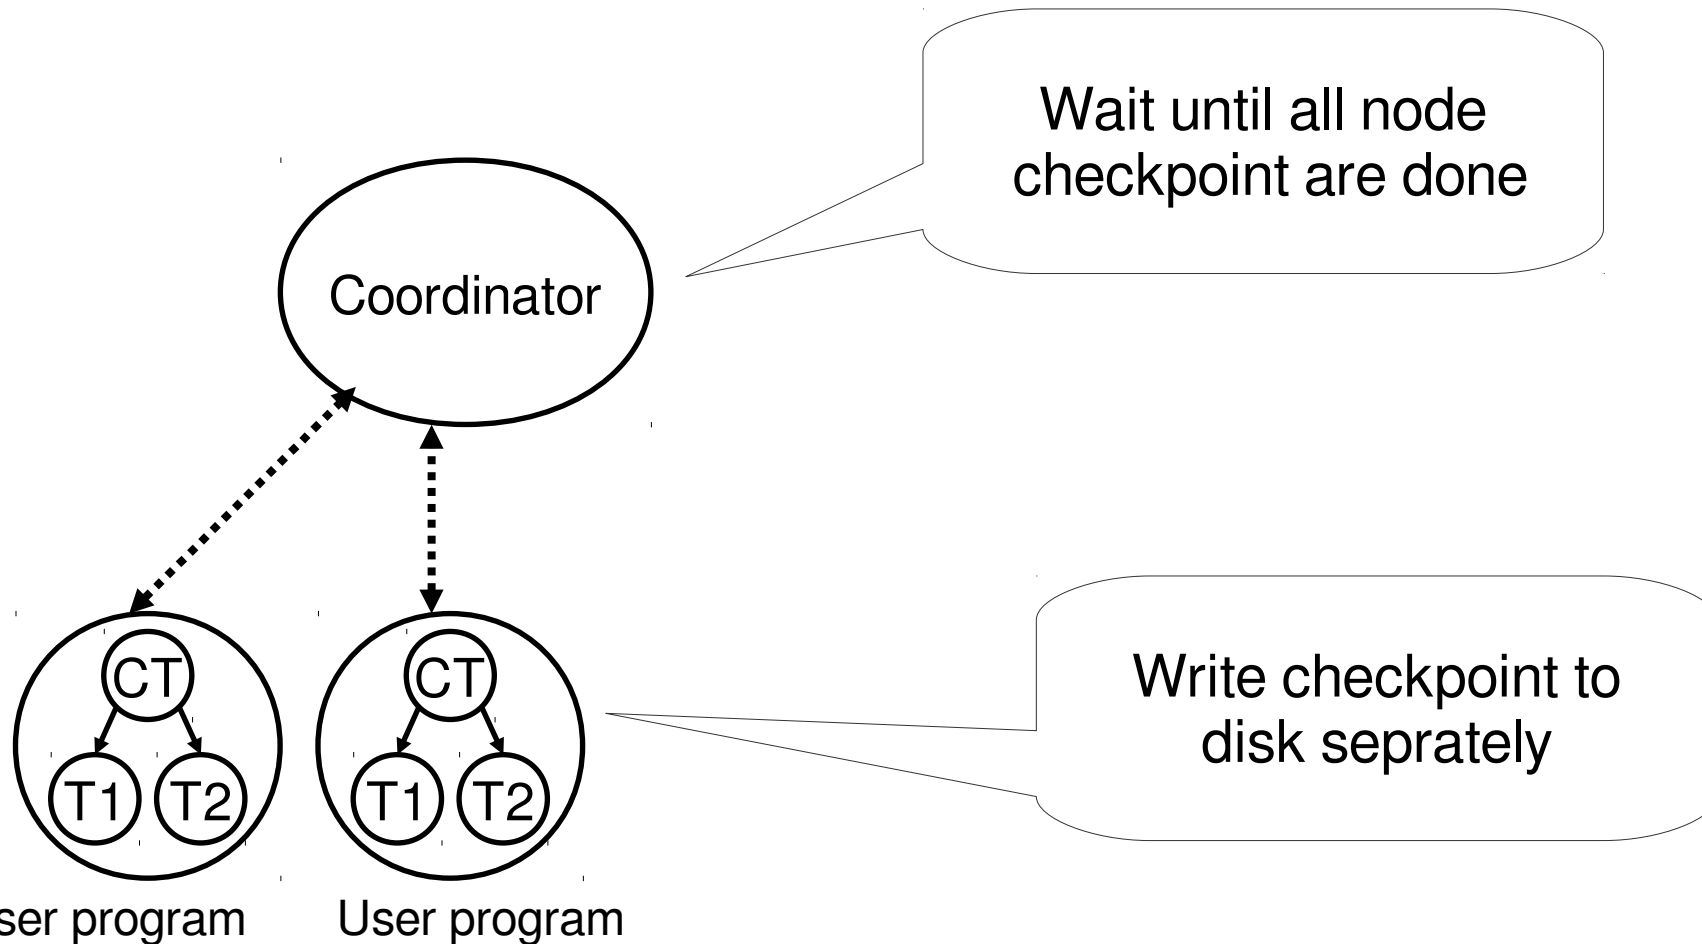
Checkpoint under DMTCP_(4/7)

- Pre-checkpoint stage
- Synchronize every node and elect shared file descriptor leaders.
- Drain kernel buffers and do network handshake with peers.



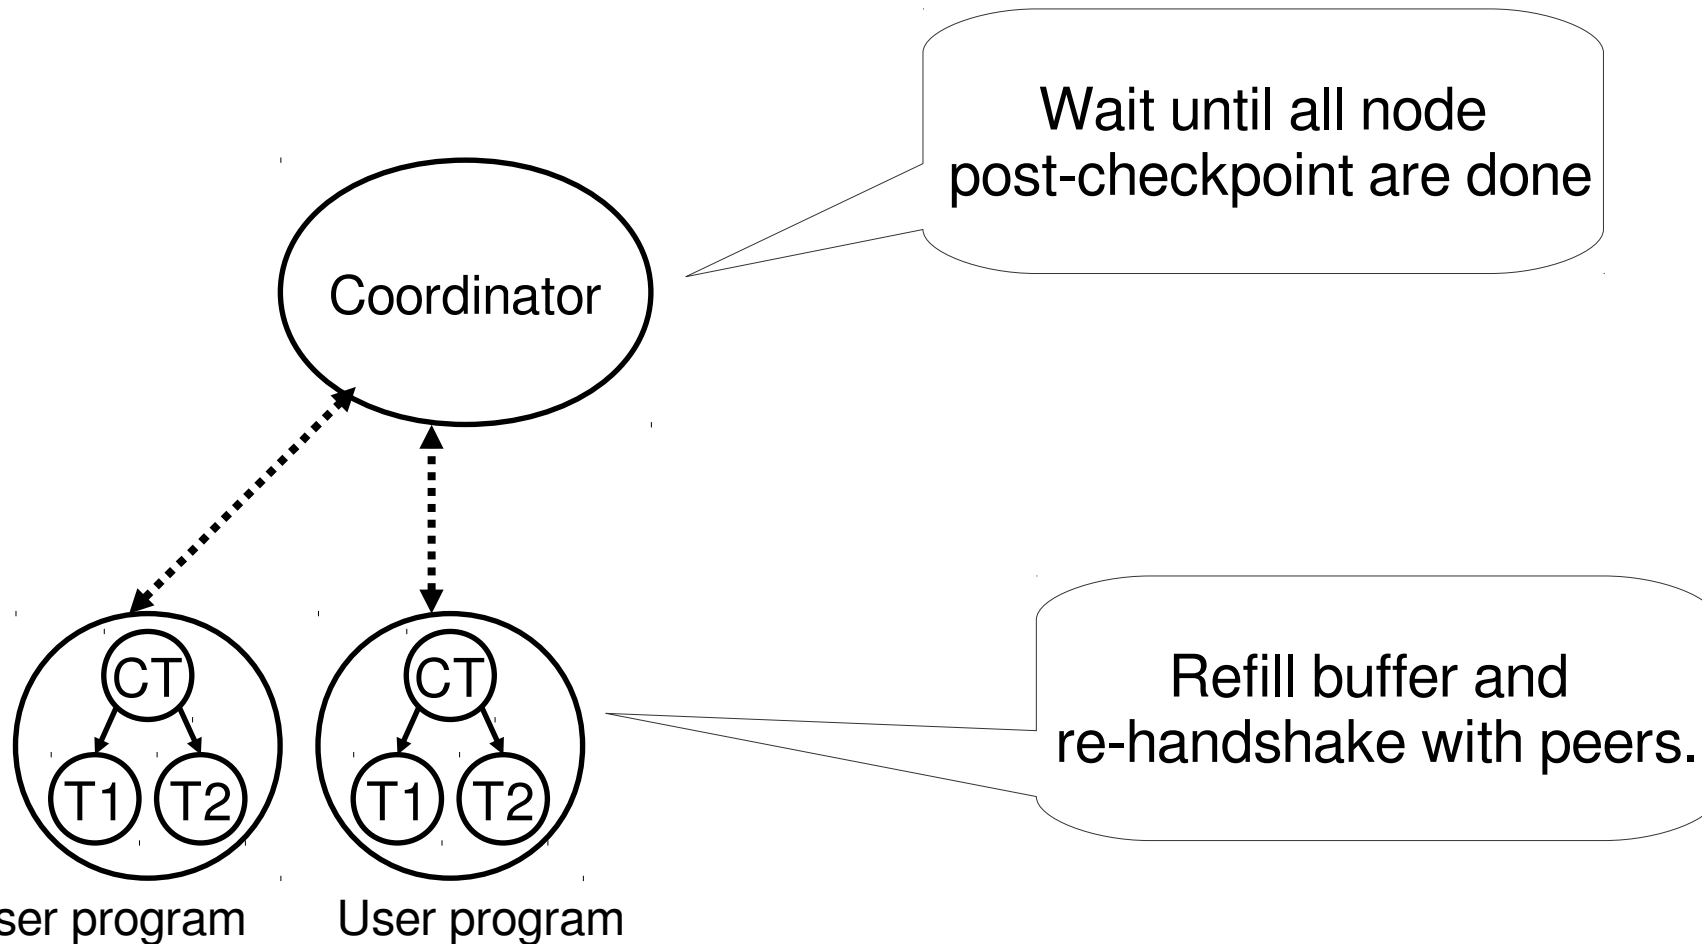
Checkpoint under DMTCP_(5/7)

- Write checkpoint to disk
 - One checkpoint file per process
 - `ckpt_<EXE>_<uid>.dmtcp`



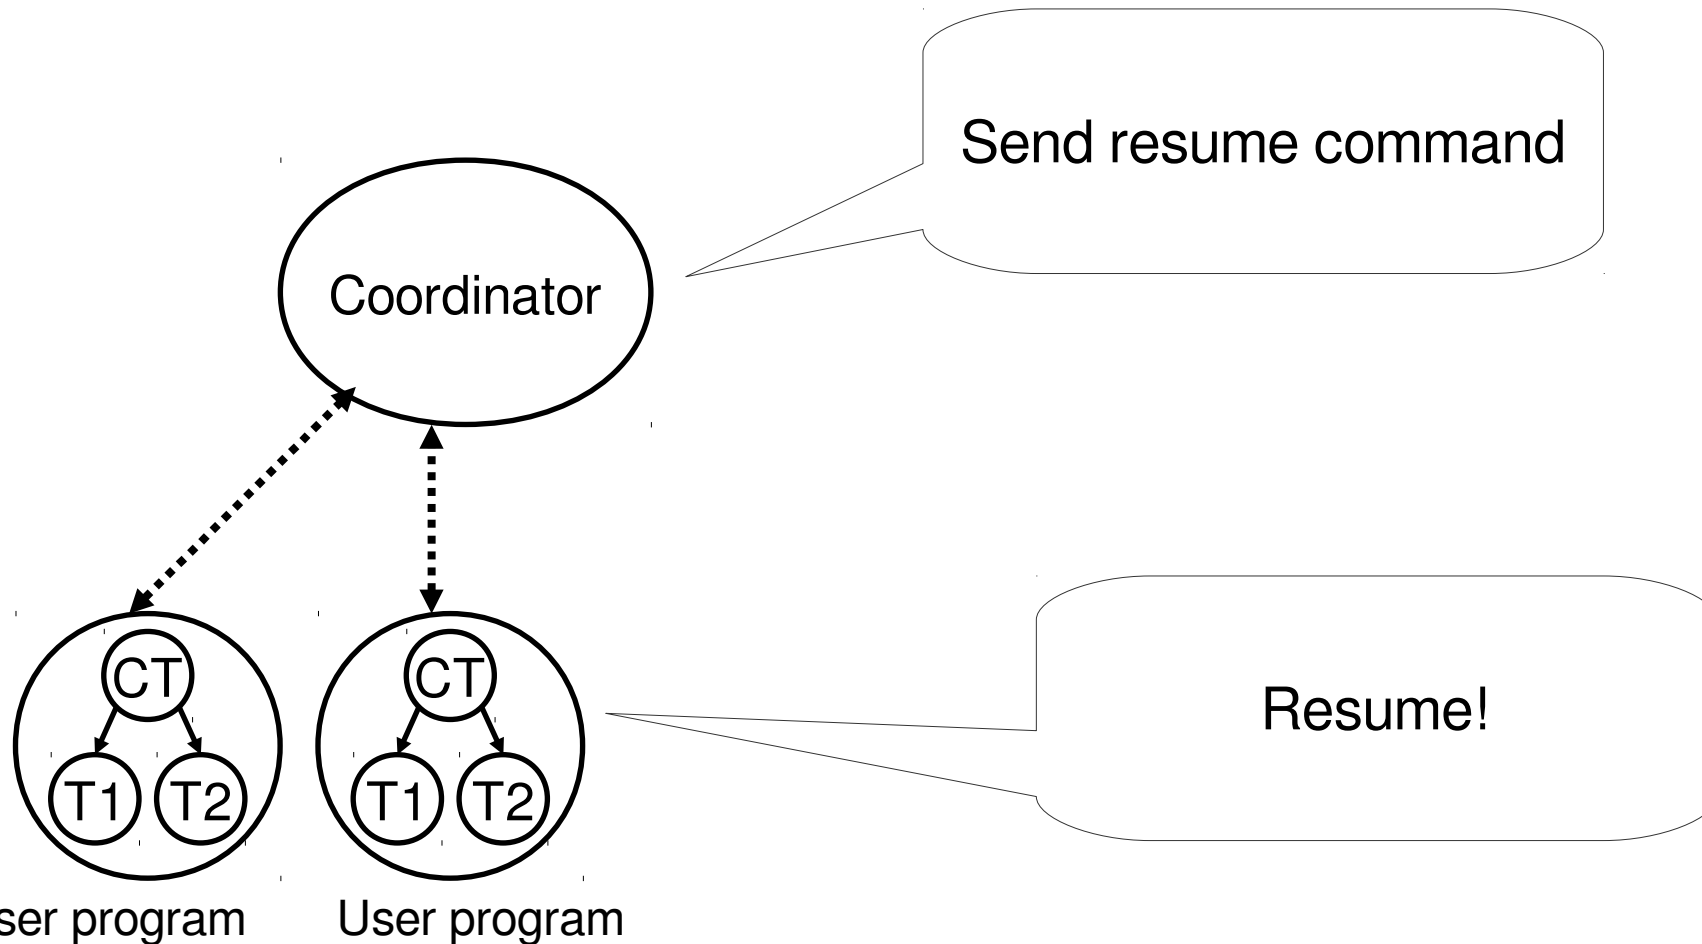
Checkpoint under DMTCP_(6/7)

- Post-Checkpoint stage
- Refill kernel buffers



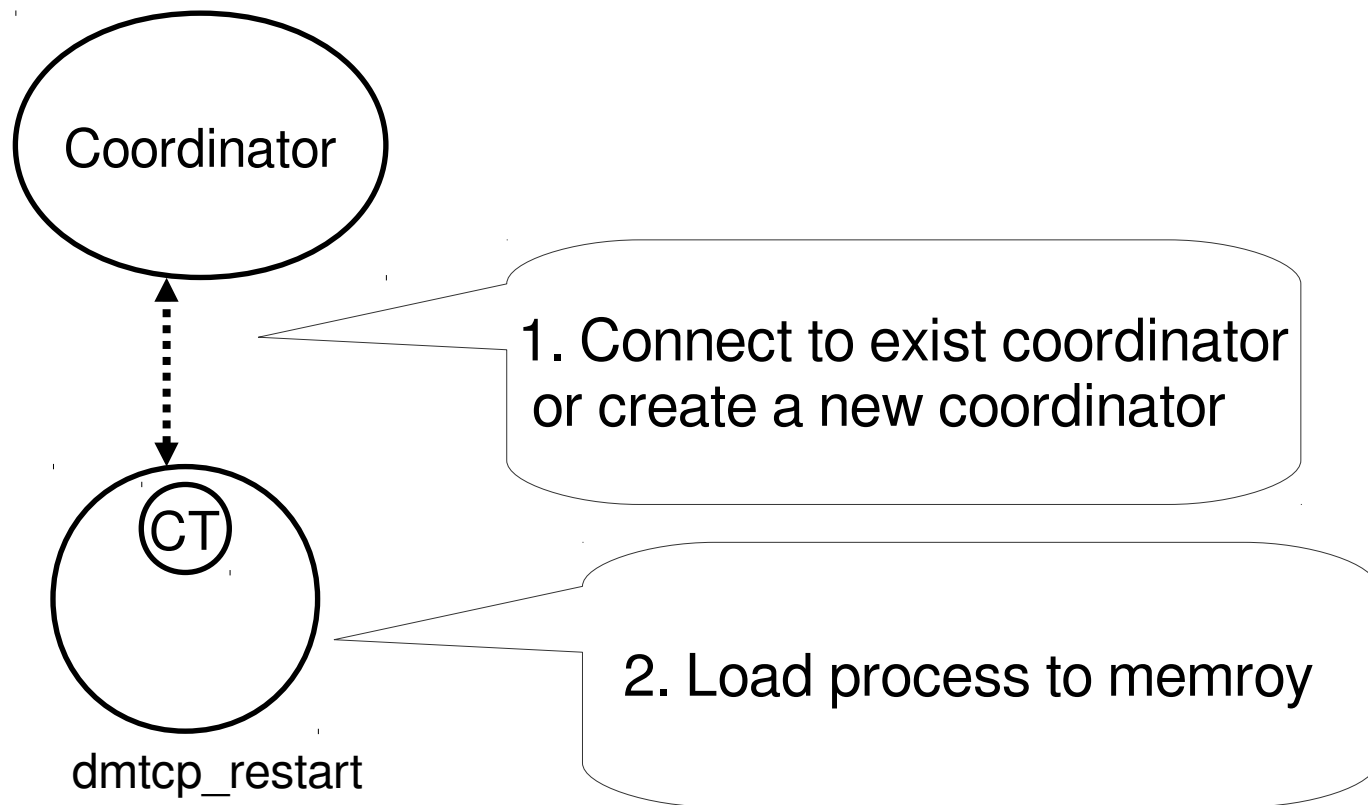
Checkpoint under DMTCP_(7/7)

- Resume user threads.




Restart under DMTCP_(1/6)

- Restart Process loads in memory.
 - `dmtcp_restart ckpt_<EXE>_<uid>.dmtcp`



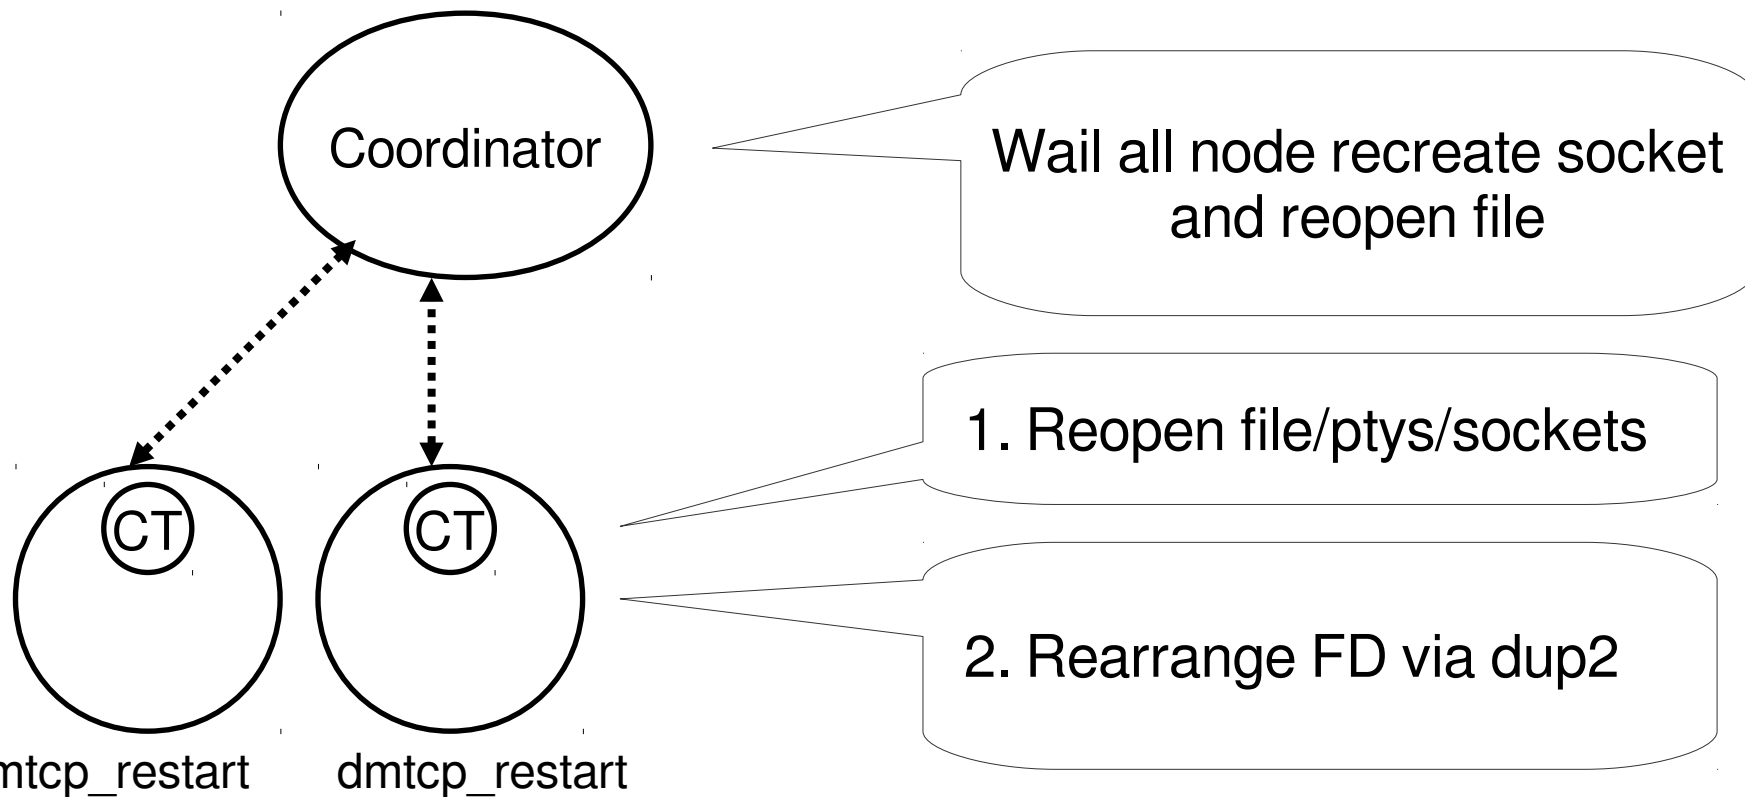
Restart under DMTCP_(2/6)

- Fork user program



Restart under DMTCP_(3/6)

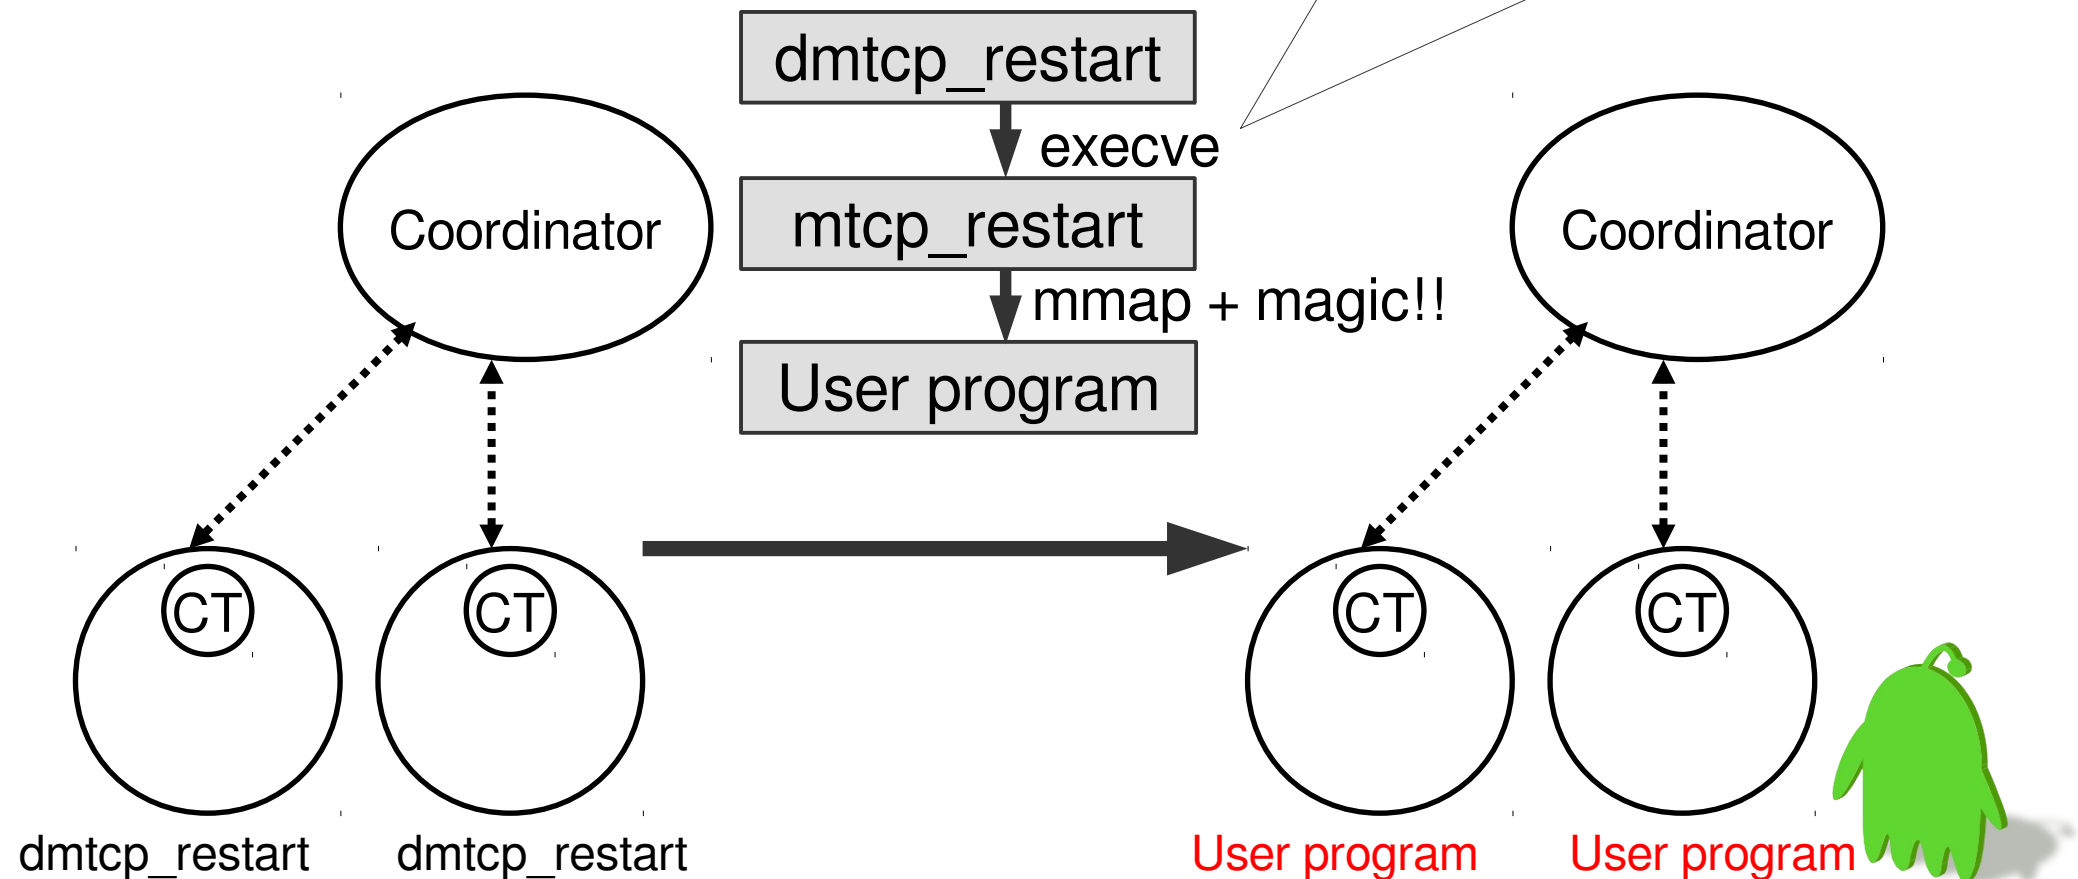
- Reopen files and recreate ptys
- Recreate and reconnect sockets
- Rearrange file descriptors to initial layout



Restart under DMTCP_(4/6)

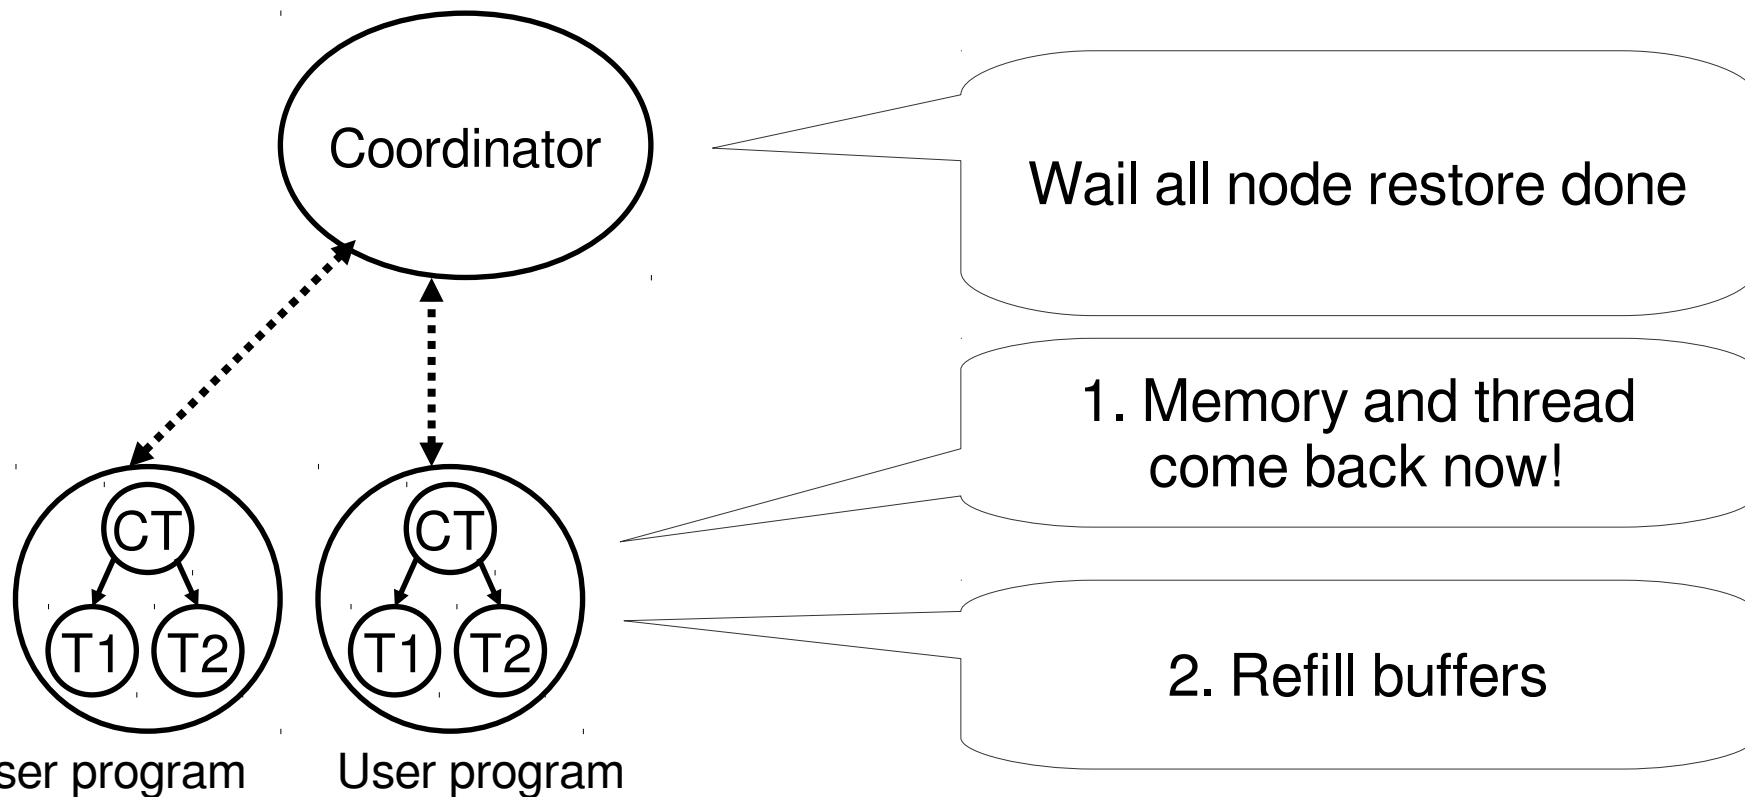
- Restore memory content.
- Restore stack status for checkpoint thread.

FDs will be preserved across execve!



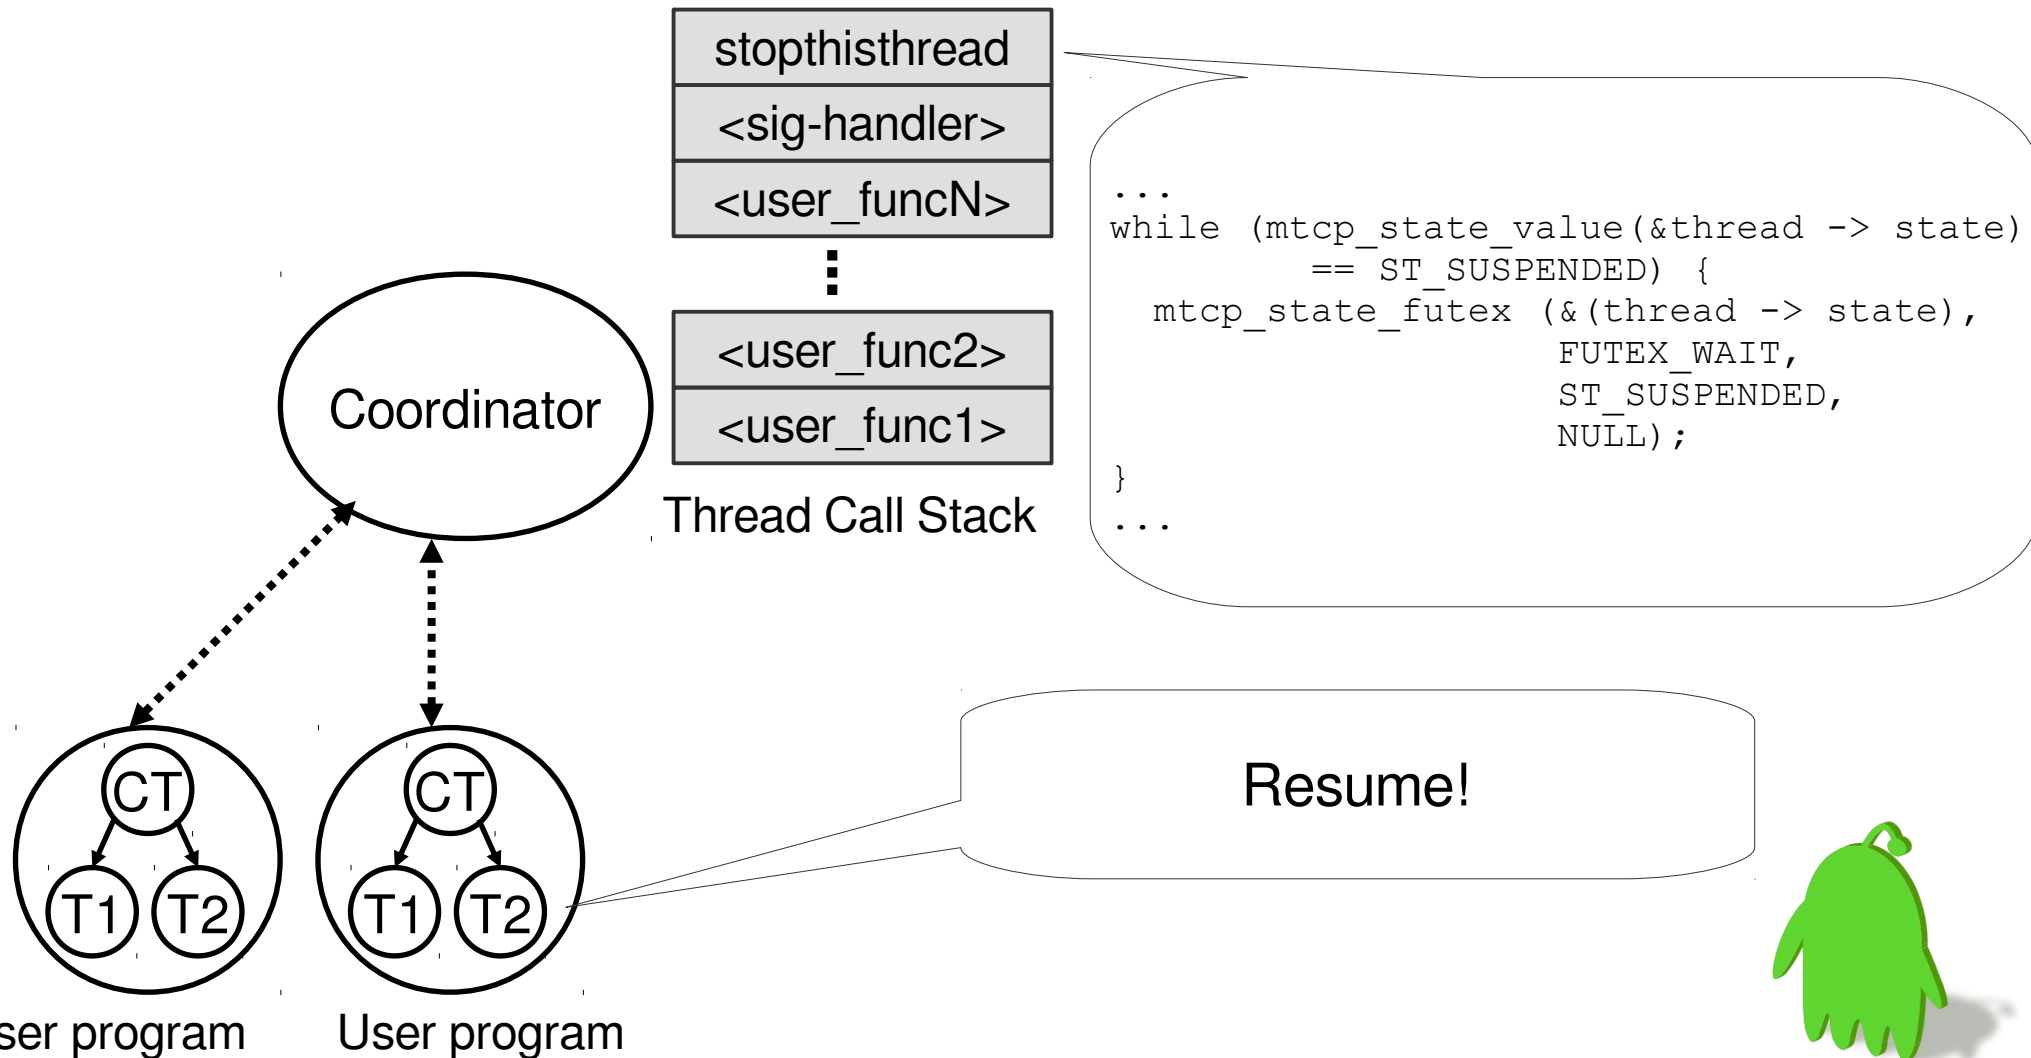
Restart under DMTCP_(5/6)

- Restore other threads.
 - Recreate thread and restore stack and context.
 - Restore back to the post-checkpoint stage
- Refill kernel buffer

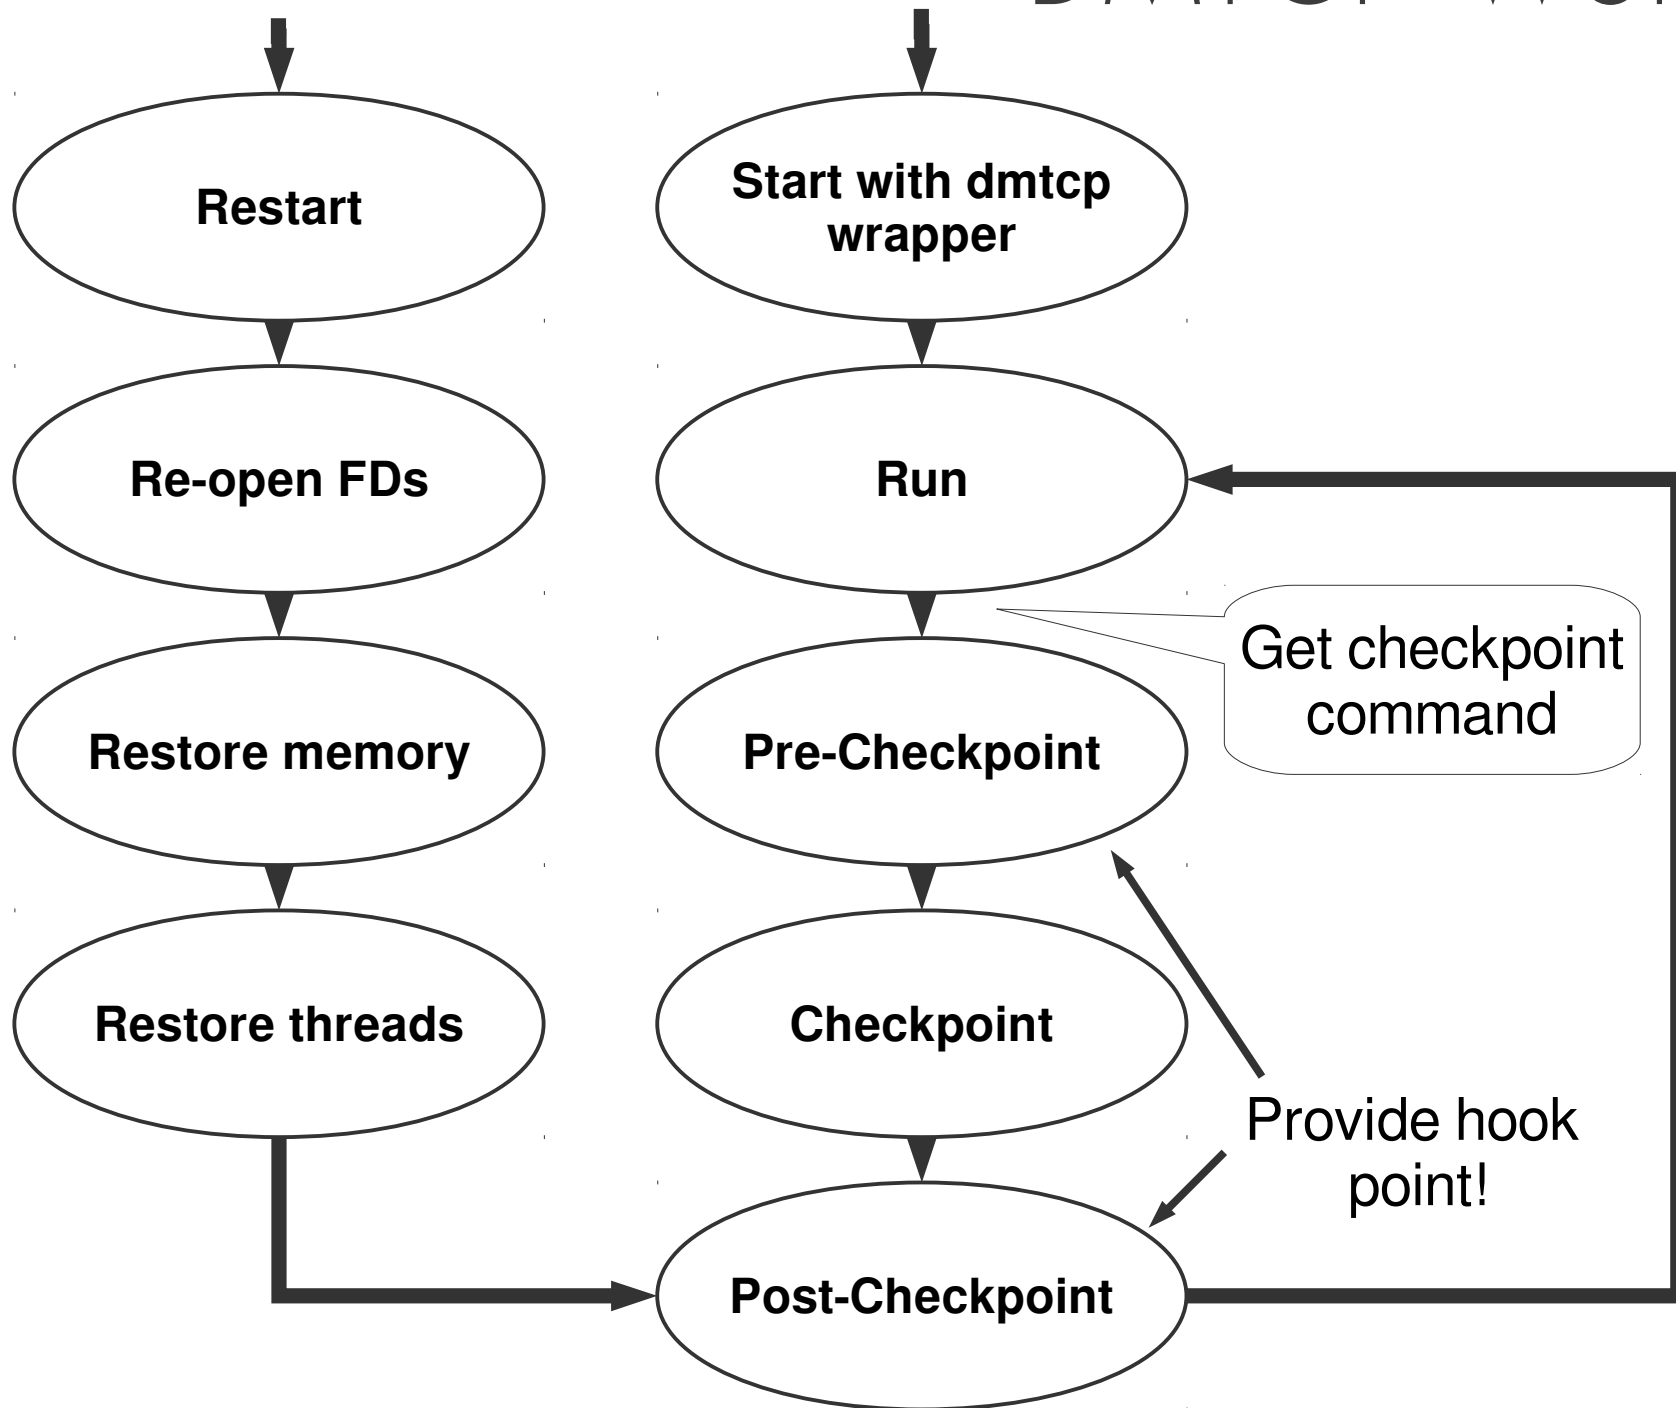


Restart under DMTCP_(6/6)

- Resume user threads



DMTCP Workflow



OS Features supported by DMTCP

- Threads, mutexes/semaphores, fork, exec
- Shared memory (via mmap), TCP/IP sockets, UNIX domain sockets, pipes, ptys, terminal modes, ownership of controlling terminals, signal handlers, open and/or shared fds, I/O (including the readline library), parent-child process relationships, process id & thread id virtualization, session and process group ids, and more...



DMTCP/Android: Additional Features

(LGPL; separated from Android)

- Binder IPC
 - Client-part: supported
 - Server-part: not supported
- Ashmem
 - Supported
- Logger
 - Supported
- Properties
 - Supported
- Wakelocks
 - Not supported



Support new FD type in DMTCP

- In DMTCP, every FD has a corresponding `Connection`
 - `TcpConnection`, `FileConnection`, `PtyConnection`
- Implement a new subclass of `Connection` if you want to support a new FD type for Android:
 - `AshmemConnection`, `BinderConnection`,
`LoggerConnection`, `PropertyConnection`
- And implement the `preCheckpoint`, `postCheckpoint`, and any others if needed.




Interface of Connection

```
class Connection {
public:
    virtual void preCheckpoint (const dmtcp::vector<int>&,
                                KernelBufferDrainer&);
    virtual void postCheckpoint (const dmtcp::vector<int>&,
                                bool);
    virtual void restore (const dmtcp::vector<int>&,
                          ConnectionRewirer *);
    virtual bool isDupConnection (const Connection&,
                                dmtcp::ConnectionToFds&);
    virtual void doLocking (const dmtcp::vector<int>&);
    virtual void saveOptions (const dmtcp::vector<int>&);
    virtual void restoreOptions (const dmtcp::vector<int>& );
    virtual void doSendHandshakes (const dmtcp::vector<int>&,
                                   const dmtcp::UniquePid&);
    virtual void doRecvHandshakes (const dmtcp::vector<int>&,
                                   const dmtcp::UniquePid&);
    virtual void restartDup2 (int, int);
protected:
    virtual void serializeSubClass (jalib::JBinarySerializer&);
};
```



Binder support for DMTCP

- BinderConnection
 - Reopen `/dev/binder` and reset option which is set via `ioctl`.
 - Restore the mmap region
- Hijack the whole libbinder!
 - Prevent libbinder from interpreting data twice
 - Implement necessary DMTCP hooks: `preCheckpoint`, `postCheckpoint`, `postRestart`
 - Re-initialize libbinder in `postRestart`
- The server part is not supported at the moment because binder server will call a blocked `ioctl` and block the whole checkpoint process.
 - We'll implement an early checkpoint stage to suspend such kind of thread soon.



Checkpoint for Zygote

- Experiment environment:
 - Android-x86 ICS emulated by VirtualBox
 - Host: Thinkpad x200 (Intel Core 2 Due @ 2.4 GHz)

	with gzip	without gzip
Checkpoint time	~10s	~5.5s
Restart time	~0.5s	~0.2s
Image size	~3M	~17M



Observations from logcat

<pre>----- beginning of /dev/log/system I/Vold (1270): Vold 2.1 (the revenge) firing up D/Vold (1270): Volume usb state changing -1 (Initializing) -> 0 (No-Media) I/Netd (1271): Netd 1.0 starting I/ (1275): ServiceManager: 0x8062b50 I/ (1276): ServiceManager: 0x804fb98 I/AudioFlinger(1276): Loaded primary audio interface from LEGACY Audio HAL I/AudioFlinger(1276): Using 'LEGACY Audio HW HAL' (audio.primary) as the ... D/AudioHardware(1276): ### setVoiceVolume: 1.000000 I/AudioPolicyService(1276): [1276]Loaded audio policy from LEGACY Audio HAL E/BatteryService(1382): usbOnlinePath not found D/AndroidRuntime(1902): D/AndroidRuntime(1902): >>>>> AndroidRuntime START com.android.systemui D/AndroidRuntime(1902): CheckJNI is ON I/SamplingProfilerIntegration(1902): Profiling disabled. I/Zygote (1902): Preloading classes... D/dalvikvm(1902): GC_EXPLICIT freed 35K, 85% free 399K/2560K, paused 0ms ... I/Zygote (1902): ...preloaded 379 resources in 548ms. D/dalvikvm(1902): GC_EXPLICIT freed 20K, 1% free 6417K/6467K, paused 0ms I/Zygote (1902): ...preloaded 31 resources in 13ms. D/dalvikvm(1902): GC_EXPLICIT freed 14K, 1% free 6418K/6467K, paused 0ms D/dalvikvm(1902): GC_EXPLICIT freed 5K, 1% free 6412K/6467K, paused 0ms D/dalvikvm(1902): GC_EXPLICIT freed <1K, 1% free 6412K/6467K, paused 0ms I/dalvikvm(1902): System server process 1911 has been created</pre>	<pre>----- beginning of /dev/log/system I/Vold (1270): Vold 2.1 (the revenge) firing up D/Vold (1270): Volume usb state changing -1 (Initializing) -> 0 (No-Media) I/Netd (1271): Netd 1.0 starting I/ (1275): ServiceManager: 0x8062b50 I/ (1276): ServiceManager: 0x804fb98 I/AudioFlinger(1276): Loaded primary audio interface from LEGACY Audio HAL I/AudioFlinger(1276): Using 'LEGACY Audio HW HAL' (audio.primary) as the D/AudioHardware(1276): ### setVoiceVolume: 1.000000 I/AudioPolicyService(1276): [1276]Loaded audio policy from LEGACY Audio HAL D/dalvikvm(1373): GC_EXPLICIT freed 14K, 1% free 6418K/6467K, paused 0ms D/dalvikvm(1373): GC_EXPLICIT freed 5K, 1% free 6412K/6467K, paused 0ms D/dalvikvm(1373): GC_EXPLICIT freed <1K, 1% free 6412K/6467K, paused 0ms I/dalvikvm(1373): System server process 1382 has been created</pre>
--	---

Normal bootup log message

Bootup log message with restart



Android Boottime with DMTCP

	1	2	3	4	5	6	7	8	9	10	avg
Boottime	27.96	27.95	32.89	26.59	32.33	32.36	33.22	32.99	36.47	32.85	31.56
Boottime with DMTCP (w/ gzip)	15.02	15.86	15.13	14.88	14.57	14.43	14.73	14.22	13.97	14.61	14.74
Boottime with DMTCP (w/o gzip)	14.98	15.13	14.61	13.90	14.72	14.84	15.46	15.06	15.32	15.39	14.94

Measured by uptime in `onStart()` of Launcher2



Reversible Debugger based on DMTCP

- URDB is a universal reversible debugger that can add reversibility gained through the use of DMTCP.
- A user debugging with gdb would then ask URDB to go back to a point in time to when:
 - the expression had a correct value
 - the next statement would cause the expression to take on an incorrect value.
- Project page: <http://sourceforge.net/projects/urdb/>



Reference

- “DMTCP: An New Linux Checkpointing Mechanism for Vanilla Universe Job”, Condor Project, University of Wisconsin-Madison
- “Checkpointing using DMTCP, Condor, Matlab and FReD”, Gene Cooperman Northeastern University, Boston
- URDB: Universal Reversible Debugger
<http://arxiv.org/abs/0910.5046>
- Cruz: Application-Transparent Distributed Checkpoint-Restart on Standard Operating Systems, G. (John) Janakiraman, Jose Renato Santos, Dinesh Subhraveti, Yoshio Turner, HP Labs





<http://0xlab.org>