



實作輕量級 RTOS 網路堆疊

Apr 24, 2009

Jim Huang(黃敬群 /jserv)

Email: <jserv.tw@gmail.com>

Blog: <http://blog.linux.org.tw/jserv/>

自我介紹

Phone/PDA, GPS, Mobile TV/Digital TV 代工設計

參與自由軟體開發 / 社群組織

新酷音輸入法

Kaffe.org

GNU Classpath

FreeDesktop

OpenMoko

...

Xenomai

Orz Microkernel

pcmanx

LXDE

OpenAVS

...

Debian Taiwan

KDE Taiwan

TOSSUG

0xlab (4/27)

警告：本議程僅探討實務面，忽略理論部份

提綱

動機

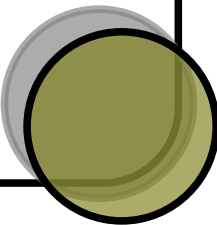
Bell 定律：每十年運算型態的移轉

RTOS 與嵌入式裝置的角色

實現 TCP/IP 的難題

回歸現實：剪裁與調適

技術討論



每年寫一套作業系統 當作業 (對不起，遲交)

JK (2001)

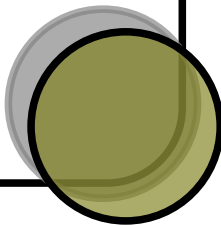
Orz Microkernel (2006)

RT nanokernel (2007)

Jamei RTOS (2007)

CuRT (2009)

??? (2009)



動機 (1)

Everything can be Orz.

Orz Microkernel



```
QEMU
sudZUZMZH2o=
jWZZ2!!"---!!HWWa [Orz/Microkernel]
<wdP" -!VZL,
.MN2" Zaaa NZL.
oZl idRV!?"SHWa 1Kh;
He .JH2C "HW: )HMc
.22" 1XC. nV: 1oZC,
.2H; )3k; s!" jHfA
12> -1Hb/ "H2C
-Zo; +!4ZMaaaaUZZWY
xHf, ~-?!!!!!!~"
HUb;
)VXL,,
+3HbC,
-1Orz,,
~~~~~

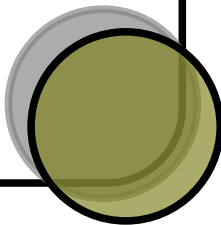
Orz Microkernel x86 pre-alpha (2006-06-01)
[jseru@sexmachine /]M
```



Orz Microkernel 的啓發

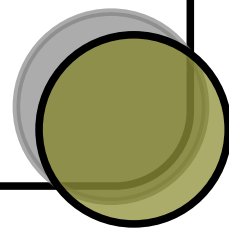
學習作業系統與相關的系統程式
該如何設計

建立自信：原來一個作業系統
只需幾 kb 的空間就實做出來



設計作業系統也 可很有趣

以實體的機器人設計作為主軸
體驗如何親手打造嵌入式系統
並著手設計相關軟硬體建設



從零到有，設計即時作業系統

杜威博士：「作中學」

RT nanokernel (OSDC.tw 2007)

Jamei (COSCUP 2007)

模仿 Linux 經典設計並建構具體而微的
RTOS

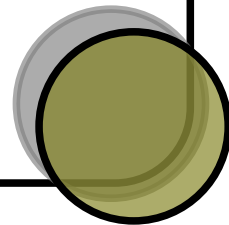


滿足自動控制系統需求

即時處理

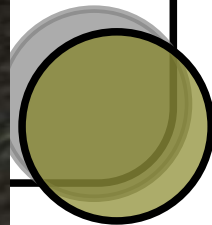
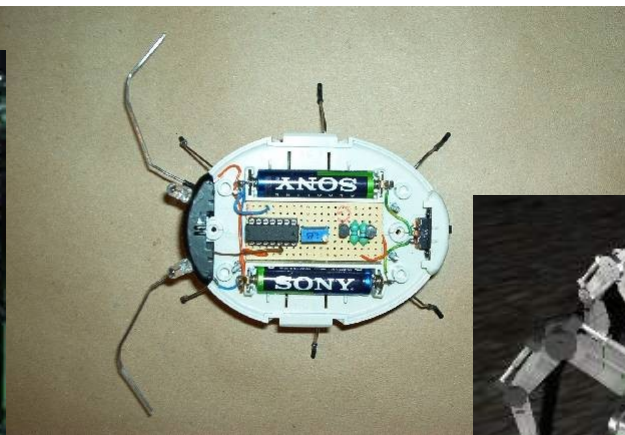
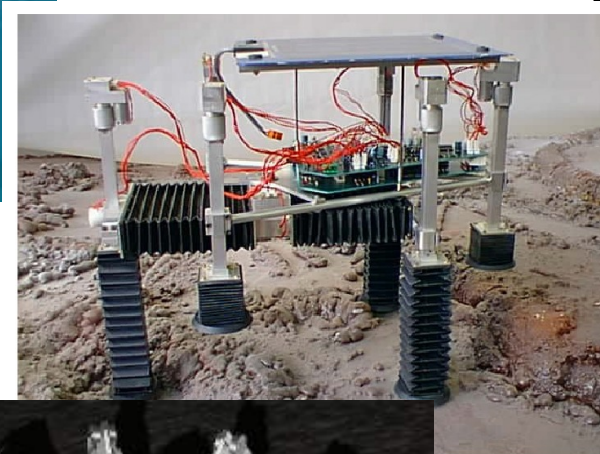
建構嵌入式環境

網路通訊能力



應用型態

國防軍事
科學搜救探索
仿生生物型態
機器人足球賽



簡化設計，適用更廣 範圍的硬體

CuRT (2009)

硬體：Marvell/Intel PXA255

特徵

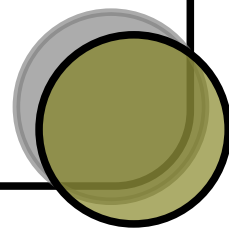
- Preemptive Multi-threading

- Priority-base Round-Robin Scheduling

- Thread Management

- Semaphore Management Support

- IPC: mailbox, message queue



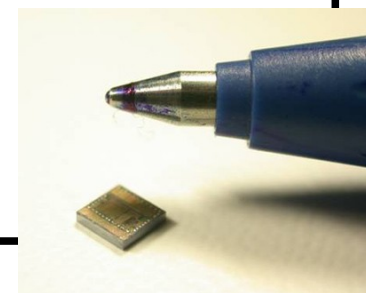
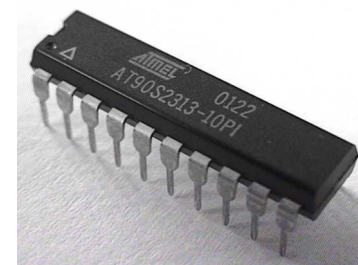
Bell 定律：每十年運算 型態的移轉

1980 年代：PC revolution

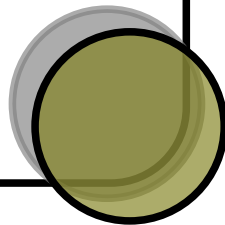
1990 年代：Internet revolution

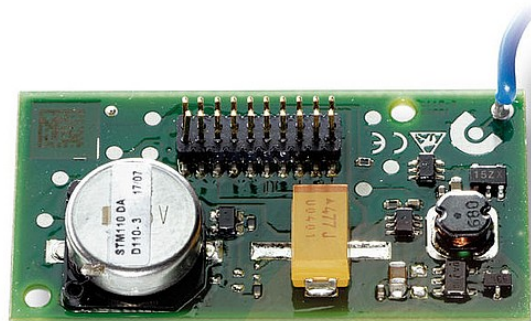
2000 年代：embedded revolution

2010 年代：embedded Internet revolution



RTOS 與嵌入式 裝置





enoclean.com



Building automation

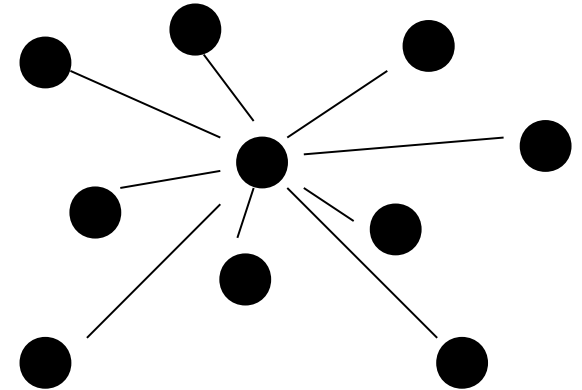


streetlinenetworks.com

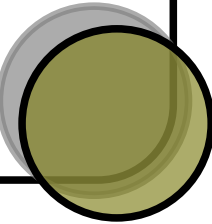
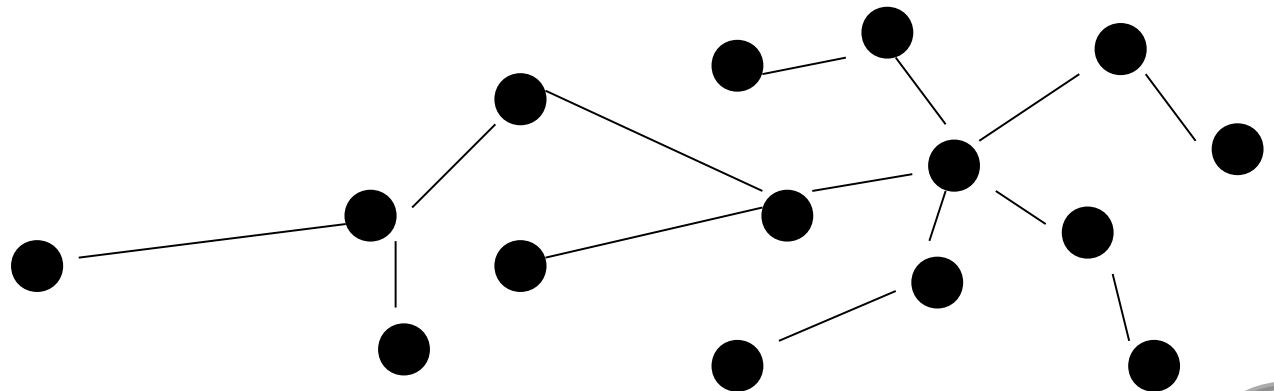
Wireless parking management

嵌入式無線網路

Star networks

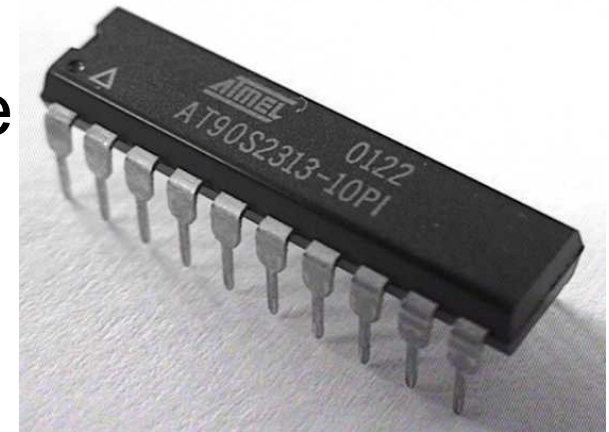


Mesh networks



典型嵌入式硬體

- Small microcontroller (microprocessor + I/O)
 - 8- and 16-bit processors
 - Typical memory size
 - 1 – 100 k Flash ROM for code
 - 0.1 – 10 k RAM for data
 - Typical speed
 - 1 – 10 MHz
- 8051, AVR, MSP430, Z80, 6502, ARM, ...



設計自己的 RTOS

Jamei = **J**ust **A**nother **M**icroprocessor **E**MBEDDED
Infrastructure

輕巧並可調整組態

仿造 Linux kernel 部份設計

arch(平台相依實做)

device driver model

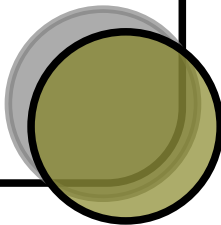
vfs

部份 POSIX Thread

部份 Realtime API (IEEE 1003.1b)

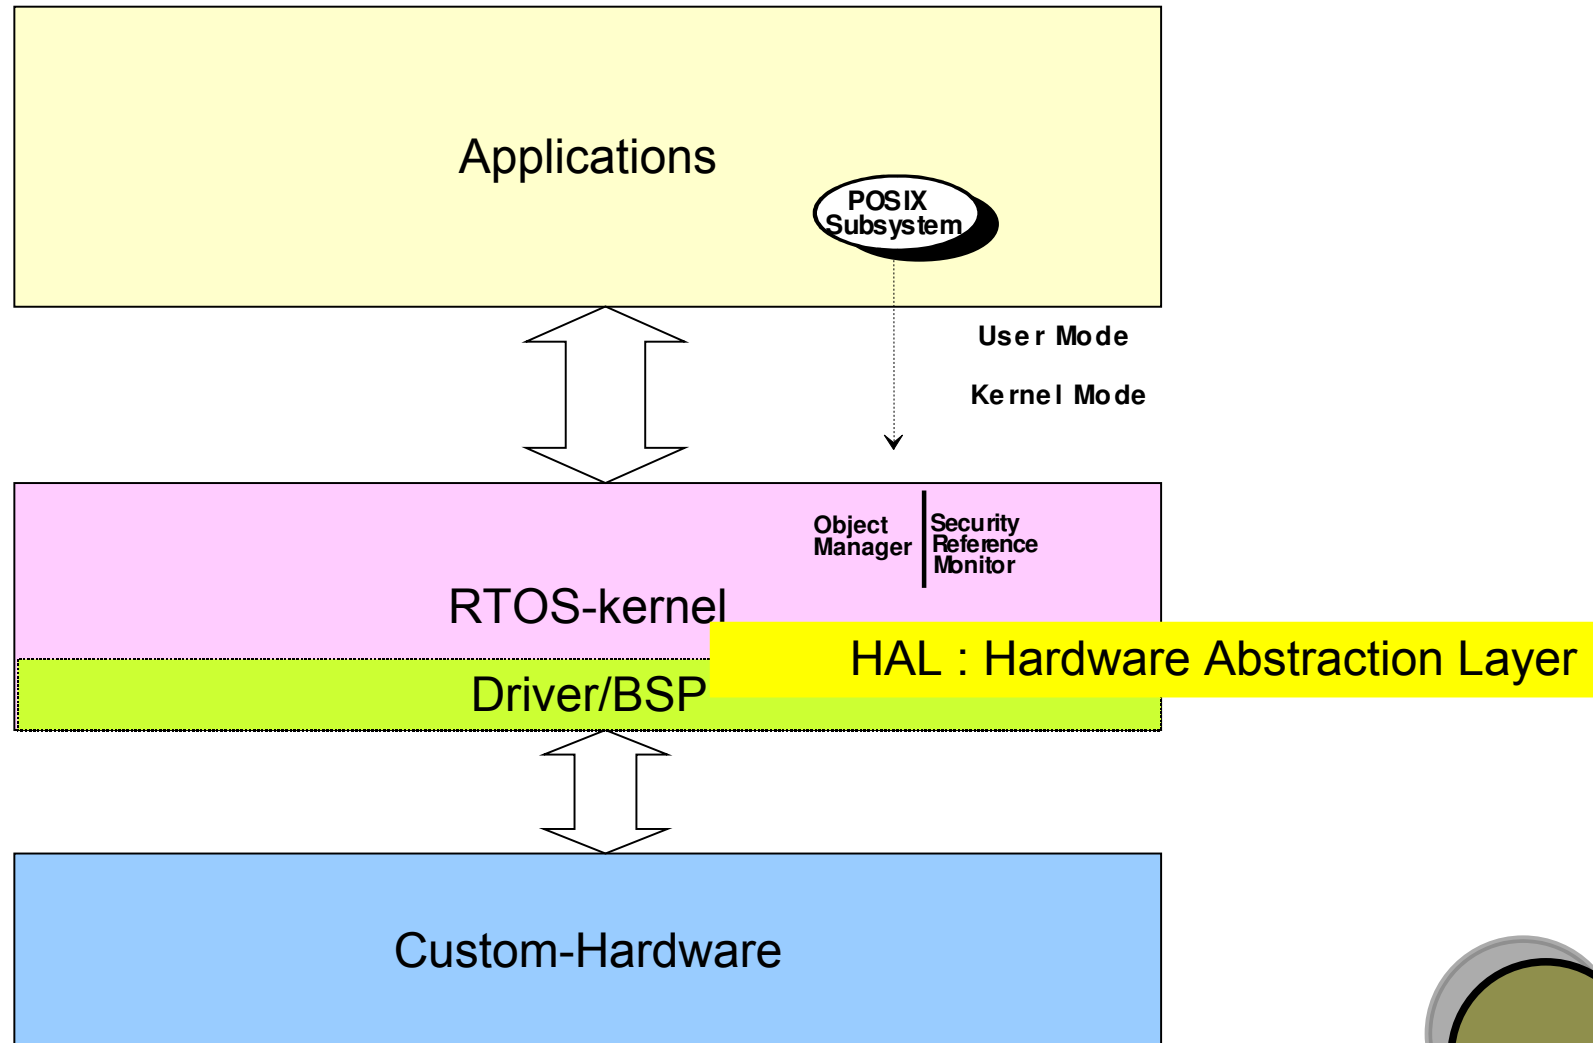
New BSD License 與 GNU GPLv2(部份)

支援 i386 與 ARM9



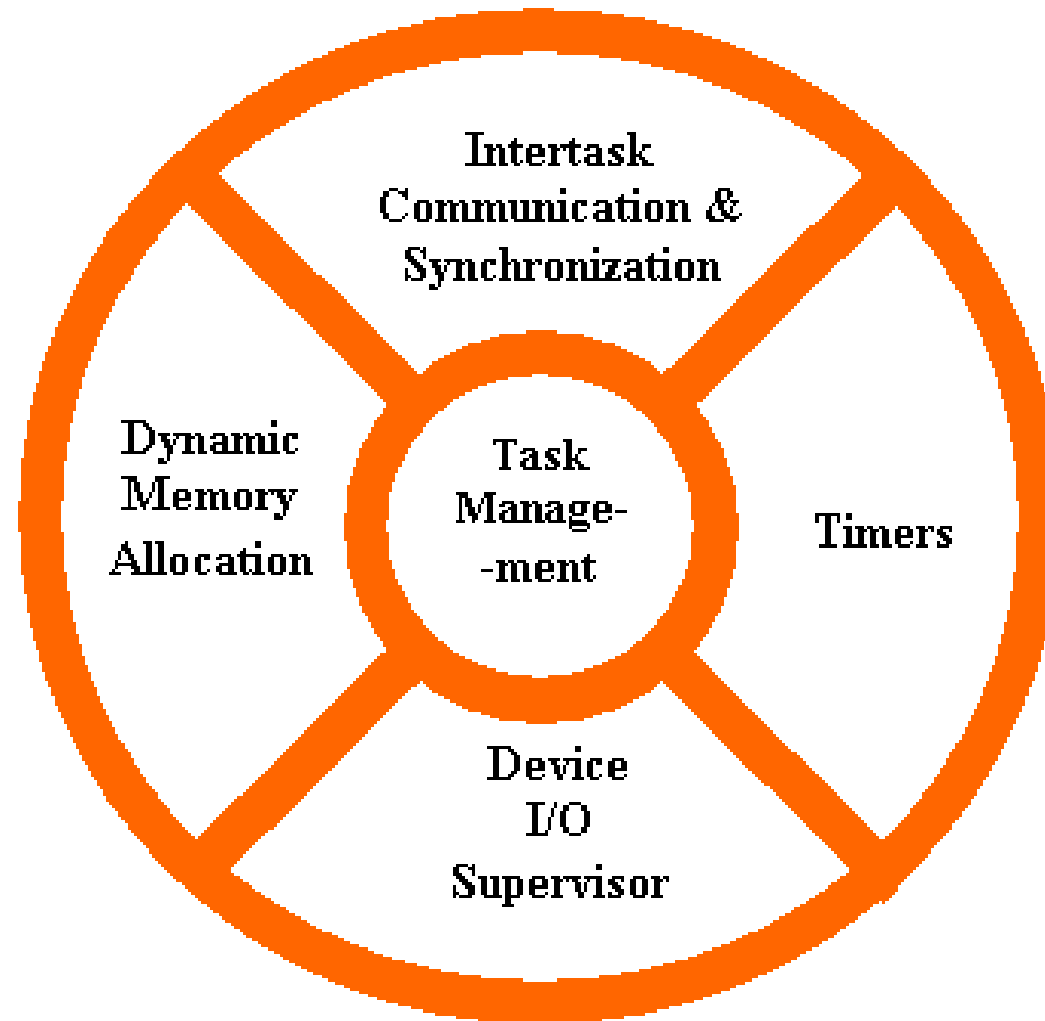
RTOS 結構

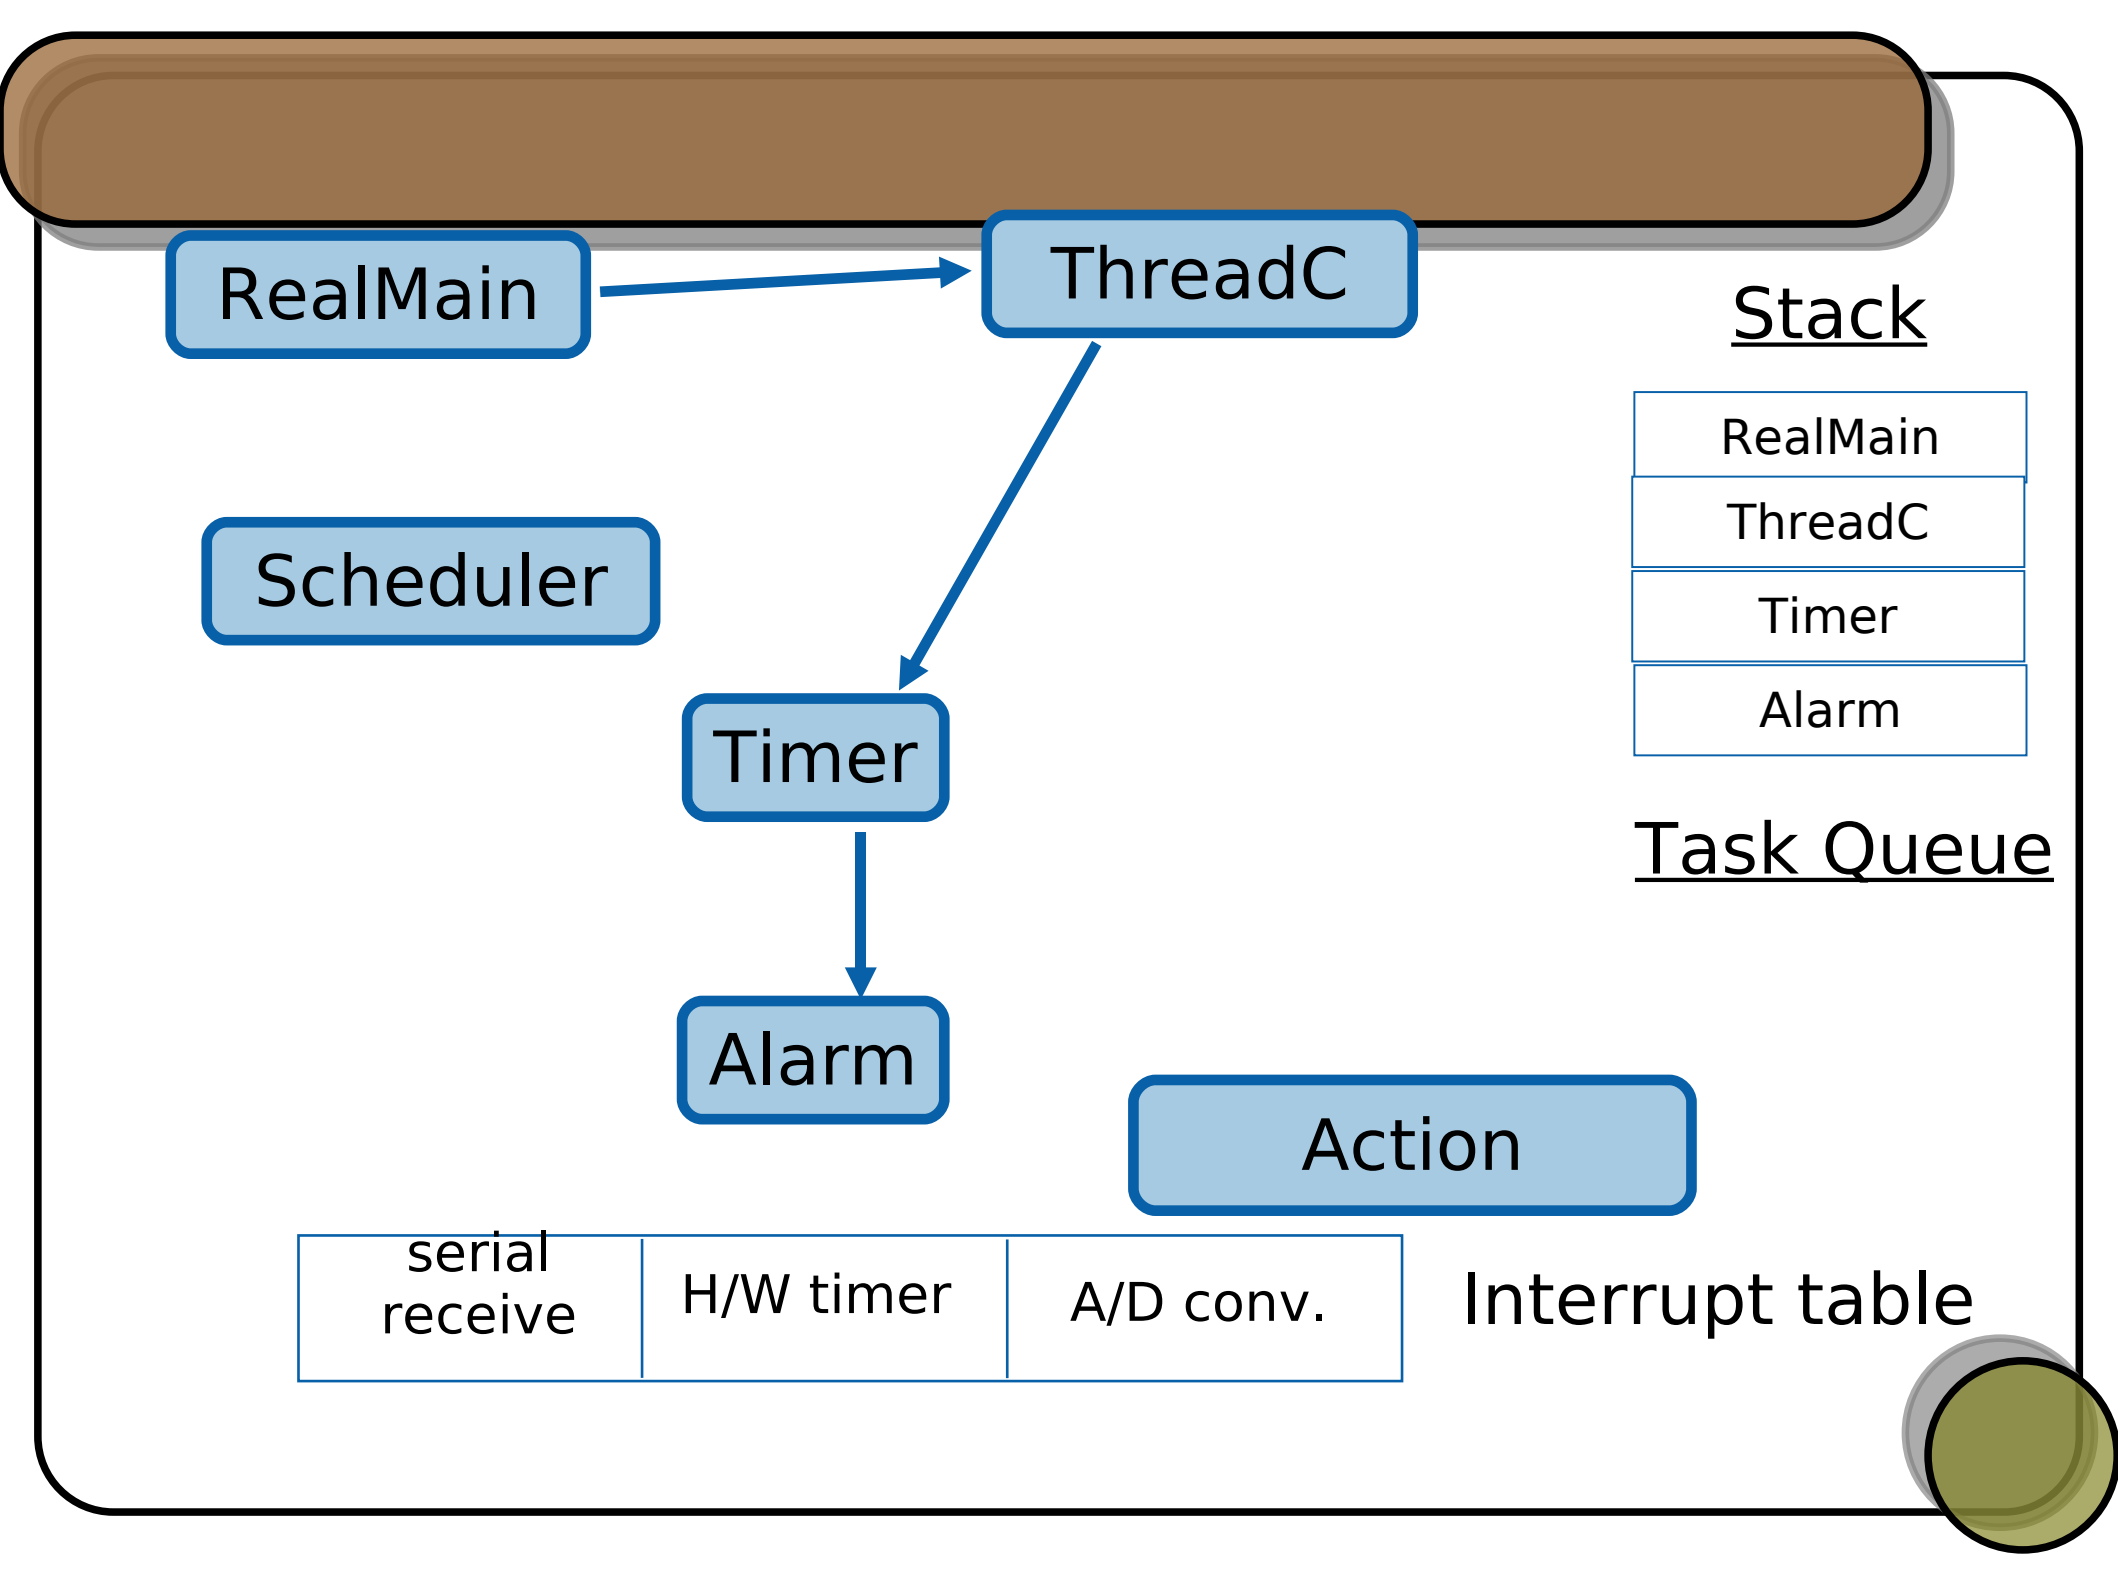
- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

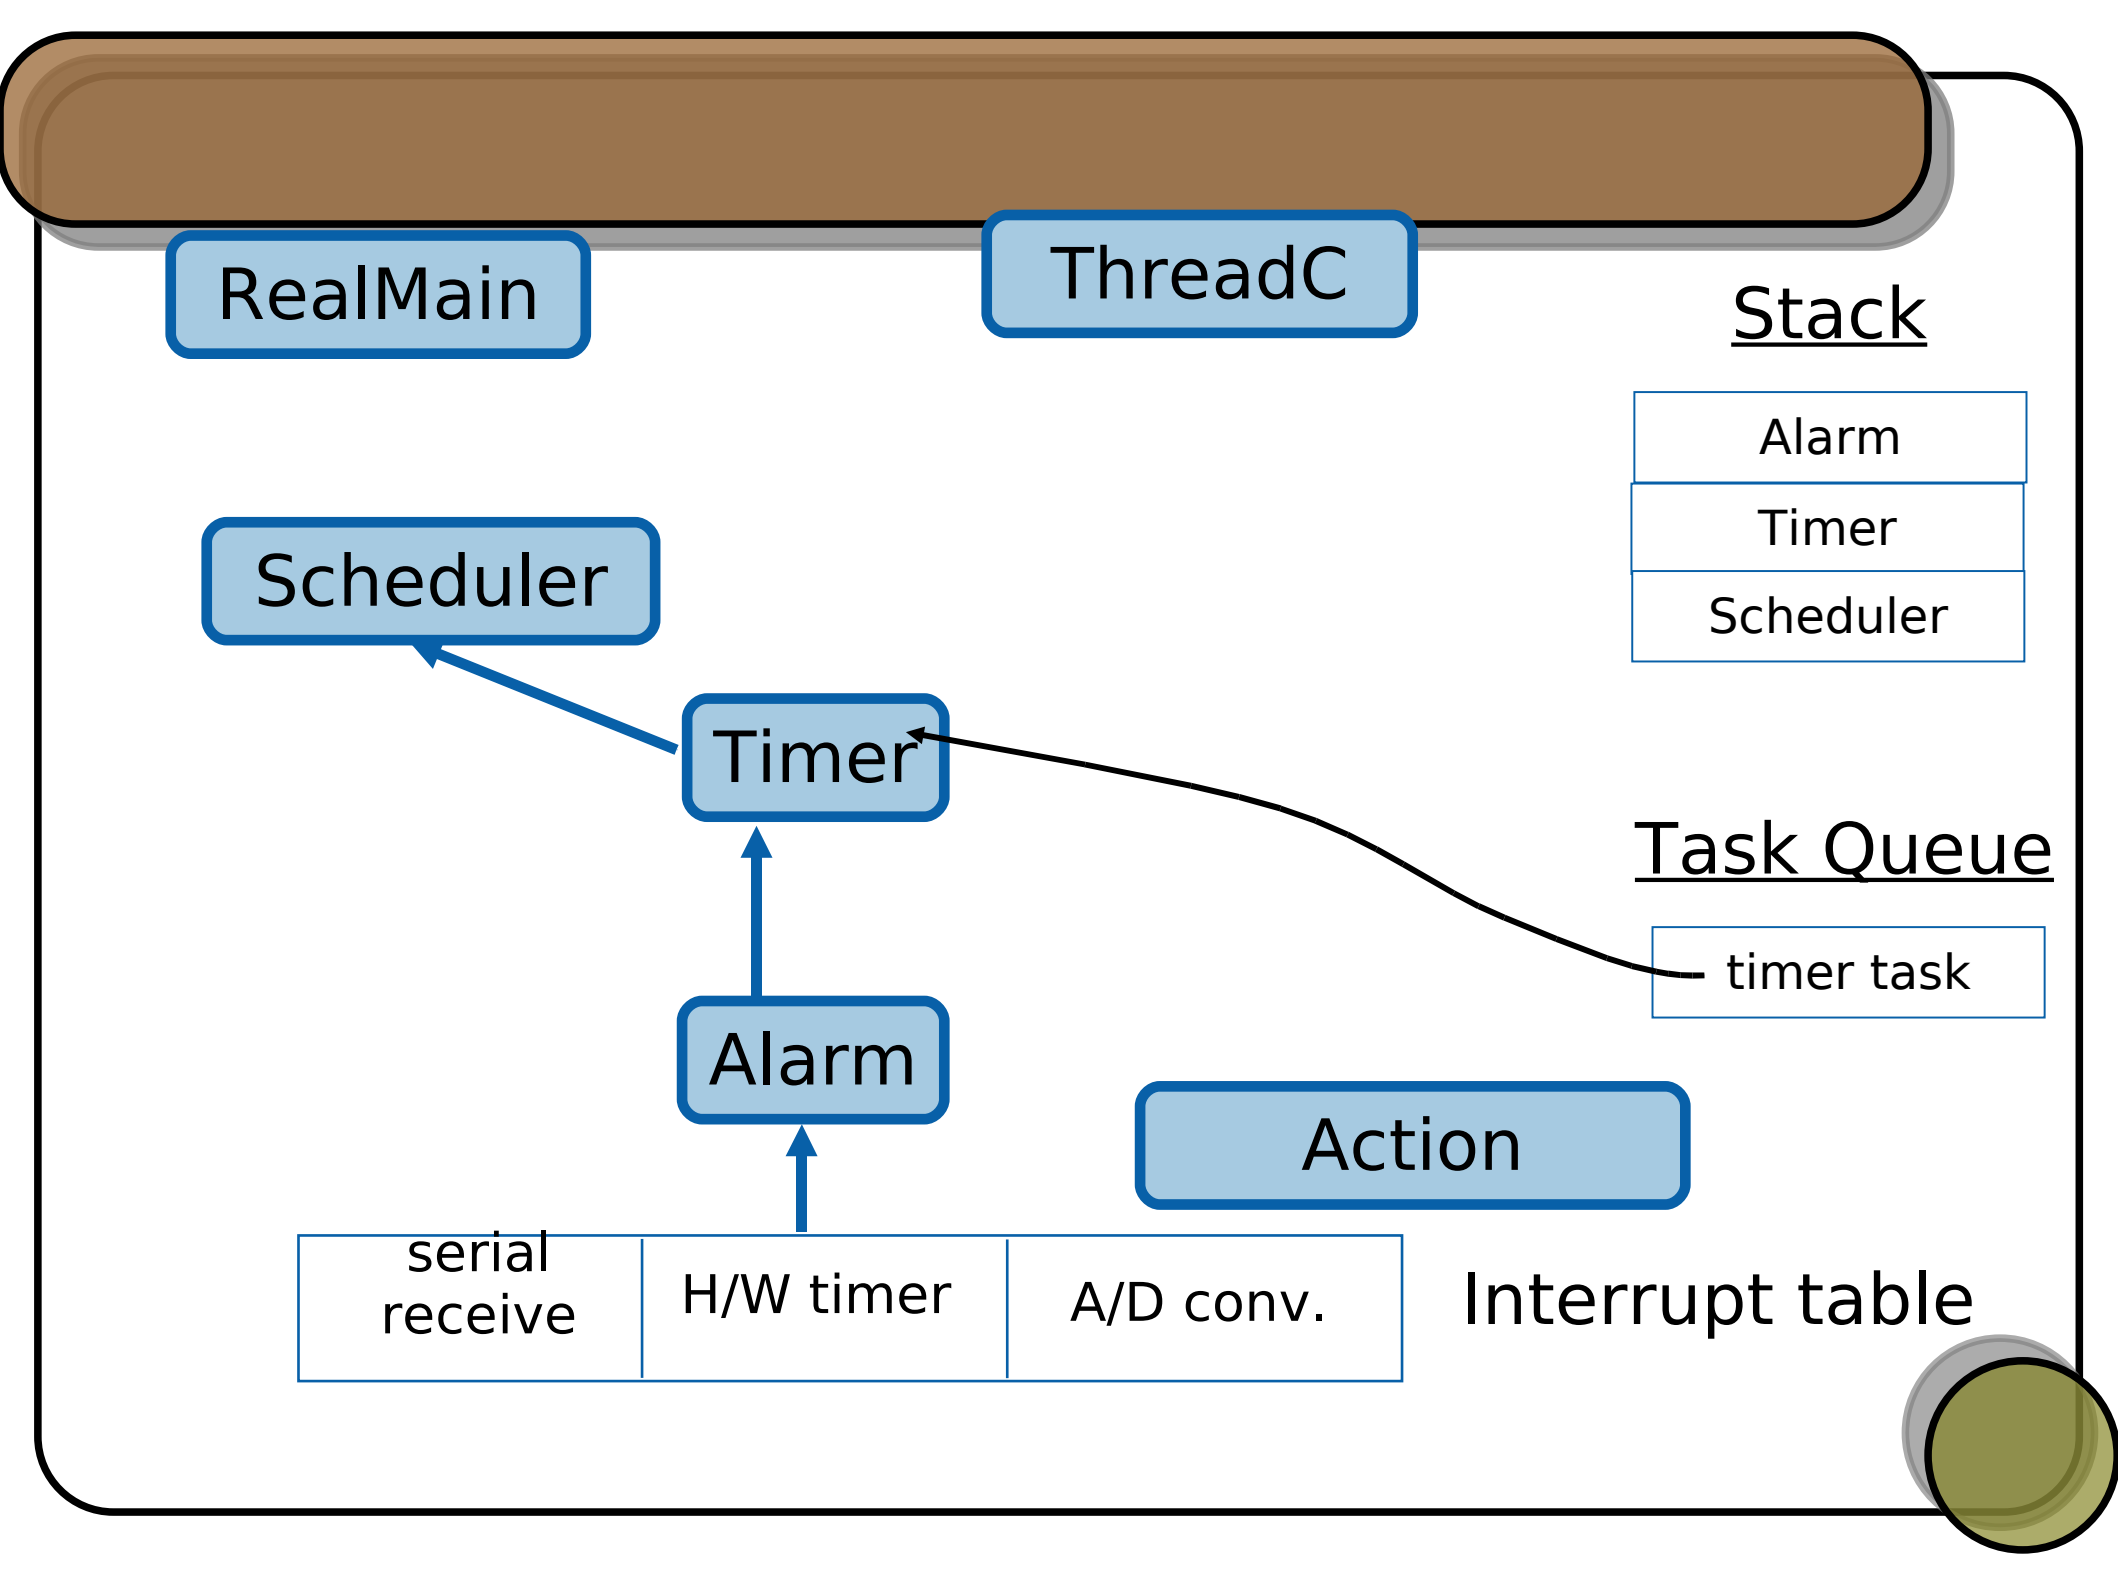


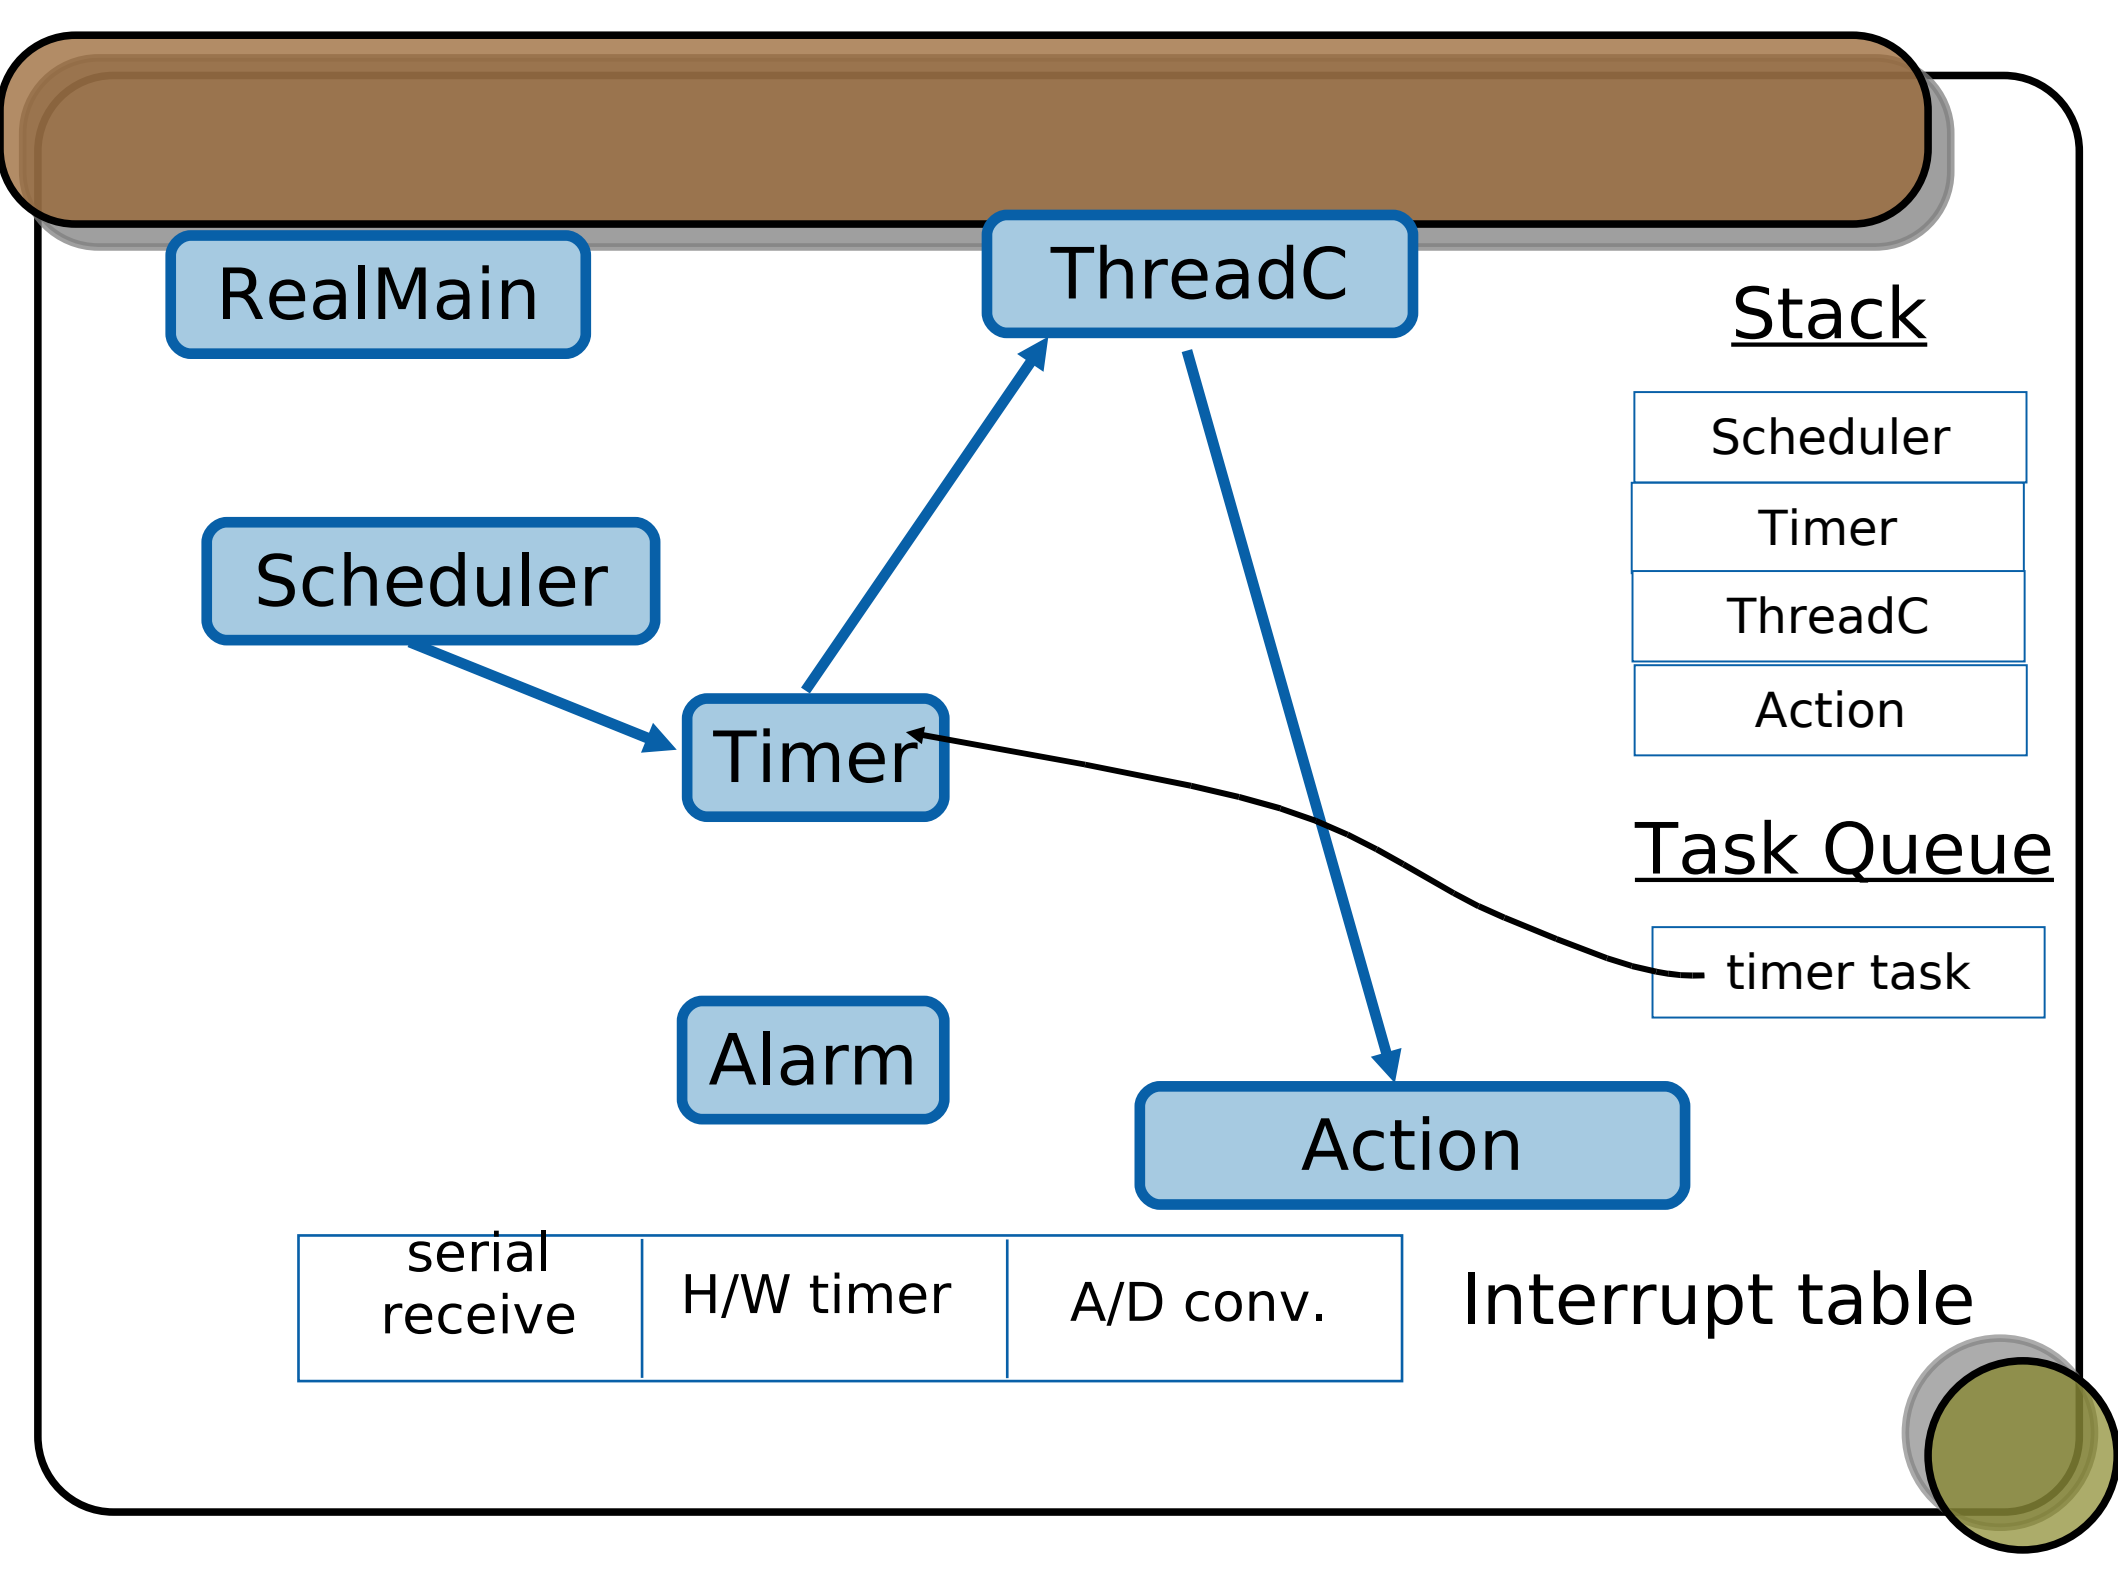
RTOS Kernel

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

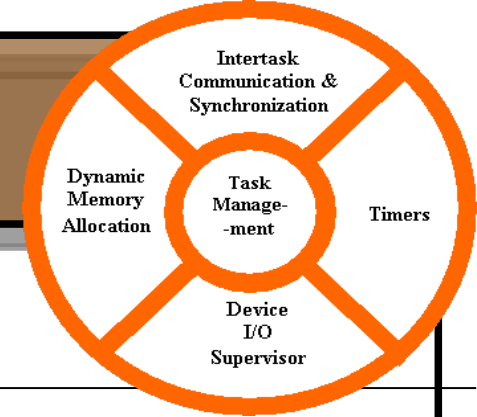








RTOS Kernel::Tasks



Task 是 RTOS 最重要的項目

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

RTOS scheduler 必須 deterministic

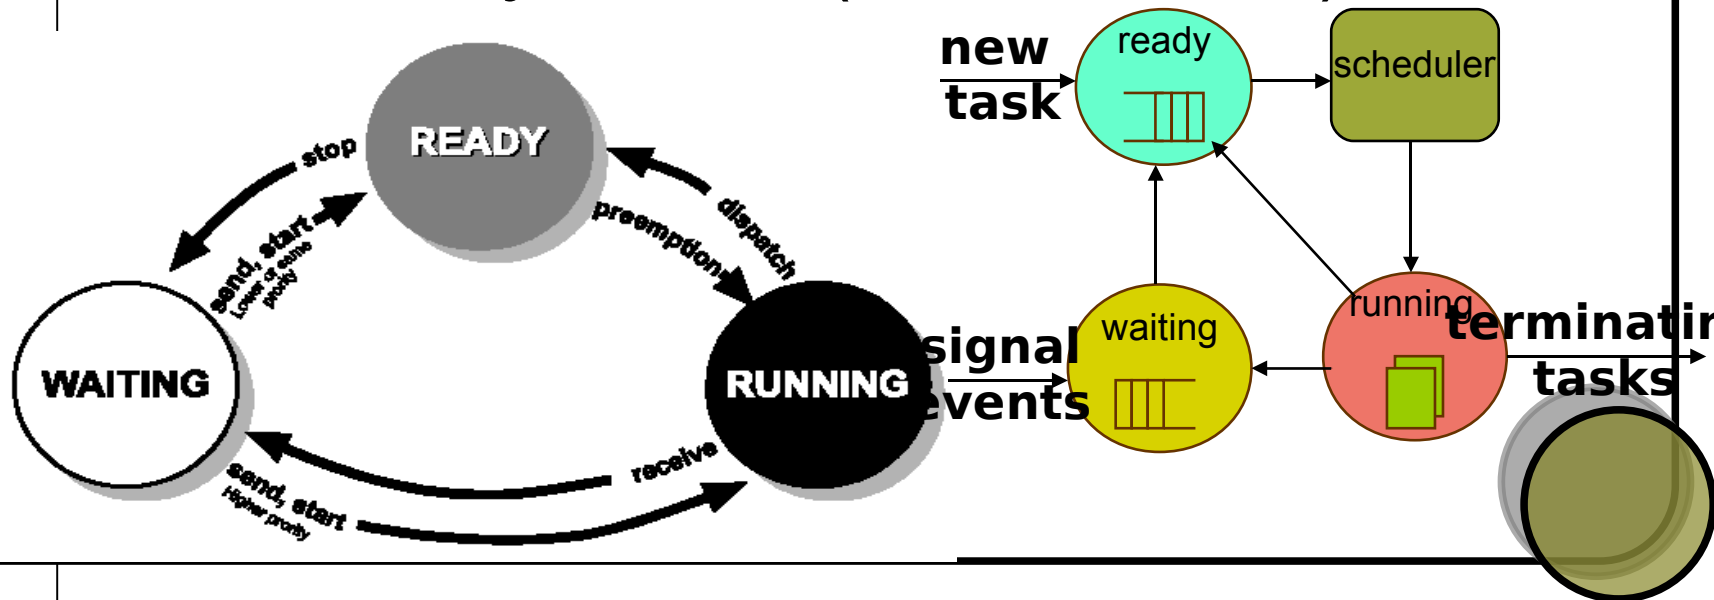
$O(1)$ or $O(n)$

Scheduling policy

Clock driven

Priority driven (RMS & EDF)

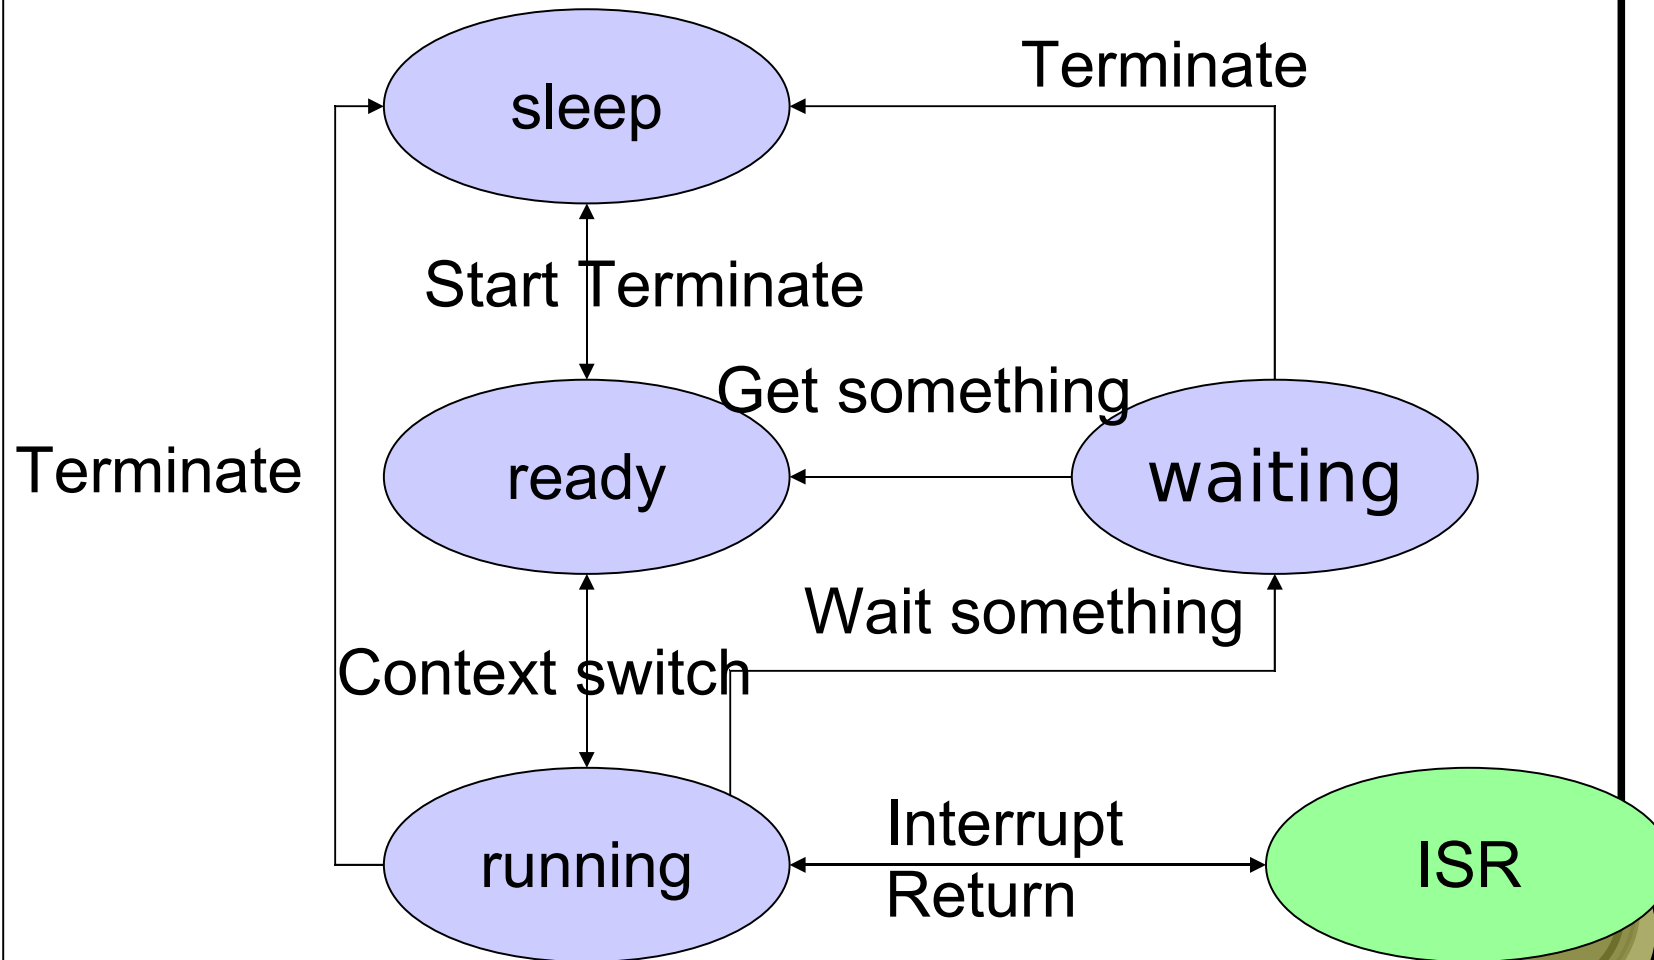
笨蛋，問題在
scheduling !



RTOS Kernel::Tasks

Task 的狀態移轉

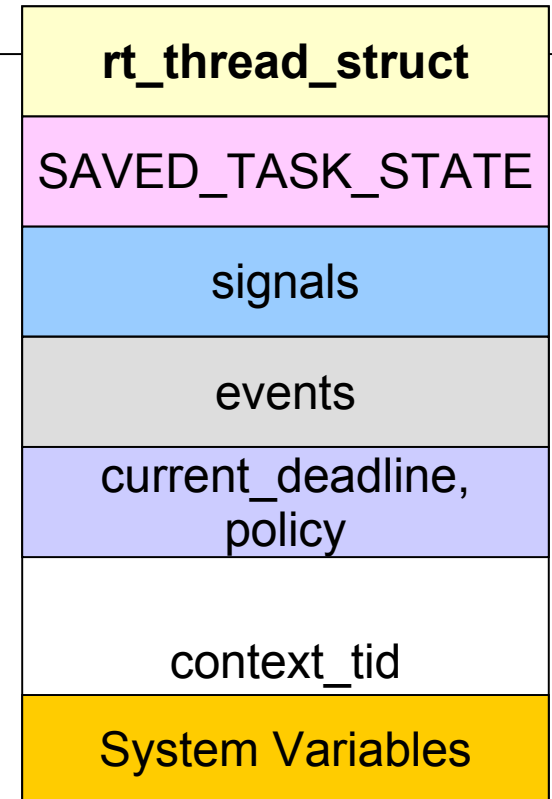
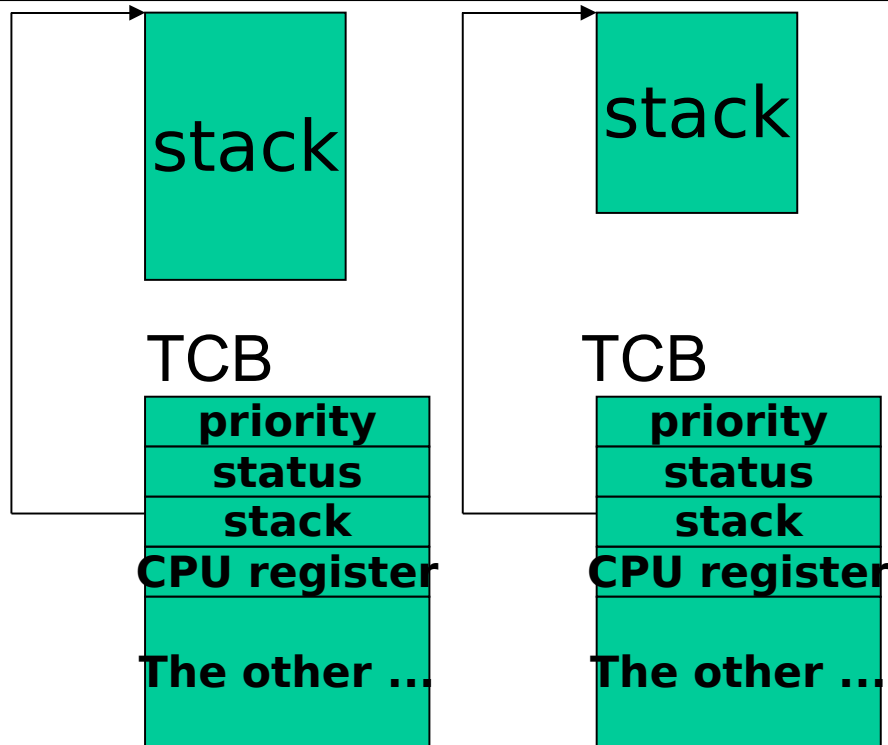
- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action



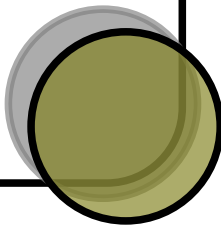
RTOS Kernel::Tasks

TCB (Task Control Block)

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action



CPU registers → context



RTOS Kernel::Tasks

Jamei 中 thread 是 Task 的單元

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

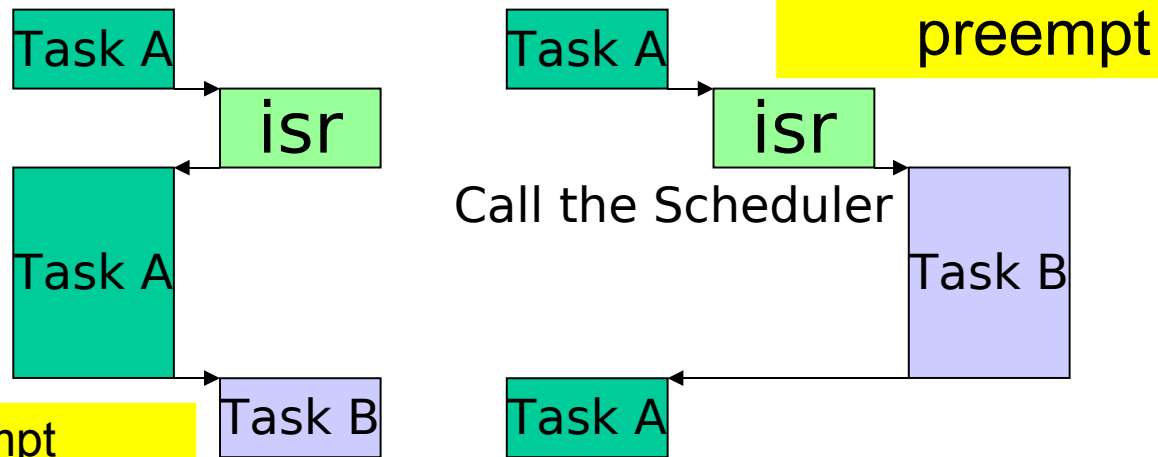
```
typedef struct rt_thread_struct *pthread_t;  
pthread_create  
pthread_exit  
pthread_kill  
pthread_wakeup_np  
pthread_suspend_np  
pthread_wait_np  
...
```

- 試圖滿足 POSIX Thread 語意
- _np]non-POSIX

RTOS Kernel::Tasks

Jamei 中 thread 是 Task 的單元

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action



non-preempt

preemptive scheduling

```
int rt_schedule(void)
{
    struct rt_thread_struct *preemptor = 0;
    ...

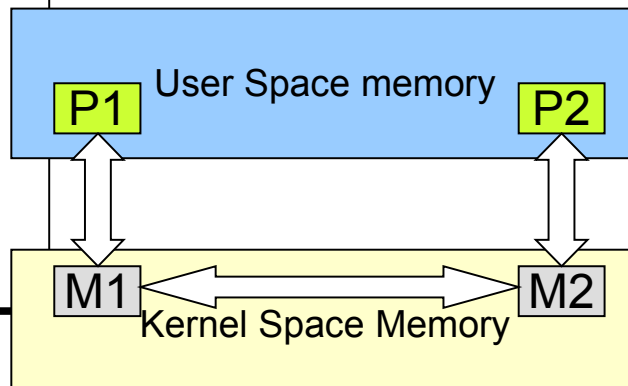
    if ((s->clock->mode == RT_CLOCK_MODE_ONESHOT)) {
        if ((preemptor = find_preemptor(s, new_task)))
        {
            (s->clock)->settimer(s->clock, preemptor->resume_time - now);
        } else {
            (s->clock)->settimer(s->clock, (HRTICKS_PER_SEC / HZ) / 2 );
        }
        set_bit (RT_SCHED_TIMER_OK, &s->sched_flags);
    }
}
```

優先權 Task A < Task B

RTOS Kernel::Memory

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

在 i386 硬體架構下，Jamei 可支援
MMU (virtual memory) 與 MPU
(memory protection)
區分 user-space 與 kernel-space
memory



```
void start_kernel(void)
{
    /* architecture-dependent */
    init_arch();
    /* NOTE:
     * - Mask all interrupts.
     * - Interrupts are mapped in 0x20-0x30 IDT entries
     */
}
```

```
#if CONFIG_KERNEL_MEMORYPROT
    init_page();
#endif
```

```
#if CONFIG_CONTEXT_MEMORYPROT
    init_context();
#endif
```

...

RTOS Kernel::Memory

Jamei 仿效 POSIX/Linux 的處理

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

APIs

kmalloc / kfree

mmap

shm (shared memory area)

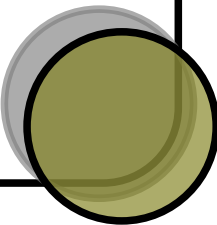
有彈性的記憶體管理機制

放任 (no protected memory)

最低限度記憶體保護機制

不可寫入 RT executive 記憶體

保護 context 執行單元記憶體



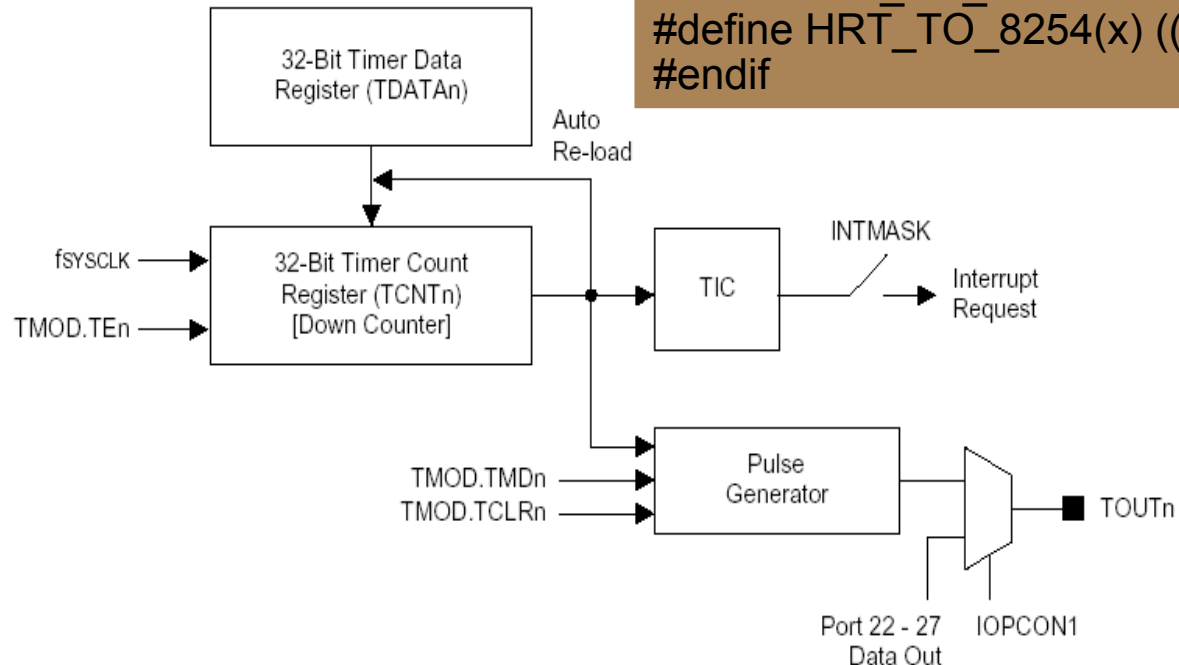
RTOS Kernel::Timer

Timer 本質上是硬體時鐘的軟體呈現

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

型態

watchdog timer
programmable
timer



```
#define HRT_FROM_NS(x) (x)
#define HRTICKS_PER_SEC 1000000000

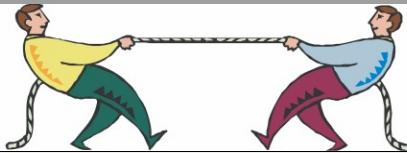
#ifndef HRT_FROM_8254
#ifndef HRT_FROM_8254
#define HRT_FROM_8254(x) ((x) * 838)
#endif
#endif

#ifndef HRT_TO_8254
#define HRT_TO_8254(x) ((x) / 838)
#endif
```

RTOS Kernel::Timer

萬惡 HZ

Throughput



High responsiveness

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

實做 HRT (High Resolution Timer)
用於 sched, sync 等
系統實作

```
struct rt_clock {  
    ...  
    int (*sethrttime)(  
        struct rt_clock *,  
        hrtime_t t);  
  
    int (*settimer)(  
        struct rt_clock *,  
        hrtime_t interval);  
  
    ...  
    hrtime_t resolution;  
    hrtime_t value;  
    hrtime_t delta;  
    pthread_spinlock_t lock;  
    struct rt_clock_arch arch;  
};
```

```
int rt_schedule(void)  
{  
    ...  
    if ((s->clock->mode == RT_CLOCK_MODE_ONESHOT)) {  
        if ((preemptor = find_preemptor(s,new_task)))  
  
        {  
            (s->clock)->settimer(s->clock, preemptor->resume_time - now);  
        } else {  
            (s->clock)->settimer(s->clock, (HRTICKS_PER_SEC / HZ) / 2 );  
        }  
    }  
    ...  
}
```


RTOS Kernel::I/O

everything is file

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

Interrupt-driven
polling / DMA
I/O mapping

memory space 與 I/O 操作的對應

APIs

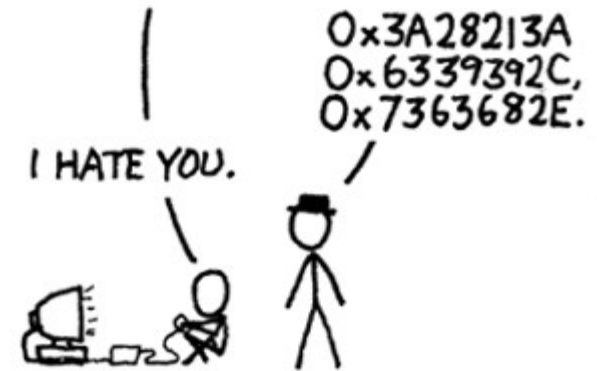
open / close

read / write

mmap

register_rtdev / unregister_rtdev

MAN, I SUCK AT THIS GAME.
CAN YOU GIVE ME
A FEW POINTERS?

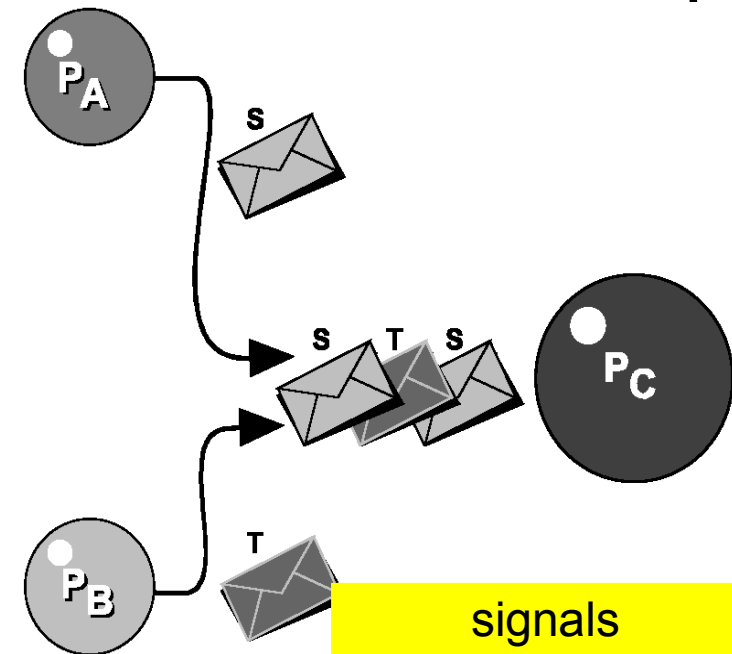
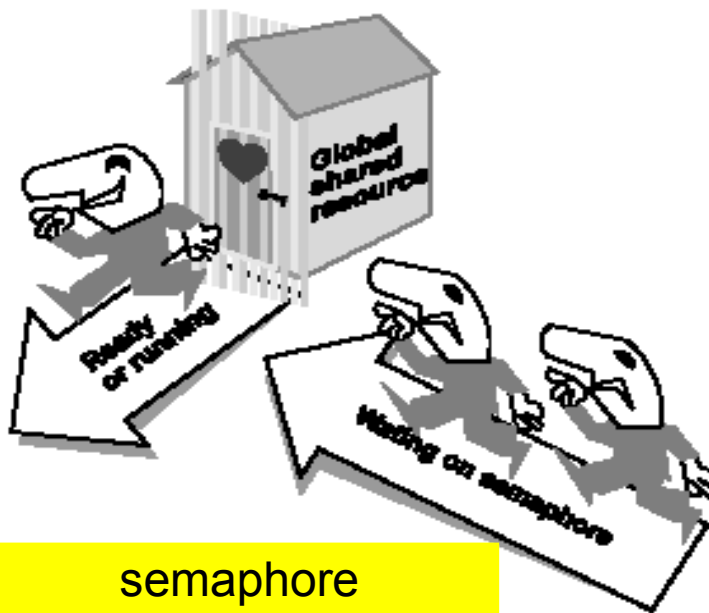


RTOS Kernel::IPC

其實是 Inter-Task Communications

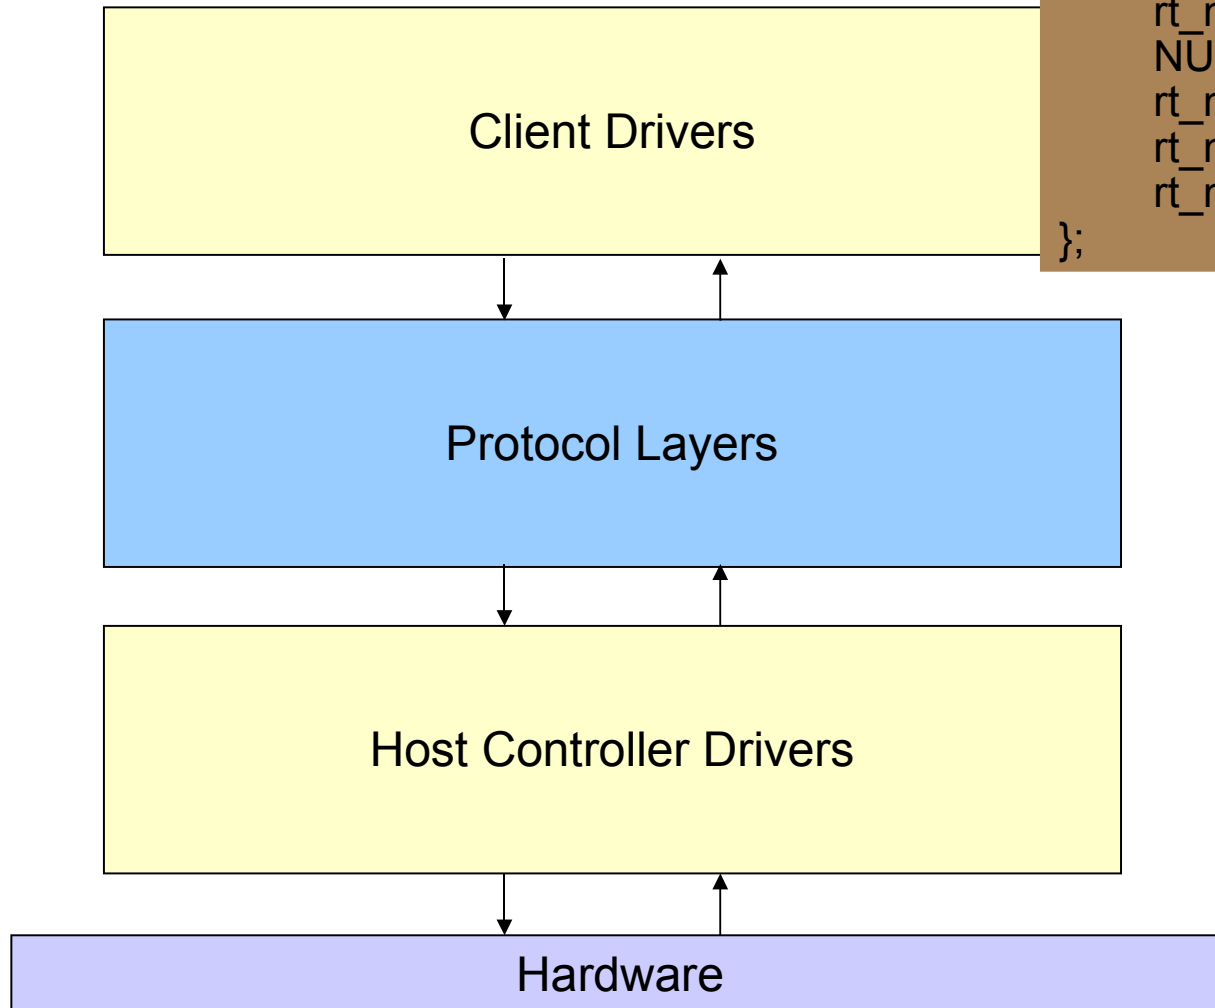
- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

可配置的組態
Signals
Semaphore

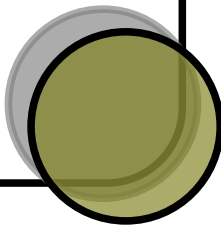


RTOS Kernel::Device Driver

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

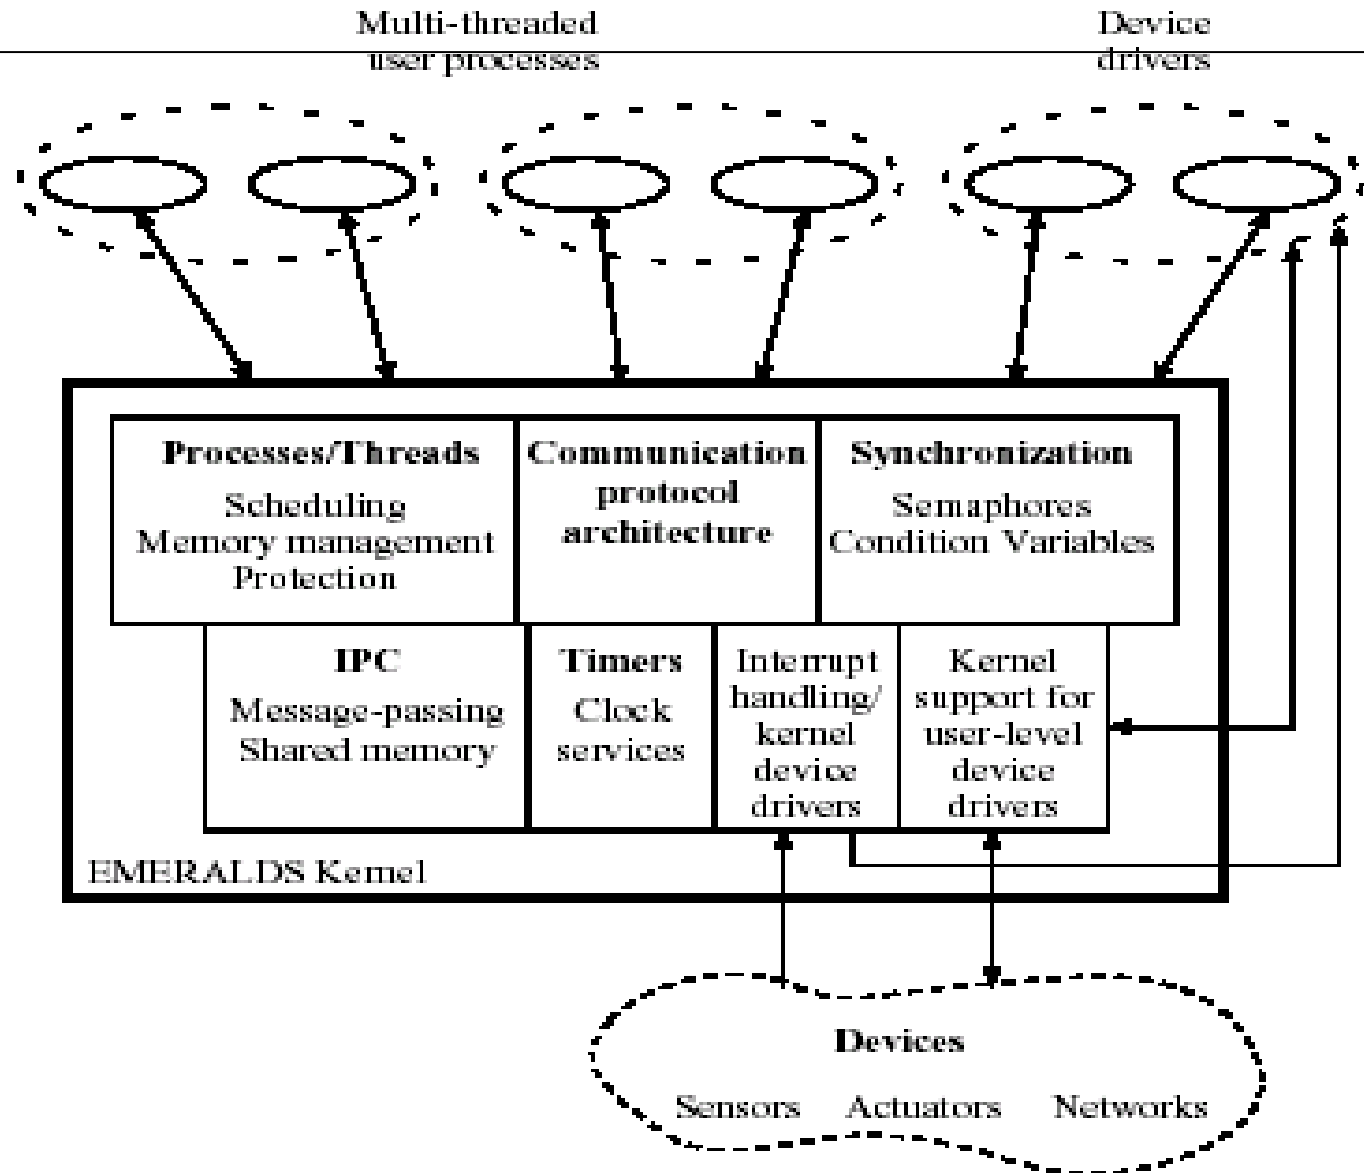


```
static struct rt_file_operations
rt_mem_fops = {
    rt_mem_llseek,
    rt_mem_read,
    rt_mem_write,
    NULL,
    rt_mem_mmap,
    rt_mem_open,
    rt_mem_release
};
```



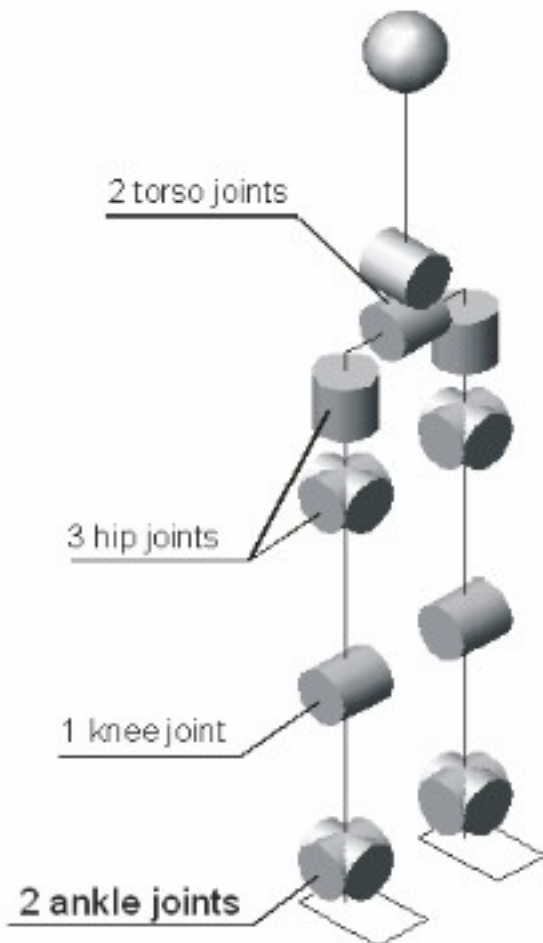
In an Action

- Introduction
- Structure
- RTOS Kernel
- Tasks
- Memory
- Timers
- I/O
- IPCs
- Device Driver
- In an Action

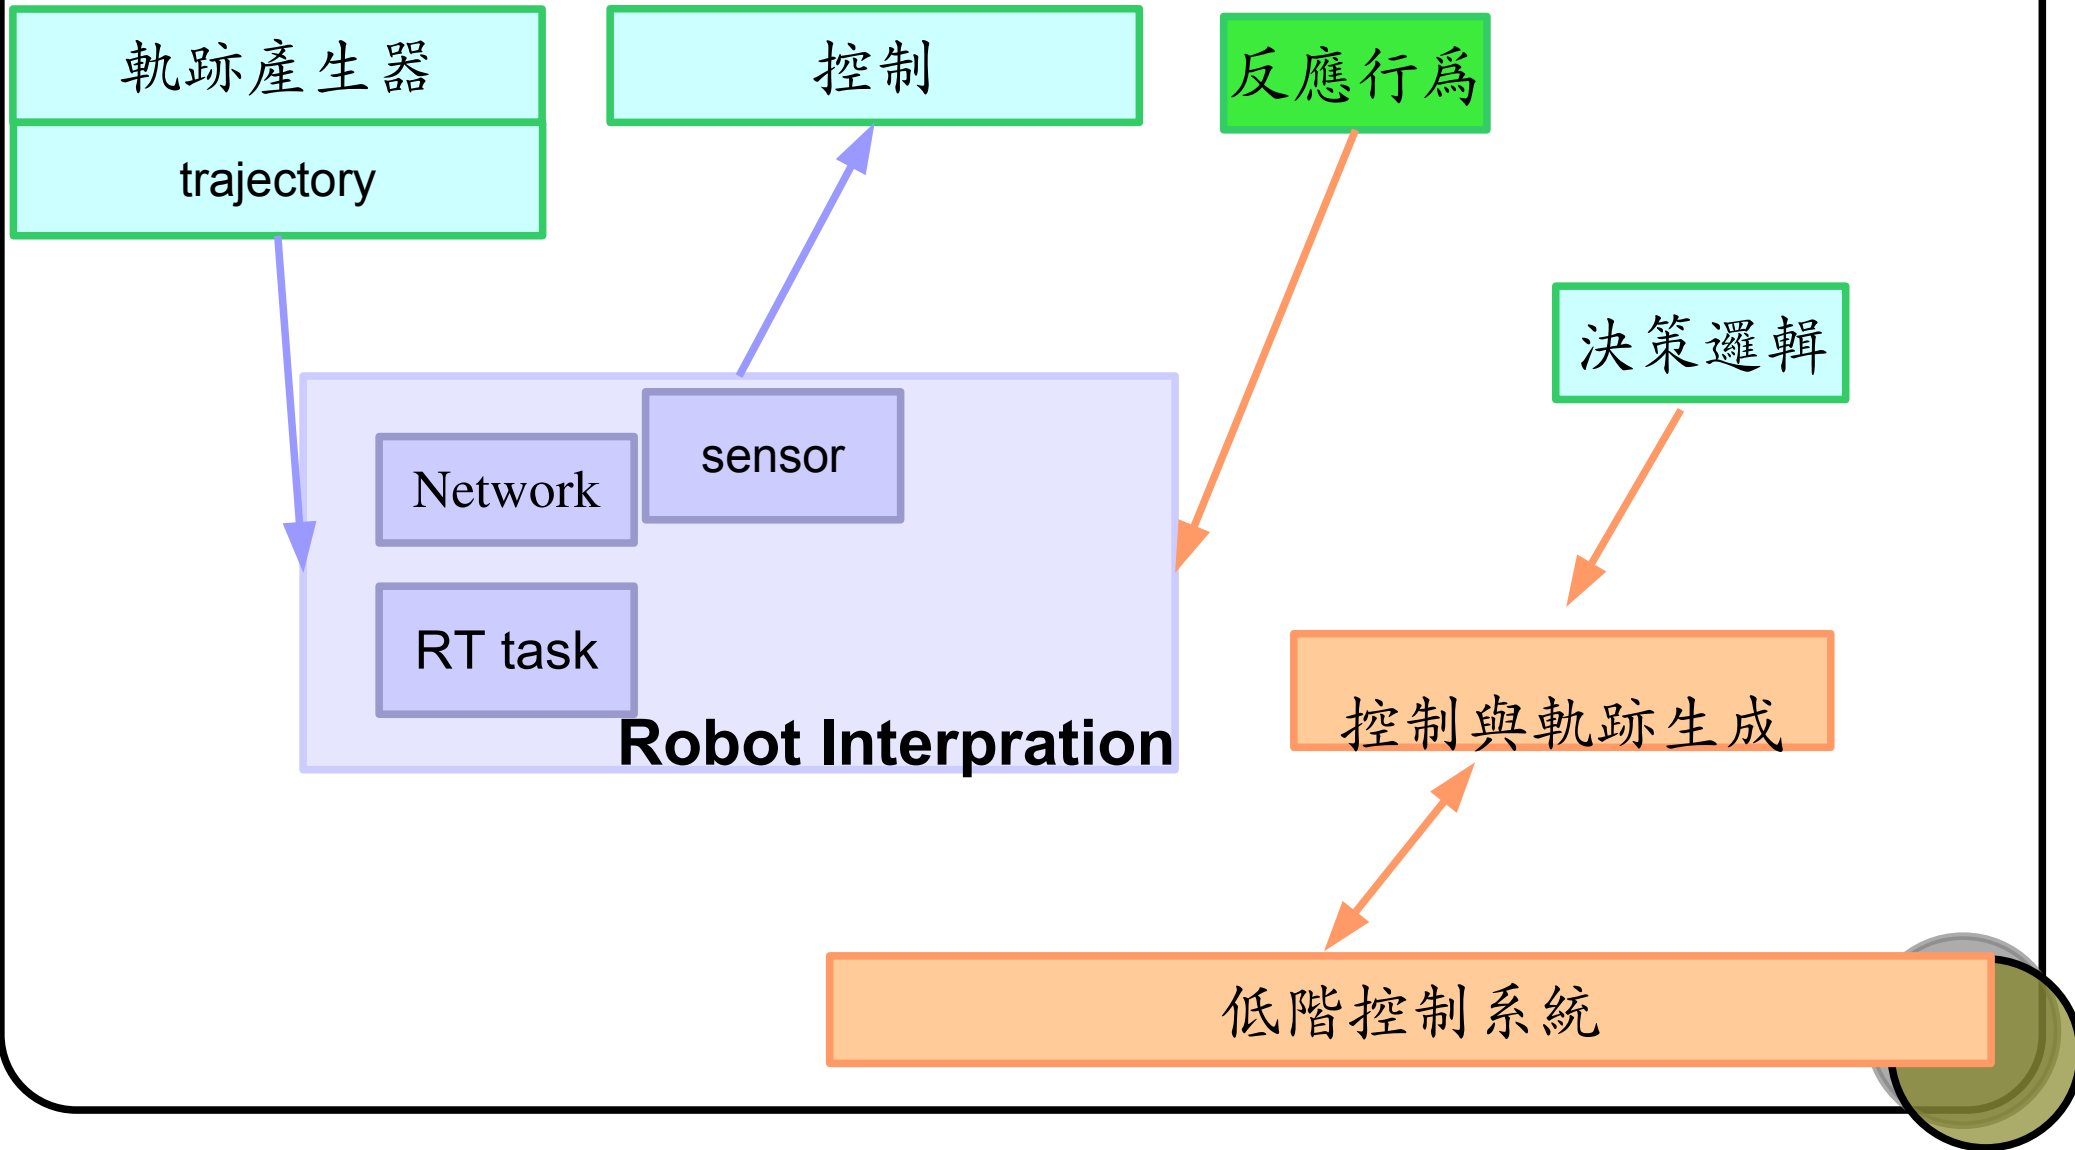


機器人硬體機構概況

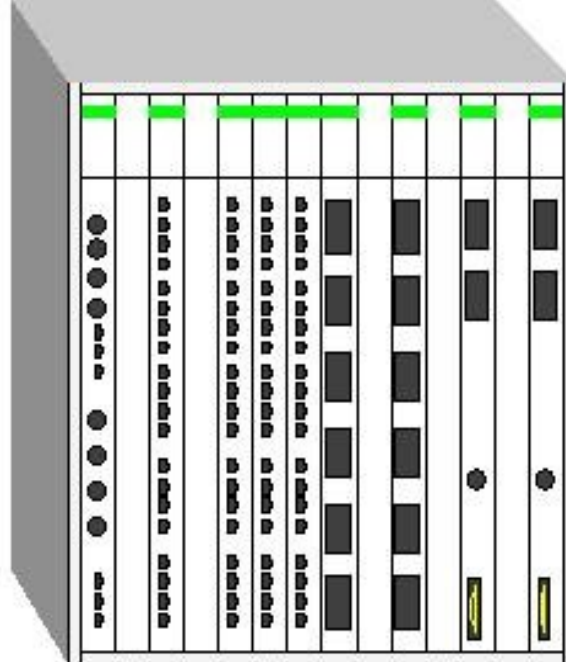
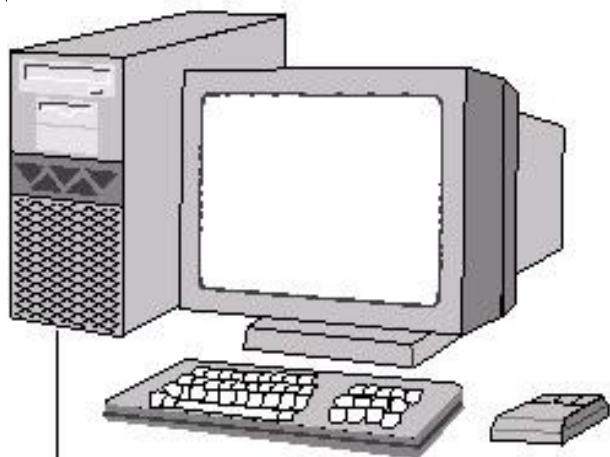
硬體：就是電腦系統中，可以讓你踢一腳的地方



控制系統



GDB or DDD

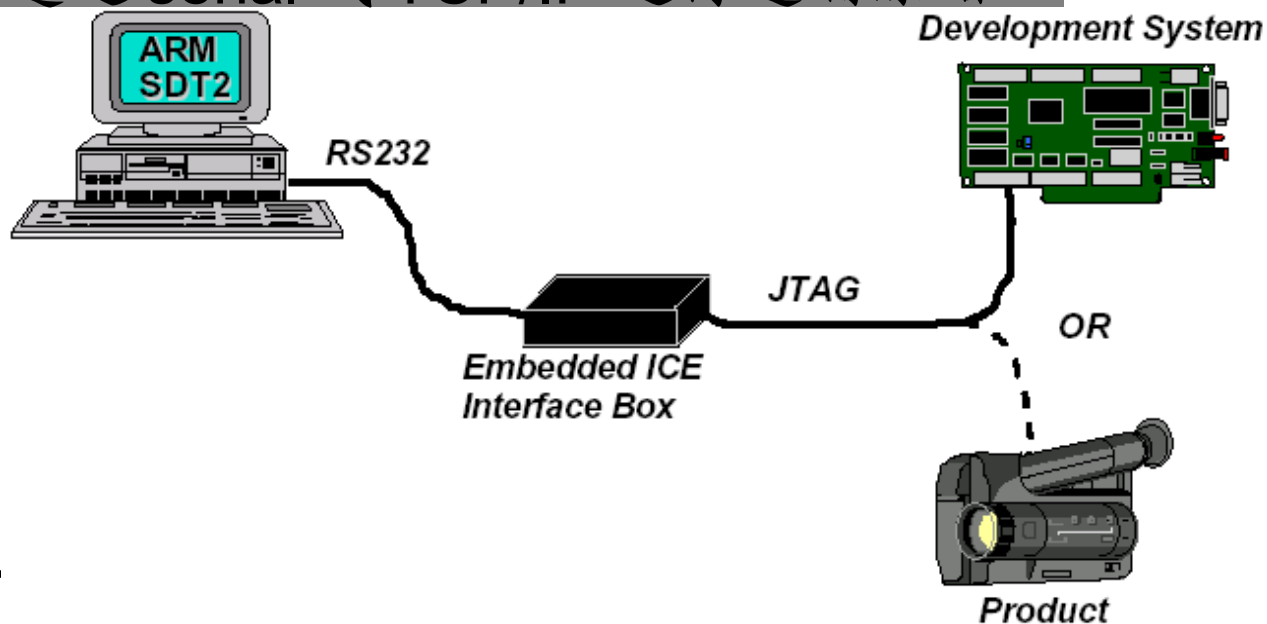


(remote protocol : some message string)

Serial **gdb stub**

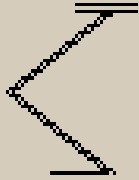
考慮在 emulation/target 模式下，該如何喚起 gdb ？

Remote Debugging：透過 serial 或 TCP/IP 進行遠端除錯



Show me the Robot

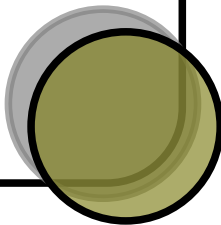
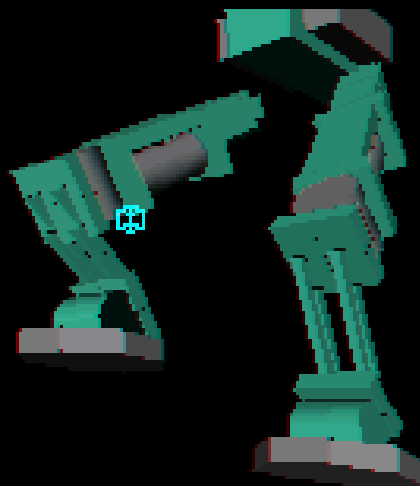
Left Leg

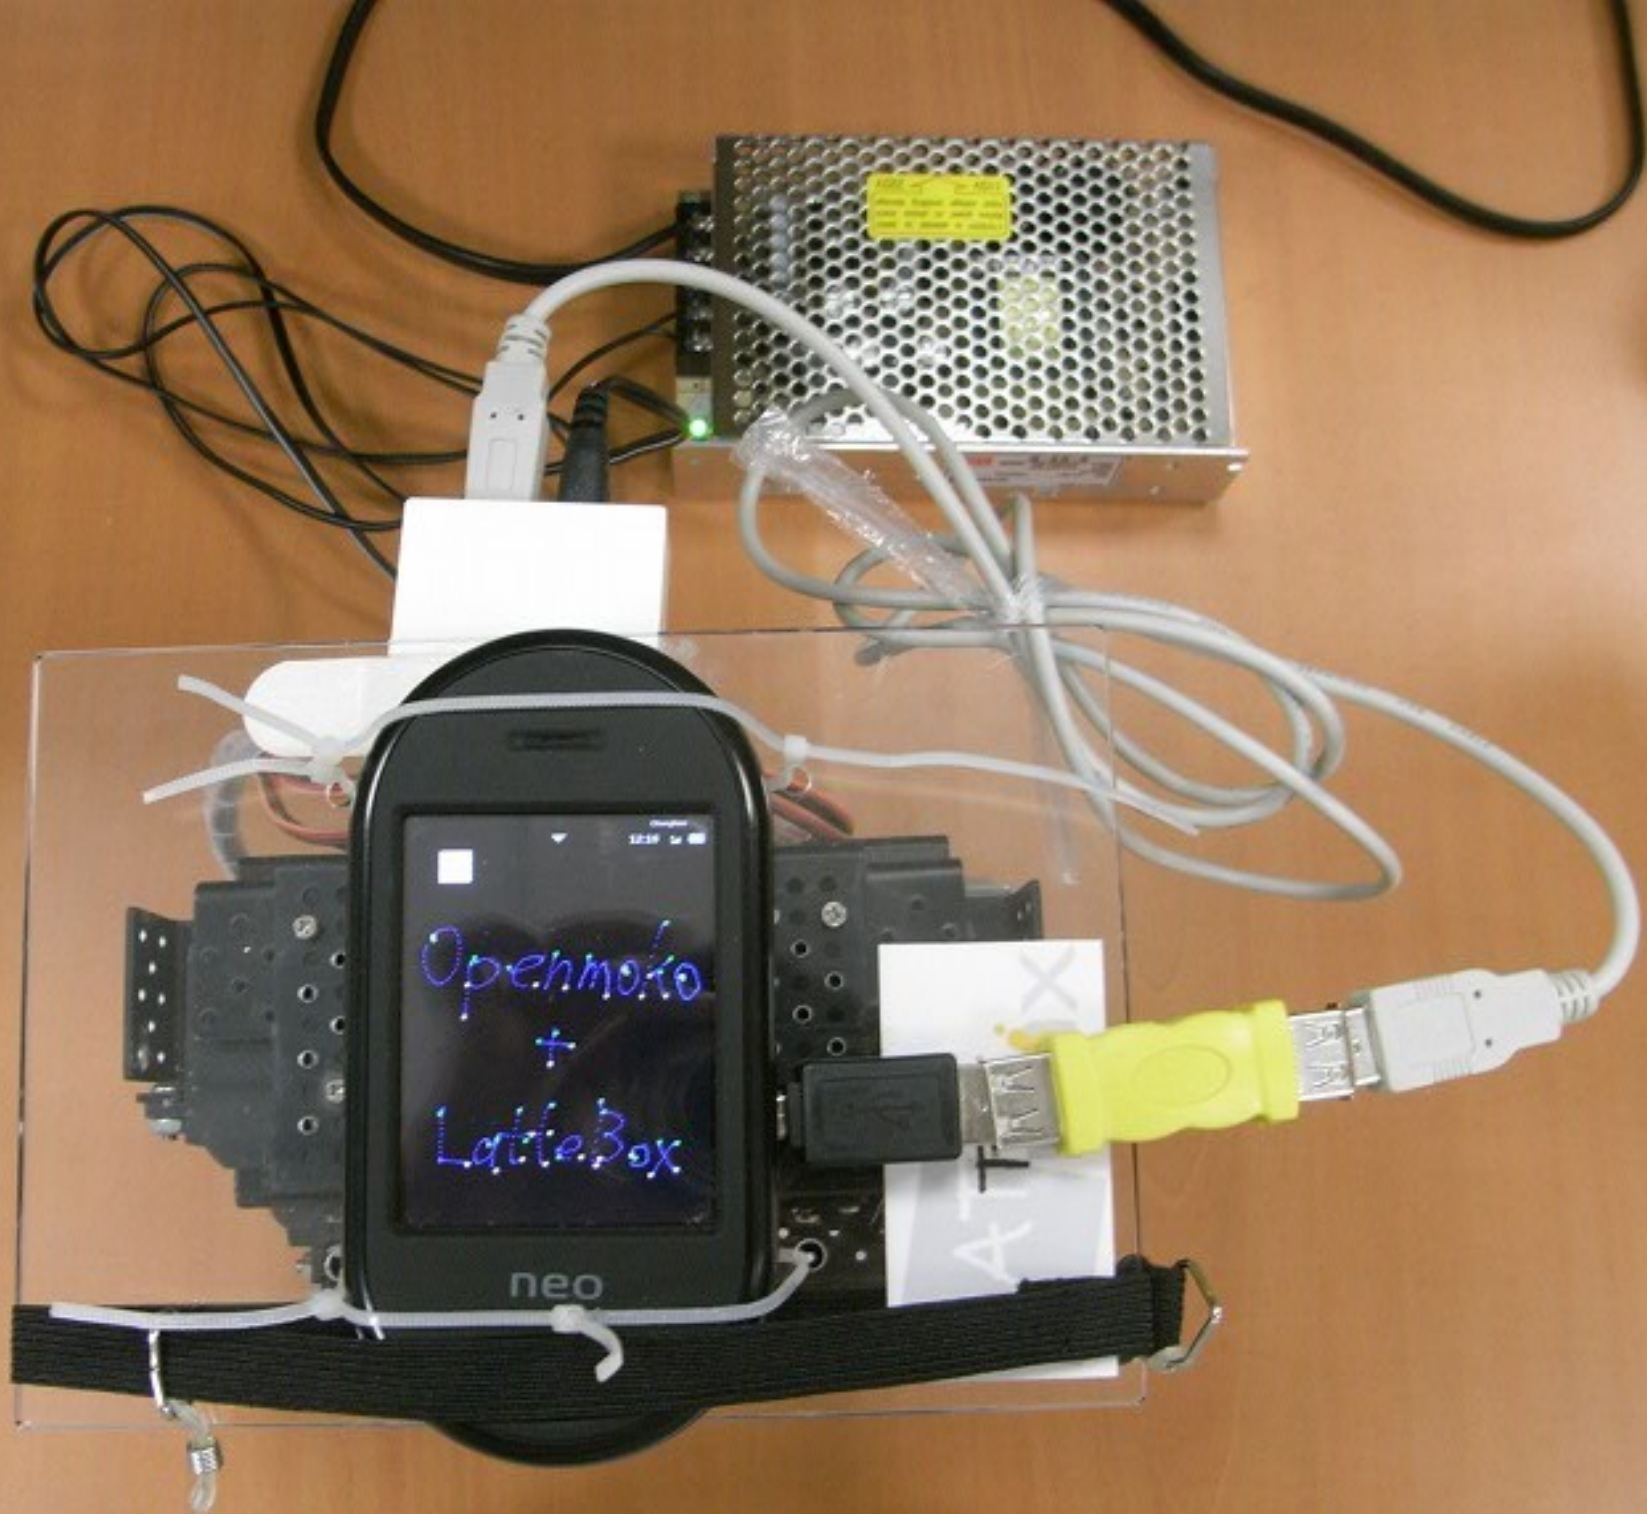


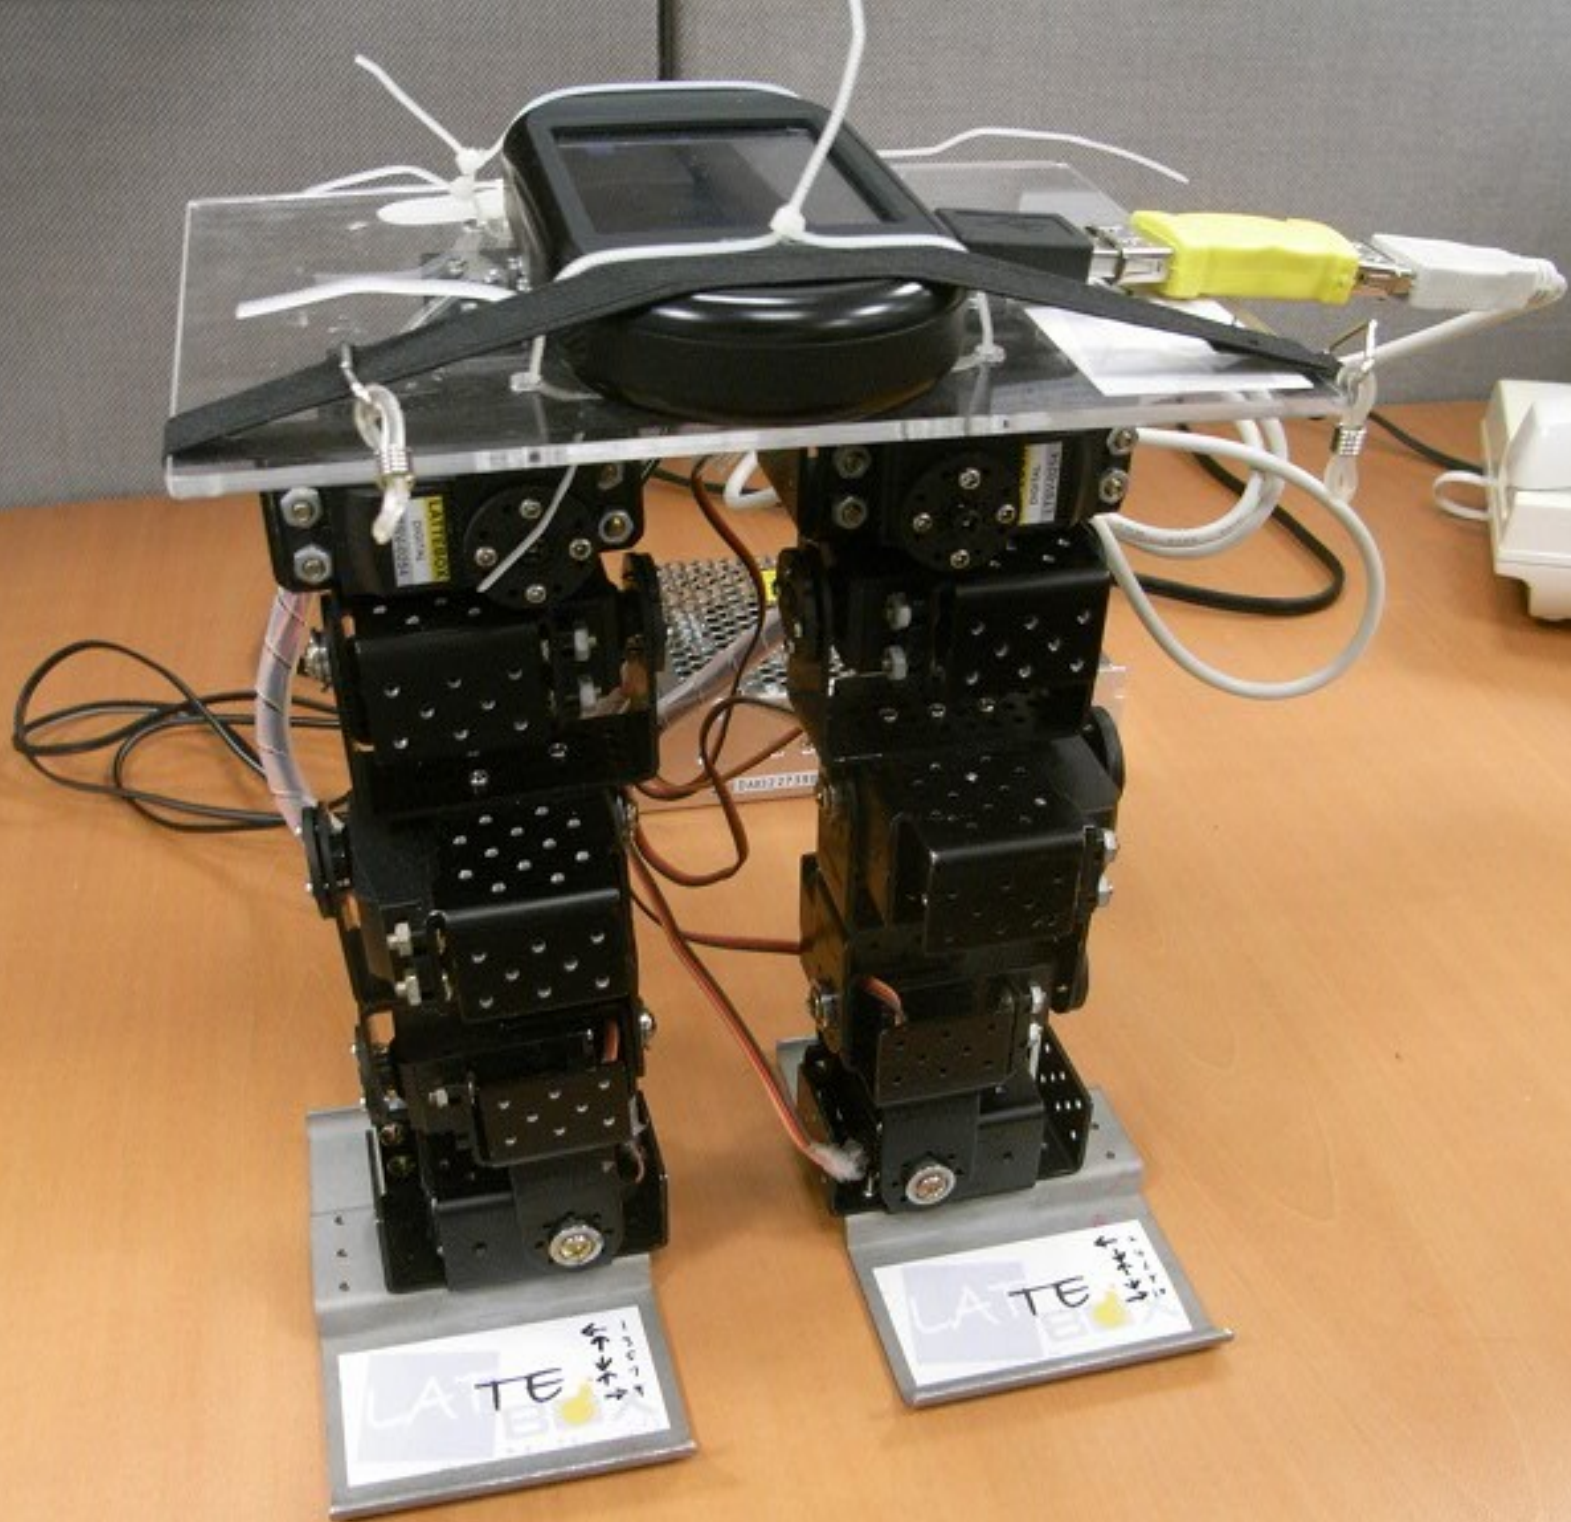
Right Left



3D View







實現 TCP/IP 的 難題

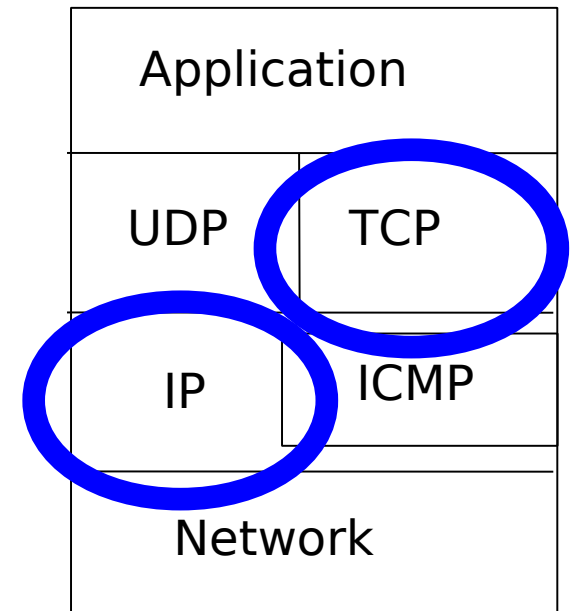
IP Network

- UDP – best-effort datagrams
- TCP – connection oriented, reliable byte-stream, full-duplex
 - Flow control, congestion control, etc
- IP – best-effort packet delivery
 - Forwarding, fragmentation

並未對資源消耗作優化

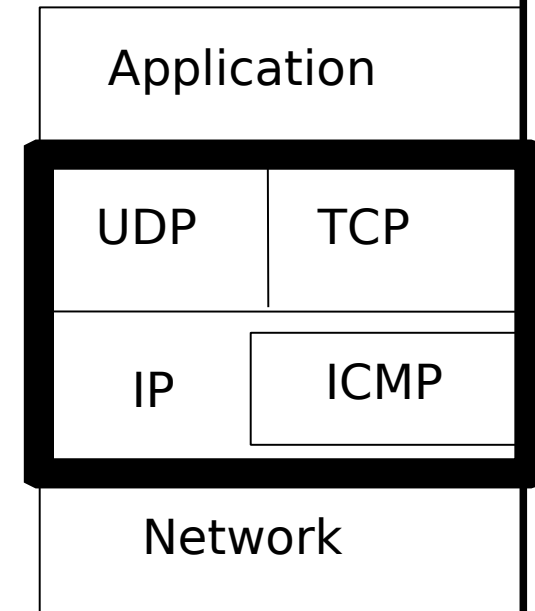
相對來說，是可行的解決方案

共通性



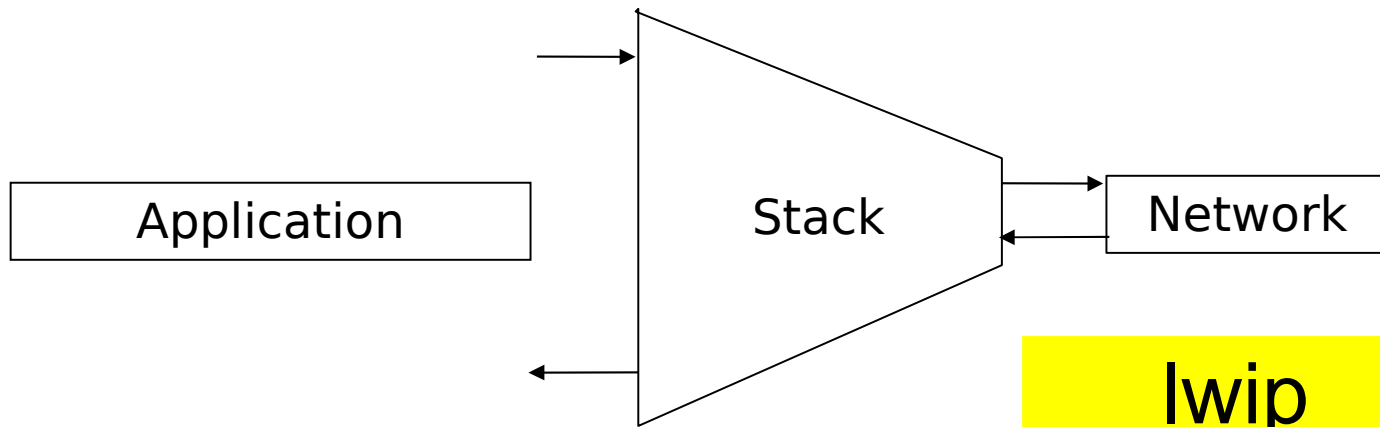
TCP/IP stack design

- lwIP – lightweight IP
 - “Application driven” : larger
- μ IP – micro IP
 - “Network driven” : smaller



TCP/IP stack design

Application driven



Application

UDP

TCP

IP

ICMP

Network

lwip

Network driven



Network

Stack

Application

uIP

資源有限、慾望無窮

成本、實體空間

Memory

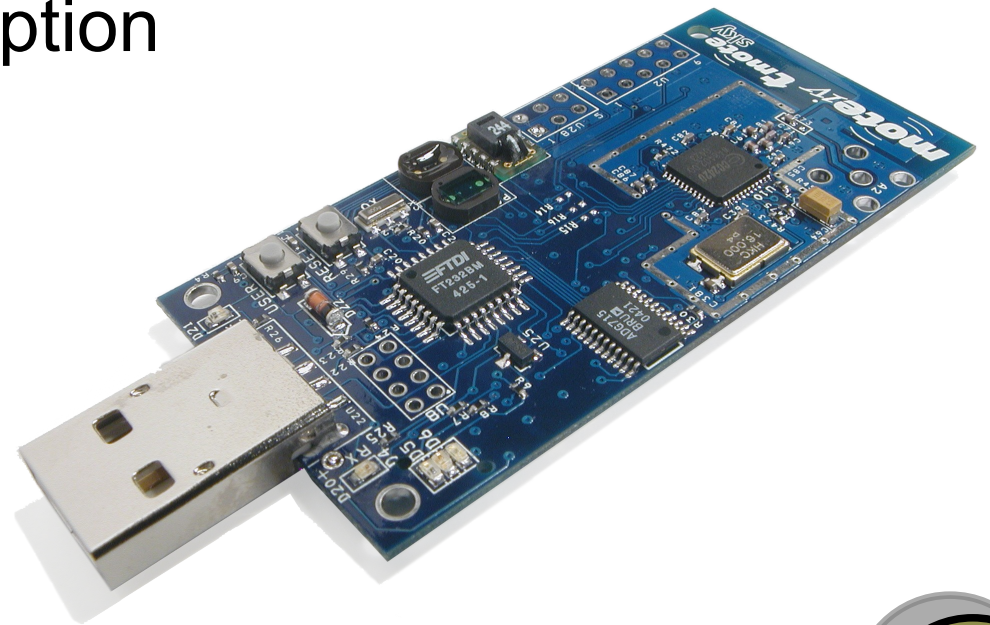
~10 k RAM, ~100 k ROM

Energy, power consumption

Batteries, ~10 mW

Bandwidth

~100 kbits/second

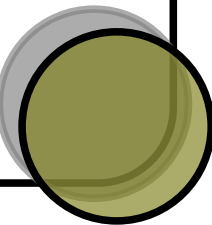


資源限制

memory

energy

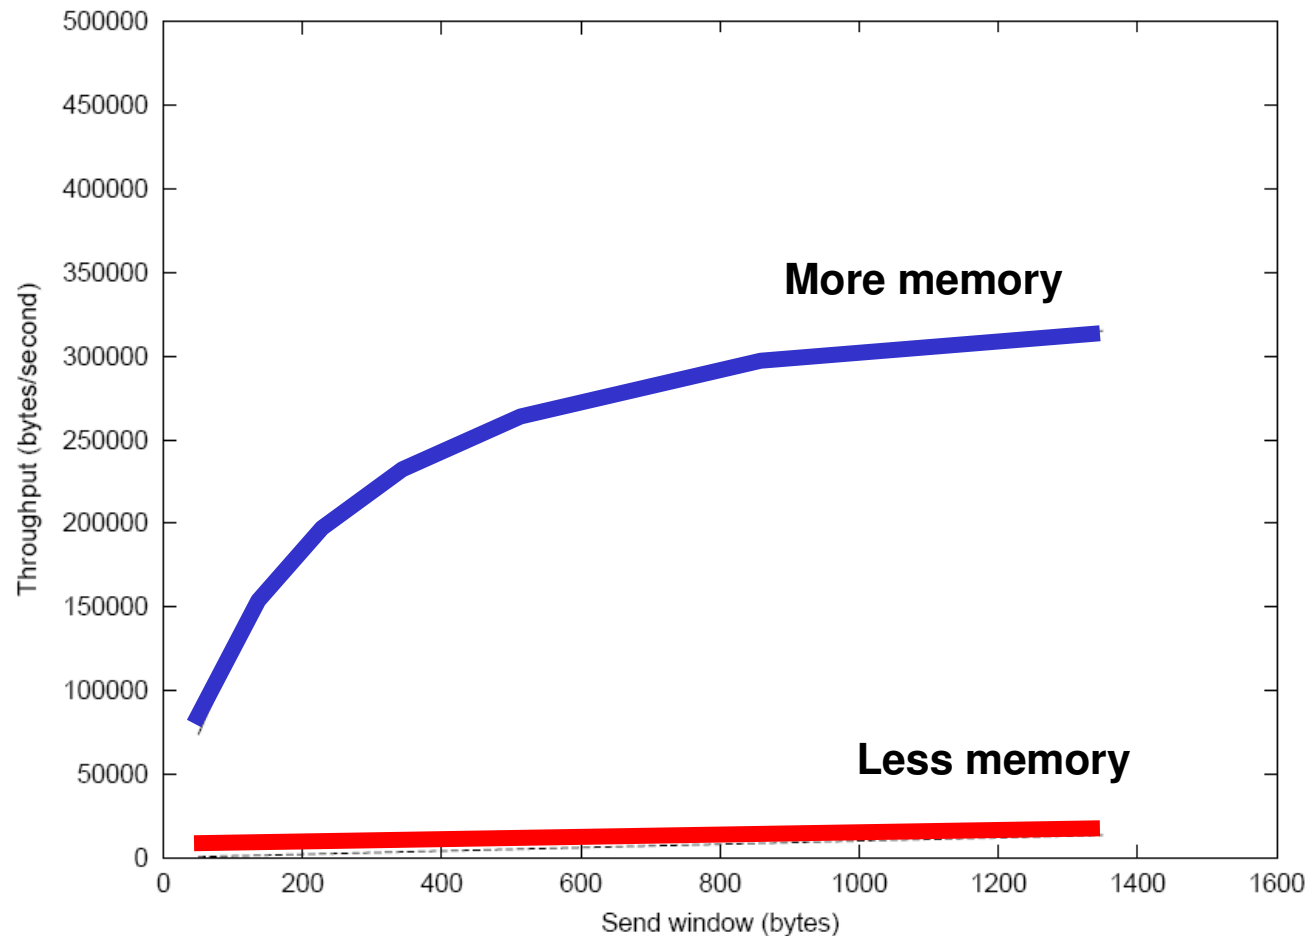
bandwidth



資源限制：Memory+TCP/IP

RAM vs throughput

基本上是兩難

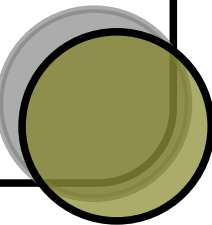


資源限制

memory

energy

bandwidth



資源限制：功耗

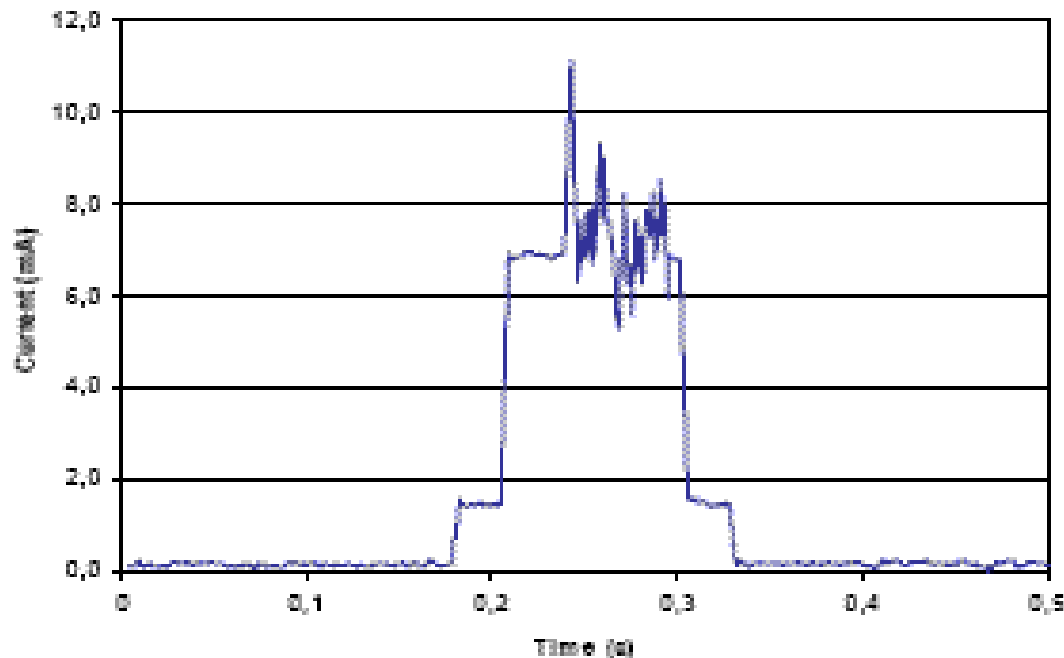


Figure 6. Power consumption: sending

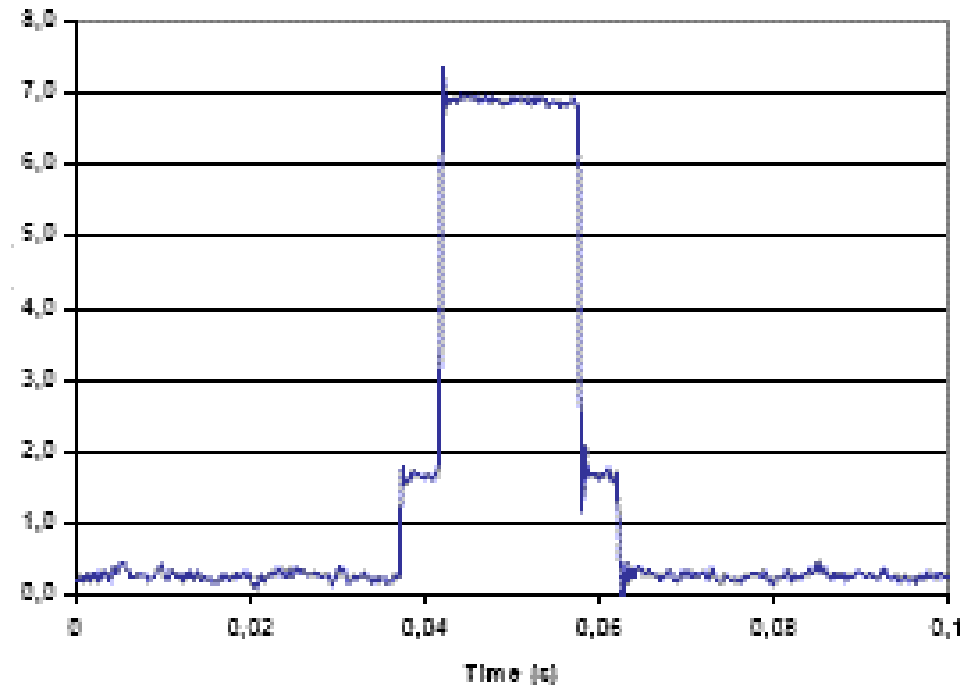
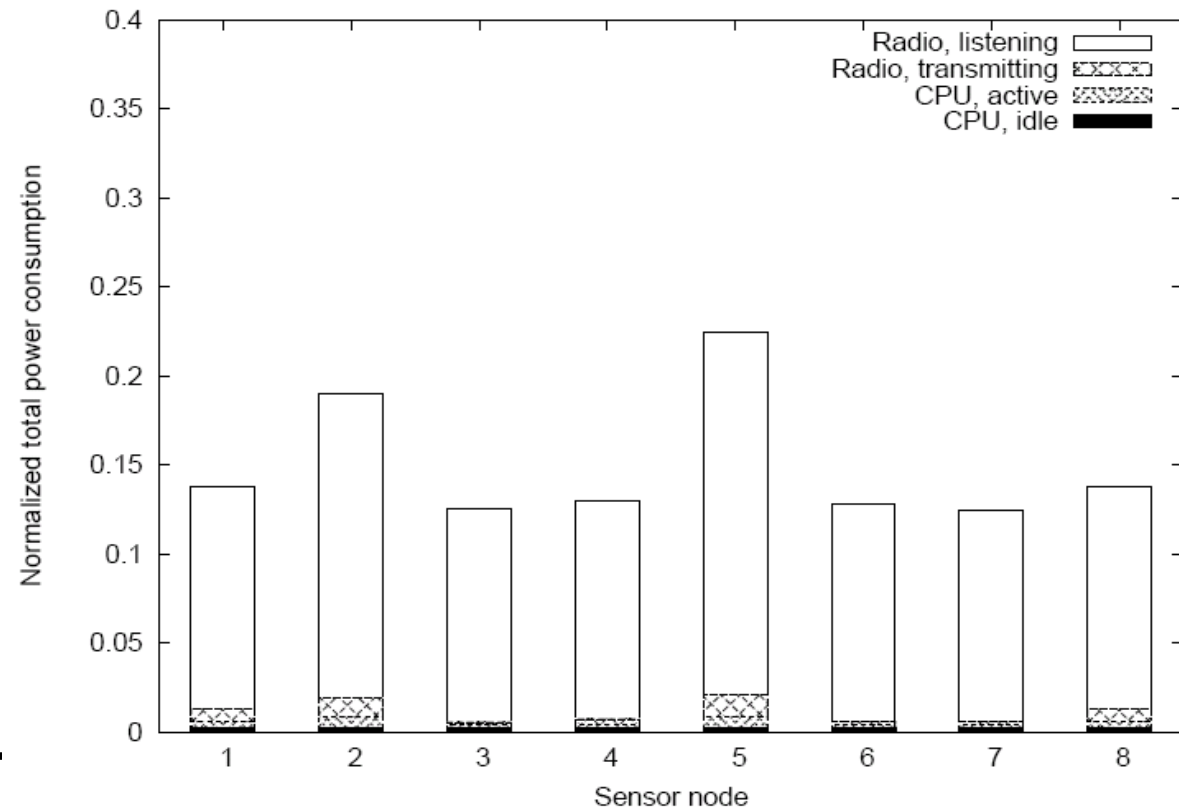


Figure 7. Power consumption: receiving

listening 與 receiving 一樣耗電

資源限制：功耗

- listening 與 receiving 一樣耗電，需透過軟體解決
- Solution 1: don't listen (ZigBee)
 - Listening nodes have a lot of energy (mains powered)
 - Problem: makes mesh networking difficult
- Solution 2: be smart about listening and energy
 - Power-saving radio mechanisms

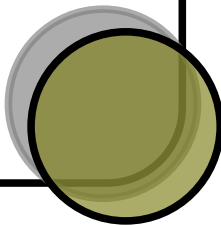


資源限制

memory

energy

bandwidth



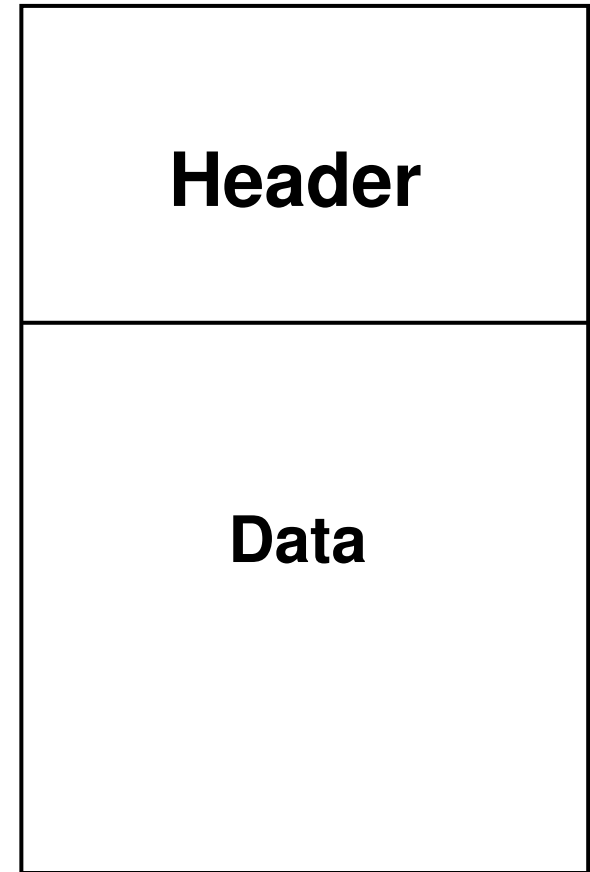
資源限制：TCP/IP bandwidth

~100 kbit/second

Messages are small

802.15.4 max size 128 bytes

TCP/IP headers are large



TCP/IP header is too long

TCP/IP headers 40 bytes

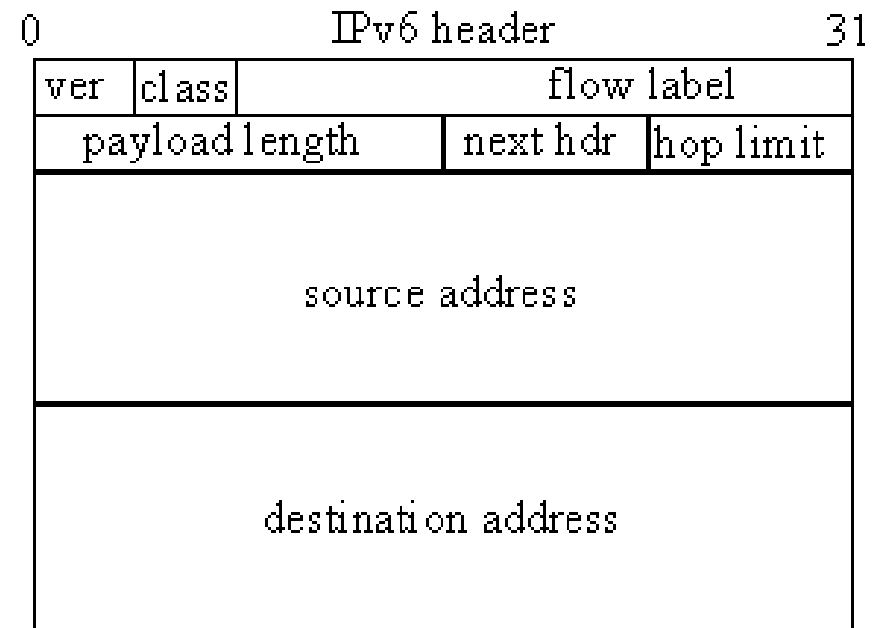
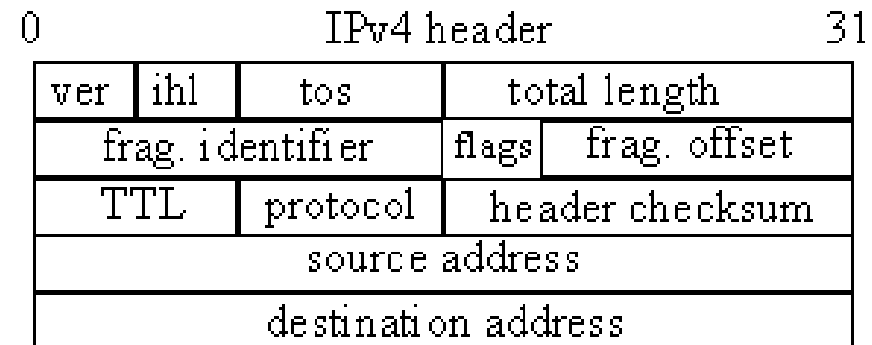
UDP/IP headers 28 bytes

802.15.4 frame size 128 bytes

27% – 36% overhead

IPv6 adds 20 bytes

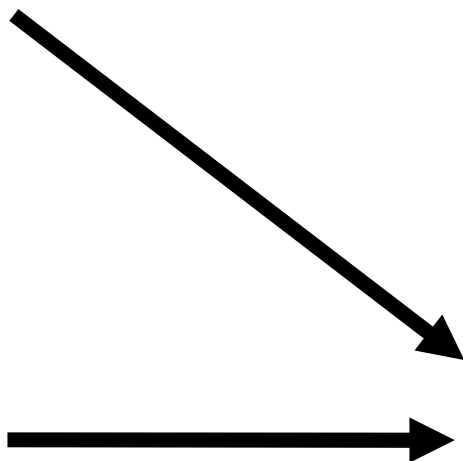
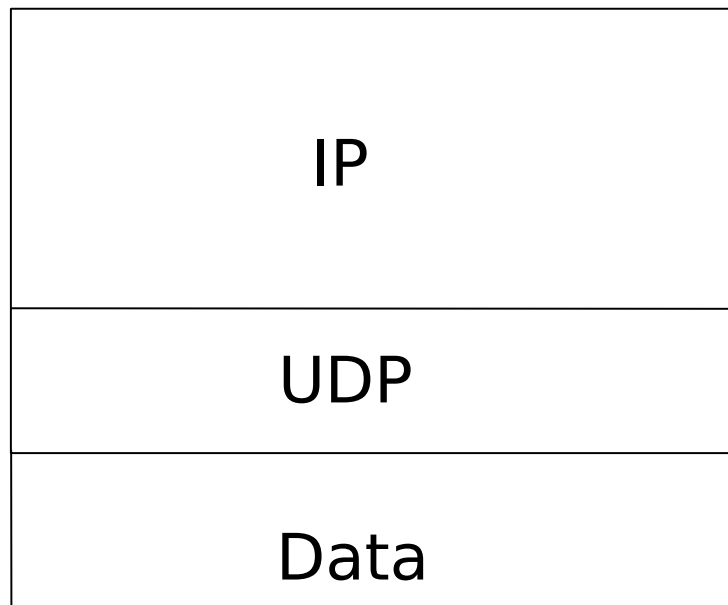
42% – 52% overhead



Header compression

僅傳遞重要資訊

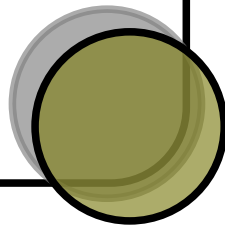
48 bytes



1-4 bytes

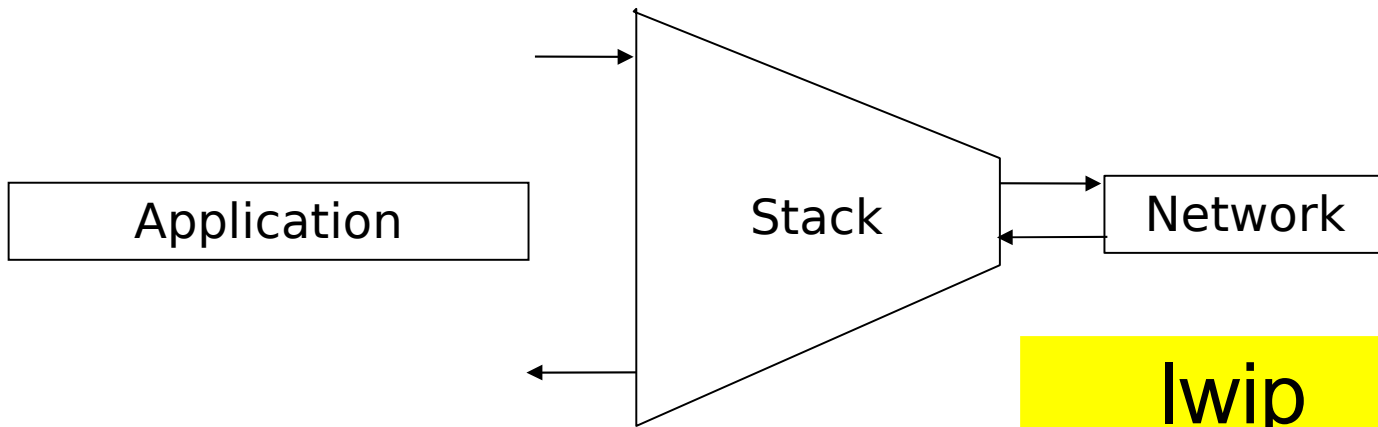


回歸現實：剪裁 與調適



TCP/IP stack design

Application driven



Application

UDP

TCP

IP

ICMP

Network

Network driven



Network

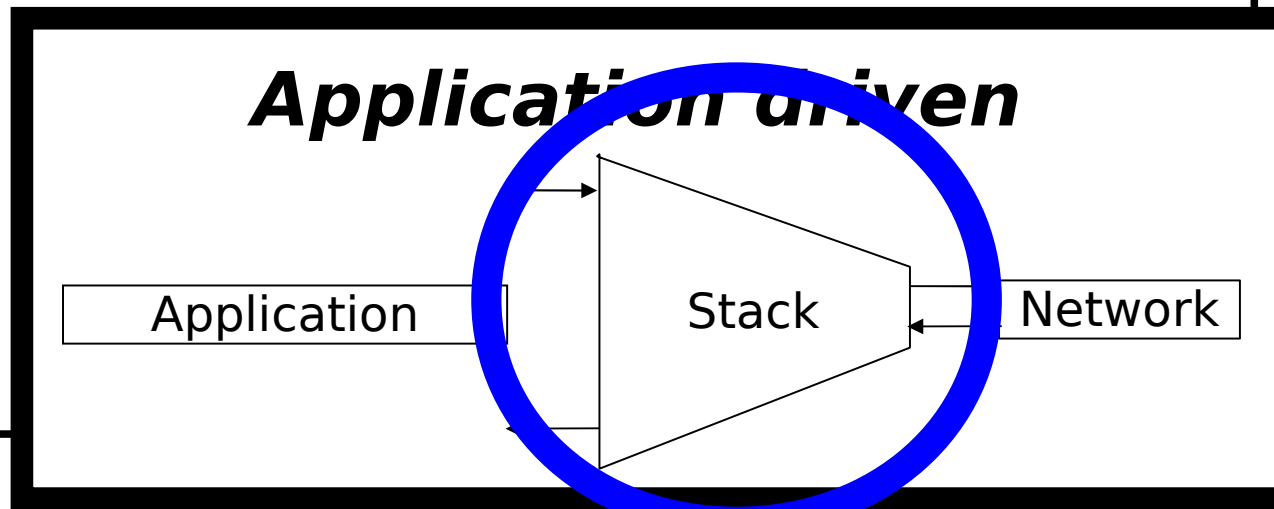
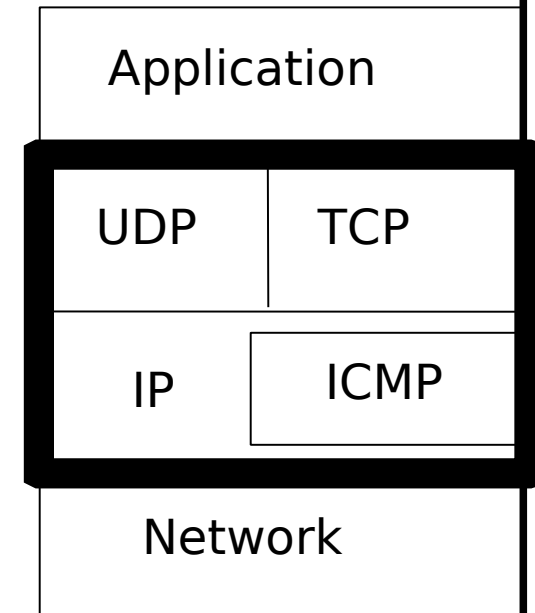
Stack

Application

uIP

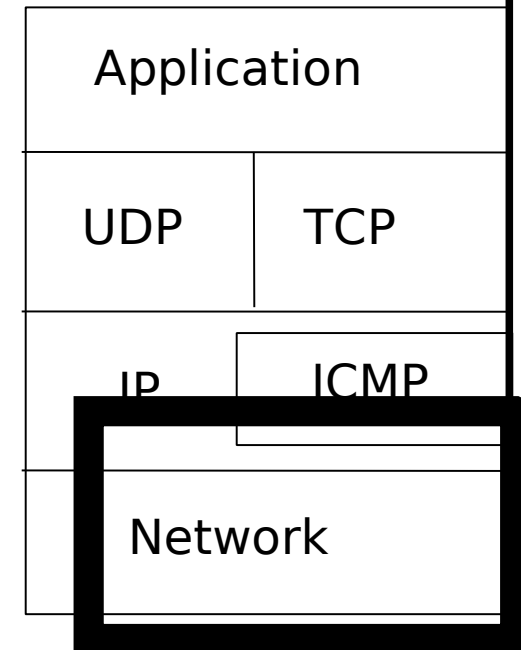
lwIP – Application driven

- <http://savannah.nongnu.org/projects/lwip/>
- lwIP – lightweight IP
- First release late 2000
- 40k code, 40k RAM
- Application driven design
 - Similar to Linux, BSD stacks
- Middle-end



μIP – smallest “full” TCP/IP stack

- <http://www.sics.se/~adam/uip/>
- First release 2001
- 3 - 5k code, 100 bytes - 2k RAM
- “Full” TCP/IP (RFC compliant)
- Used in Contiki OS

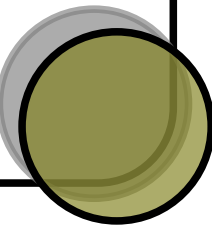


Network driven



μIP 技巧

- Shared packet buffer
- Lower throughput
- Event-driven APIs



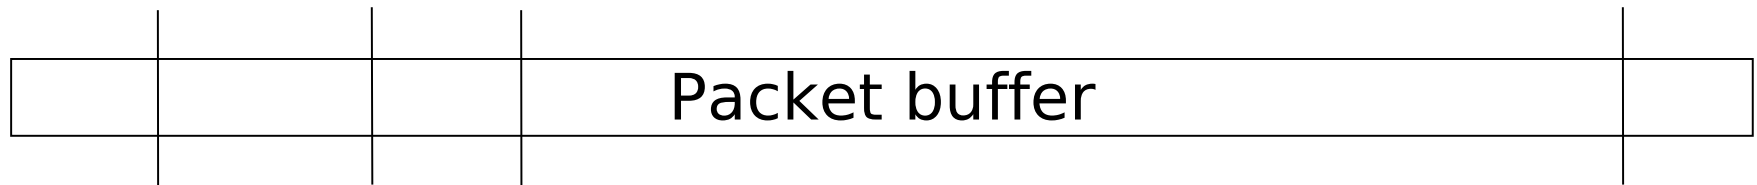
Shared packet buffer

- 所有封包 (outbound + inbound) 使用同一份緩衝區
- 緩衝區的大小決定 throughput



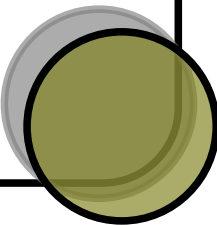
Shared packet buffer

- Implicit locking: single-threaded access
 - 1) Grab packet from network – put into buffer
 - 2) Process packet
 - Put reply packet in the same buffer
 - 3) Send reply packet into network



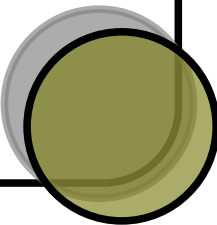
Lower Throughput

- μ IP trades throughput for RAM
 - Low RAM usage = low throughput
- 小的系統通常不會有太多資料
- 「具通訊的能力」會比 throughput 多寡
來得重要



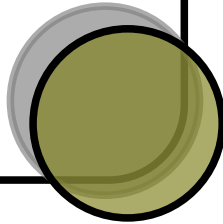
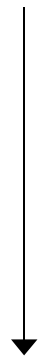
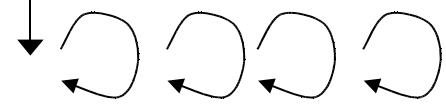
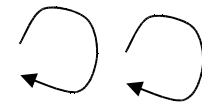
Event-driven APIs

- μ IP 沒有 BSD sockets
 - BSD sockets 建構於 threads 之上
 - Threads 引來 overhead (RAM)
- 因此，透過特製的 event-driven API
- co-routine
 - Applications are called by μ IP, call must return
- Protosockets – BSD socket-like API based on protothreads



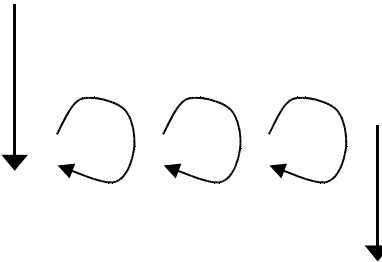
Protothread

```
int a_protothread(struct pt *pt) {  
    PT_BEGIN(pt);  
    /* ... */  
    PT_WAIT_UNTIL(pt, condition1);  
    /* ... */  
    if(something) {  
        /* ... */  
        PT_WAIT_UNTIL(pt, condition2);  
        /* ... */  
    }  
    PT_END(pt);  
}
```

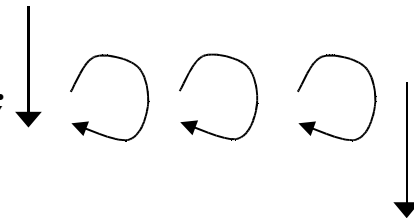


Hierarchical Protothreads

```
int a_protothread(struct pt *pt) {  
    static struct pt child_pt;  
  
    PT_BEGIN(pt);  
  
    PT_INIT(&child_pt);  
    PT_WAIT_UNTIL(pt2(&child_pt) != 0);  
  
    PT_END(pt);  
}
```

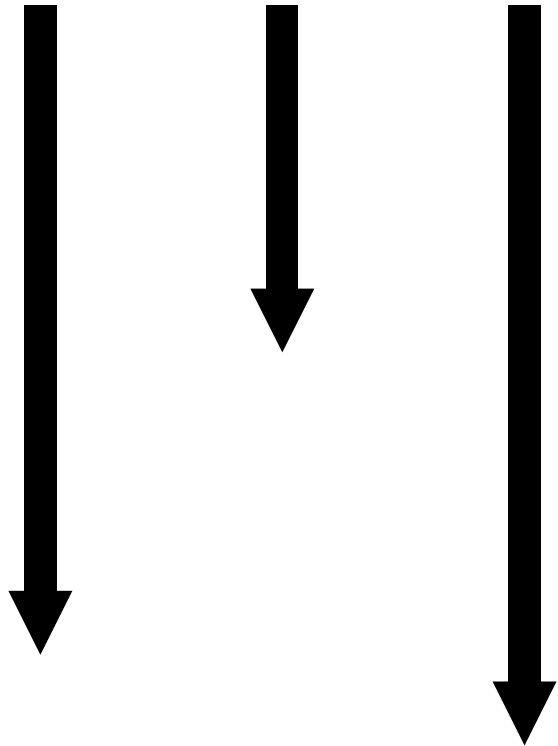


```
int pt2(struct pt *pt) {  
    PT_BEGIN(pt);  
  
    PT_WAIT_UNTIL(pt, condition);  
  
    PT_END(pt);  
}
```

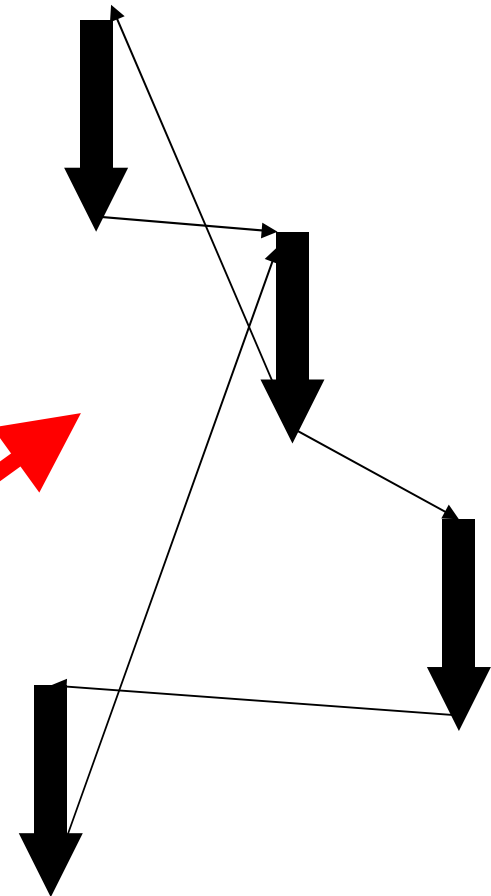


Threads vs. Events

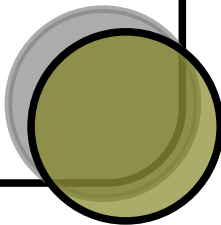
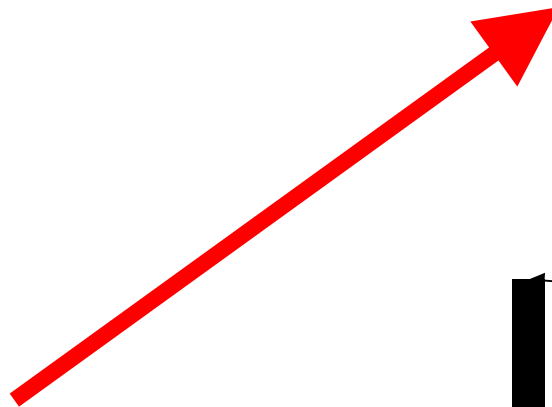
Threads: sequential code flow



Events: unstructured code flow



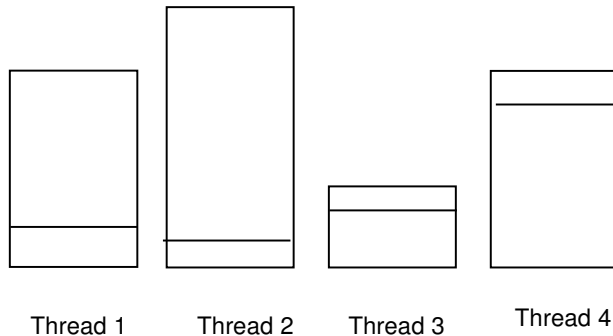
No blocking wait!



Events require one stack

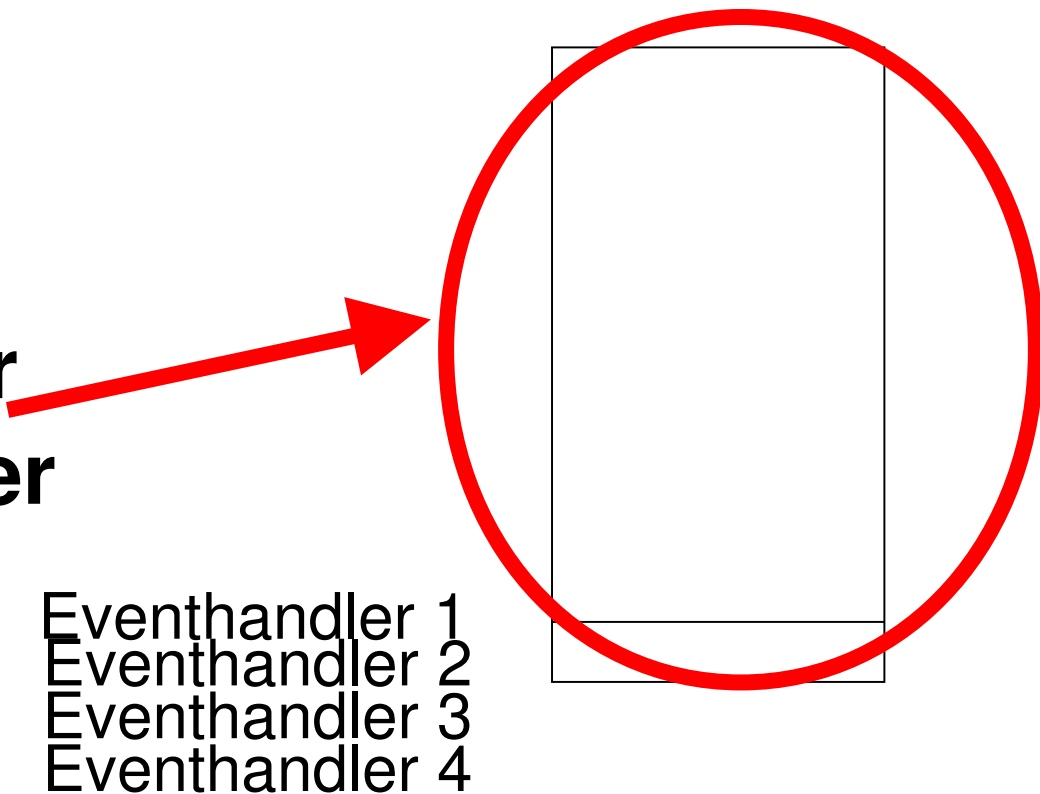
Threads require per-thread stack memory

Four threads, each with its own stack



- Four event handlers, one stack

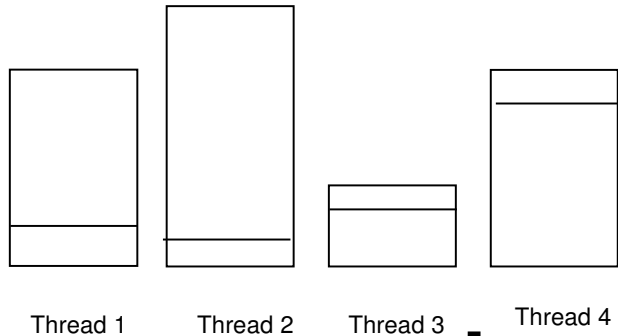
Stack is reused for every event handler



Protothreads require one stack

Threads require per-thread stack memory

Four threads, each with its own stack

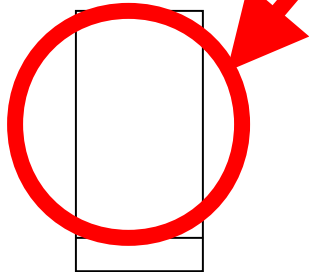


- Four protothreads, one stack

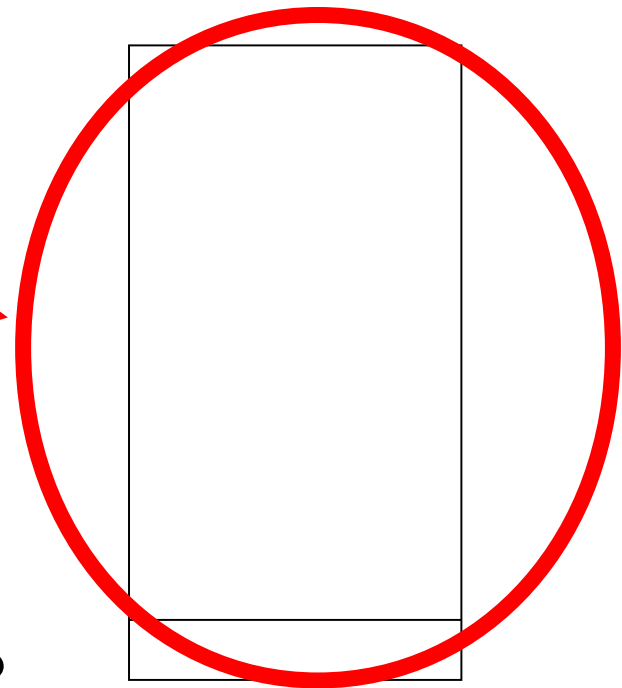
Just like events

Events require one stack

- Four event handlers, one stack



Protothread 1
Protothread 2
Protothread 3
Protothread 4



Event-driven APIs

```
void example2_app(void) {
    struct example2_state *s =
        (struct example2_state *)uip_conn->appstate;

    if(uip_connected()) {
        s->state = WELCOME_SENT;
        uip_send("Welcome!\n", 9);
        return;
    }

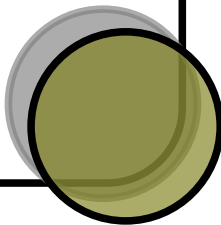
    if(uip_acked() &&
        s->state == WELCOME_SENT) {
        s->state = WELCOME_ACKED;
    }

    if(uip_newdata()) {
        uip_send("ok\n", 3);
    }

    if(uip_rexmit()) {
        switch(s->state) {
            case WELCOME_SENT:
                uip_send("Welcome!\n", 9);
                break;
            case WELCOME_ACKED:
                uip_send("ok\n", 3);
                break;
        }
    }
}
```

Event-driven APIs

- Event-driven API 不見得適用所有程式
- Protosockets: sockets-like API using protothreads
 - Extremely lightweight stackless threads
 - 2 bytes per-thread state, no stack
- Protothreads 允許“blocking”，本質是循序的
- overhead 遠小於真實的 thread



Event-driven APIs

```
PT_THREAD(smtp_protothread(void))
{
    PSOCK_BEGIN(s);

    PSOCK_READTO(s, '\n');

    if(strncmp(inputbuffer, "220", 3) != 0) {
        PSOCK_CLOSE(s);
        PSOCK_EXIT(s);
    }

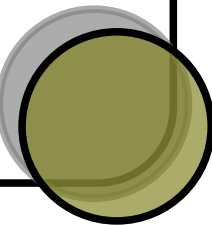
    PSOCK_SEND(s, "HELO ", 5);
    PSOCK_SEND(s, hostname, strlen(hostname));
    PSOCK_SEND(s, "\r\n", 2);

    PSOCK_READTO(s, '\n');

    if(inputbuffer[0] != '2') {
        PSOCK_CLOSE(s);
        PSOCK_EXIT(s);
    }
}
```

Proof-of-concept TCP/IP stacks

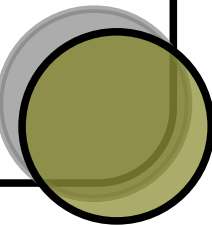
- phpstack – TCP/IP stack, webserver written in PHP
<http://www.sics.se/~adam/phpstack/>
- miniweb – TCP/IP stack, webserver using 30 bytes of RAM
<http://www.sics.se/~adam/miniweb/>



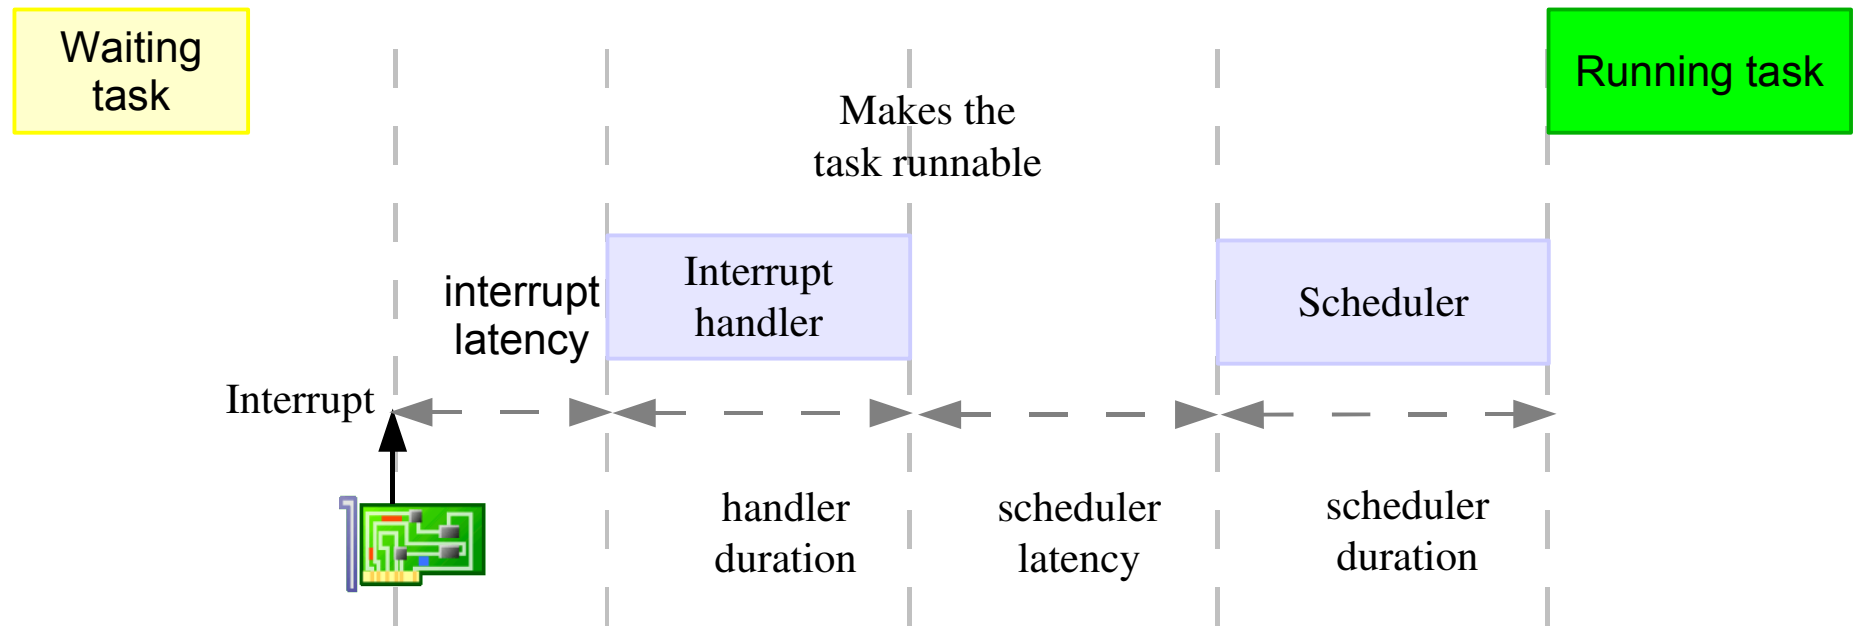
技術討論

討論

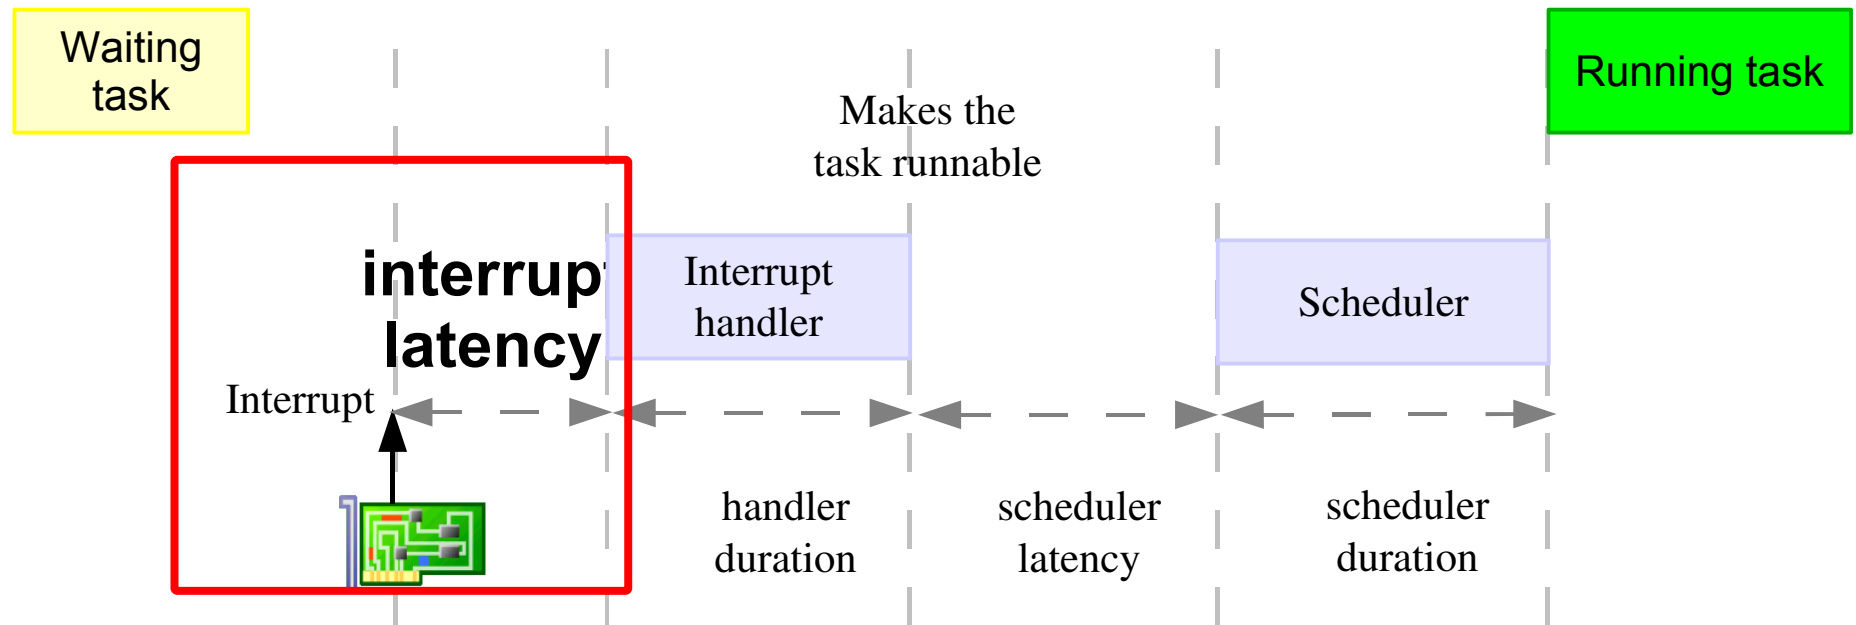
- μ IP 適用於低階的環境
- 中階或網路需求較高者，應用 lwIP 或 BSD socket
- Thread preemption



情境：Process 正等待 Device I/O(由中斷觸發) 的結束，
方可繼續執行



**kernel latency = interrupt latency +
handler duration +
scheduler latency +
scheduler duration**



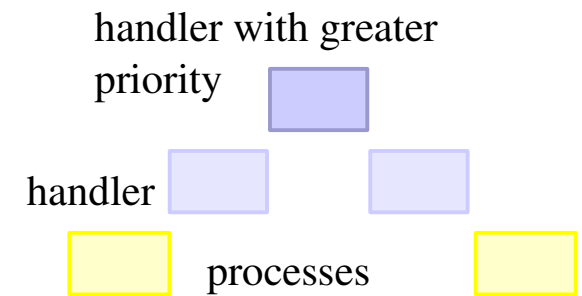
來源

Interrupts disabled by kernel code:
spinlocks, driver code...

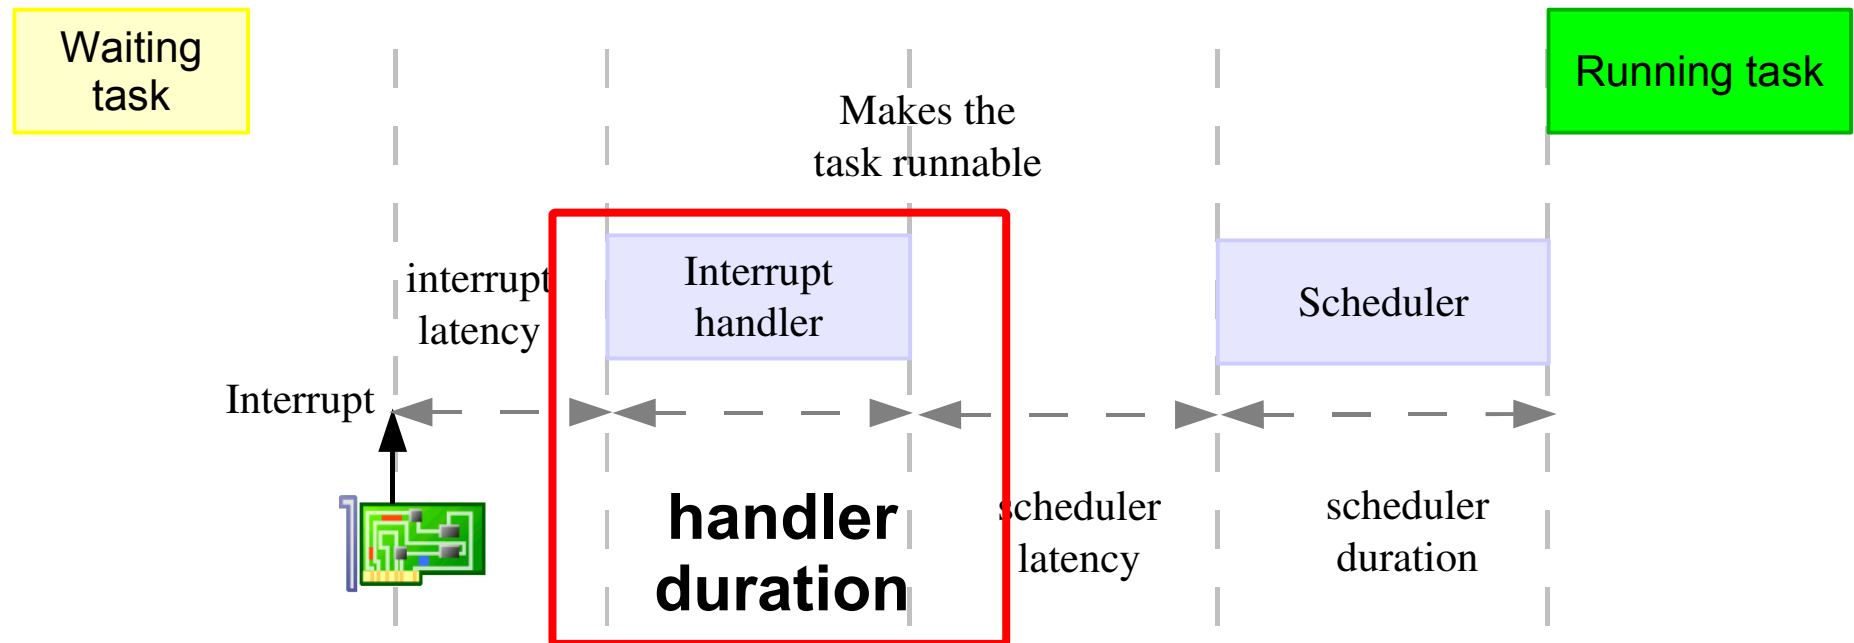
Bad driver using the fast interrupt
mode (should be reserved to timer
interrupts).

Other interrupts processed before.

Interrupt context, managed by the CPU



Process context, managed by the scheduler



執行 interrupt handler 的時間，在以下情況會更嚴重：

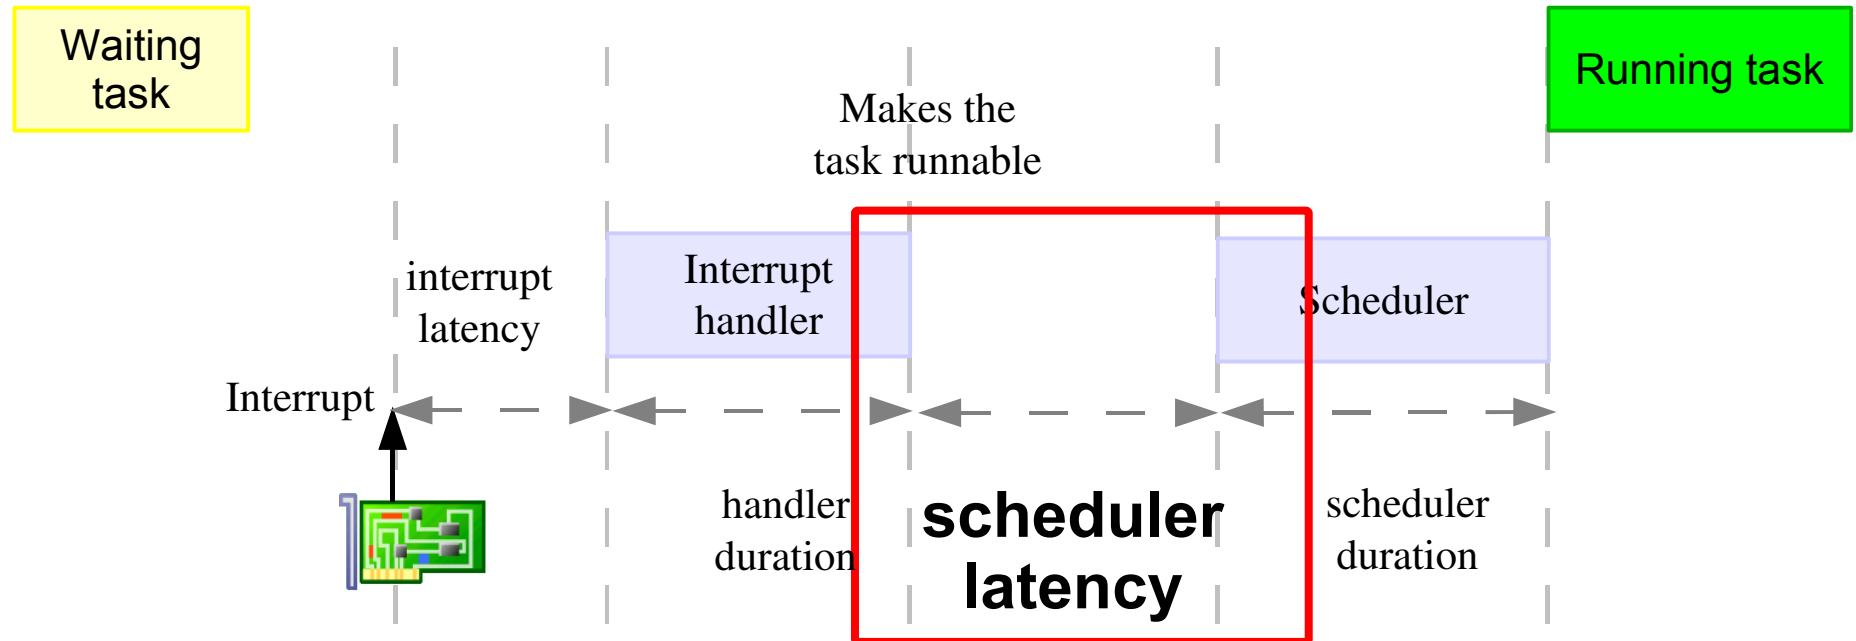
Preemption by other interrupts.

Interrupt handler with a softirq component

Interrupts disabled by kernel code

Bad driver using the fast interrupt mode

Other interrupts processed before.



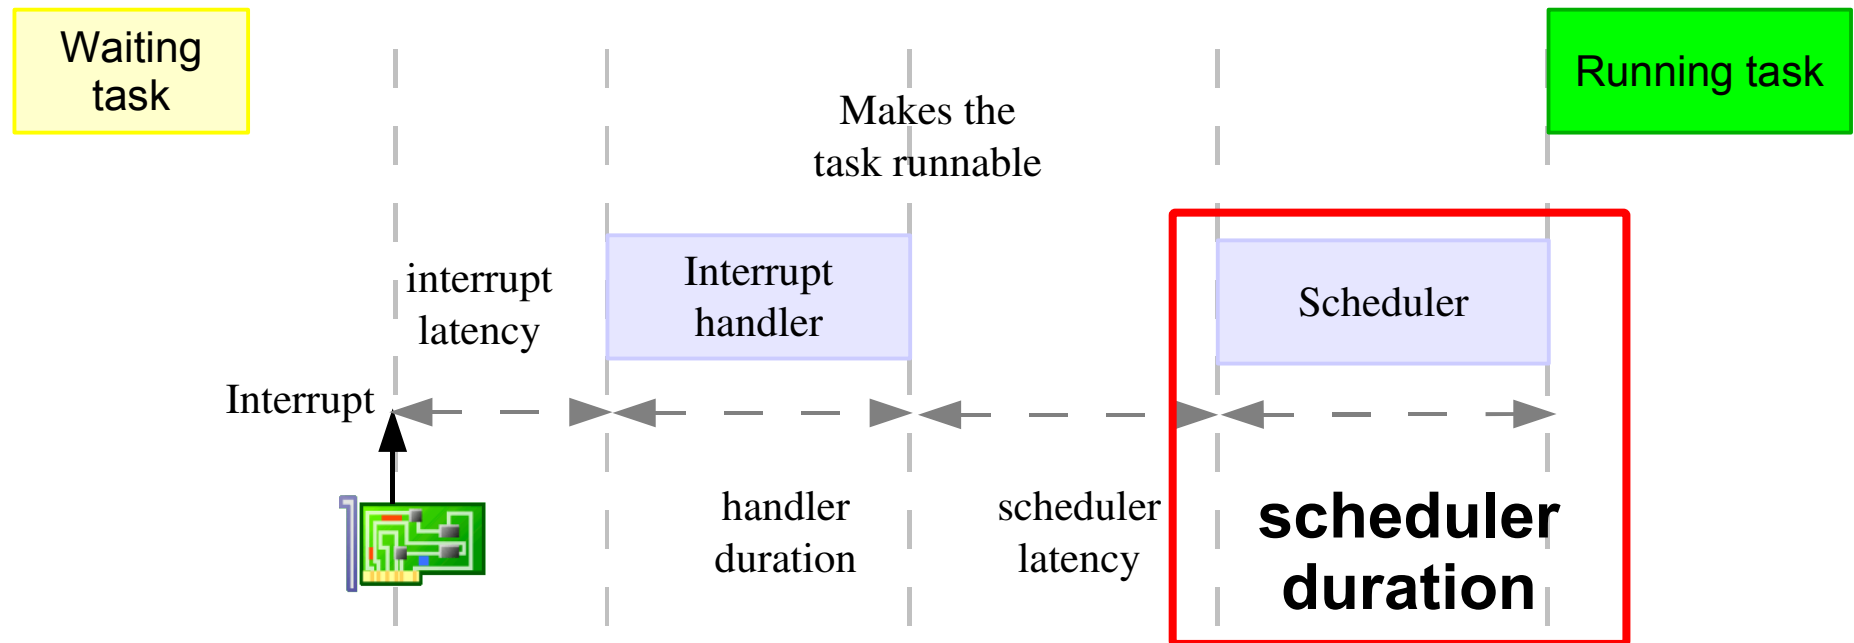
用於 scheduler 的時間，原因：

- Kernel code not preemptible.

- Need to wait for returning from interrupts or from system calls.

- Interrupts disabled for a long time in driver code.

- Can cause a timer interrupt to be missed.



執行 scheduler 並切換到新的 task 所耗費的時間

執行 scheduler 的時間為常數

context switching time

Restoring processor registers and stack for the new process.

Restoring the address space of the new process
(except for threads sharing the same address space).

結論

無所不在的嵌入式系統

Invisible Computer

機電整合與自由軟體的機會
作中學

