Faults inside System Software

Jim Huang (黃敬群) <**jserv**@0xlab.org>

June 6, 2013 / NCU, Taiwan

Rights to copy

creative commons
COMMONS DEED

Corrections, suggestions, contributions and translations are welcome!

Attribution – ShareAlike 3.0

Latest update: June 8, 2013

You are free

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Under the following conditions

Attribution. You must give the original author credit.

Share Alike. If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.

- For any reuse or distribution, you must make clear to others the license terms of this work.
- Any of these conditions can be waived if you get permission from the copyright holder.

Your fair use and other rights are in no way affected by the above.

License text: http://creativecommons.org/licenses/by-sa/3.0/legalcode

zeczec x haniboi.com

UP UP

zeczec x haniboi.com

UP UP

Goals of This Presentation

- Analysis of Large-scale system software
- Diagnose faults inside system software, especially for device drivers
- Deal with faulty device driver implementation

Agenda

- General Analysis about Faulty system software
- Approaches to Deal
  - Runtime Isolation
  - Static Analysis

General Analysis about Faulty System Software

Some statistics

- Drivers cause 85% of Windows XP crashes.
  - Michael M. Swift, Brian N. Bershad, Henry M. Levy: *"Improving the Reliability of Commodity Operating Systems"*, SOSP 2003
- Error rate in Linux drivers is 3x (maximum: 10x) higher than for the rest of the kernel
  - Life expectancy of a bug in the Linux kernel (~2.4): 1.8 years
  - Andy Chou, Junfeng Yang, Benjamin Chelf, Seth Hallem, Dawson R. Engler: *"An Empirical Study of Operating System Errors"*, SOSP 2001

Some statistics

- Causes for driver bugs
  - 23% programming error
  - 38% mismatch regarding device specification
  - 39% OS-driver-interface misconceptions
  - Leonid Ryzhyk, Peter Chubb, Ihor Kuz and Gernot Heiser: *"Dingo: Taming device drivers"*, EuroSys 2009

Anecdote: Linux e1000 NVRAM bug

- **[Aug 8, 2008]** Bug report: e1000 PCI-X network cards rendered broken by Linux 2.6.27-rc
  - overwritten NVRAM on card
- **[Oct 1, 2008]** Intel releases quickfix
  - map NVRAM somewhere else
- **[Oct 15, 2008]** Reason found:
  - dynamic ftrace framework tries to patch __init code, but .init sections are unmapped after running init code
  - NVRAM got mapped to same location
  - scary cmpxchg() behavior on I/O memory
- **[Nov 2, 2008]** dynamic ftrace reworked for Linux 2.6.28-rc3

FTrace & **NIC driver**!

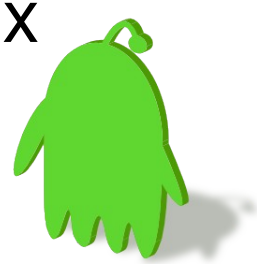

Linux Device Driver bugs

[Dingo: Taming device drivers, 2009]

Driver	#loc	#bugs
USB		
RTL8150 USB-to-Ethernet adapter	827	16
EL1210a USB-to-Ethernet adapter	710	2
KL5kusb101 USB-to-Ethernet apapter	925	15
Generic USB network driver	1028	45
USB hub	2234	67
USB-to-serial converter	989	50
USB mass storage	803	23
Firewire		
IEEE1394 Ethernet controller	1413	22
SBP-2 transport protocol	1713	46
PCI		
Mellanox InfiniHost InfiniBand adapter	11718	123
BNX2 Ethernet adapter	5412	51
i810 frame buffer	2920	16
CMI8338 audio	2660	22
		498

Linux version 3.0

- consists of
  - 7702 features
  - 893 Kconfig files
  - 31281 source files
  - 88897 #ifdef blocks

Even worse...

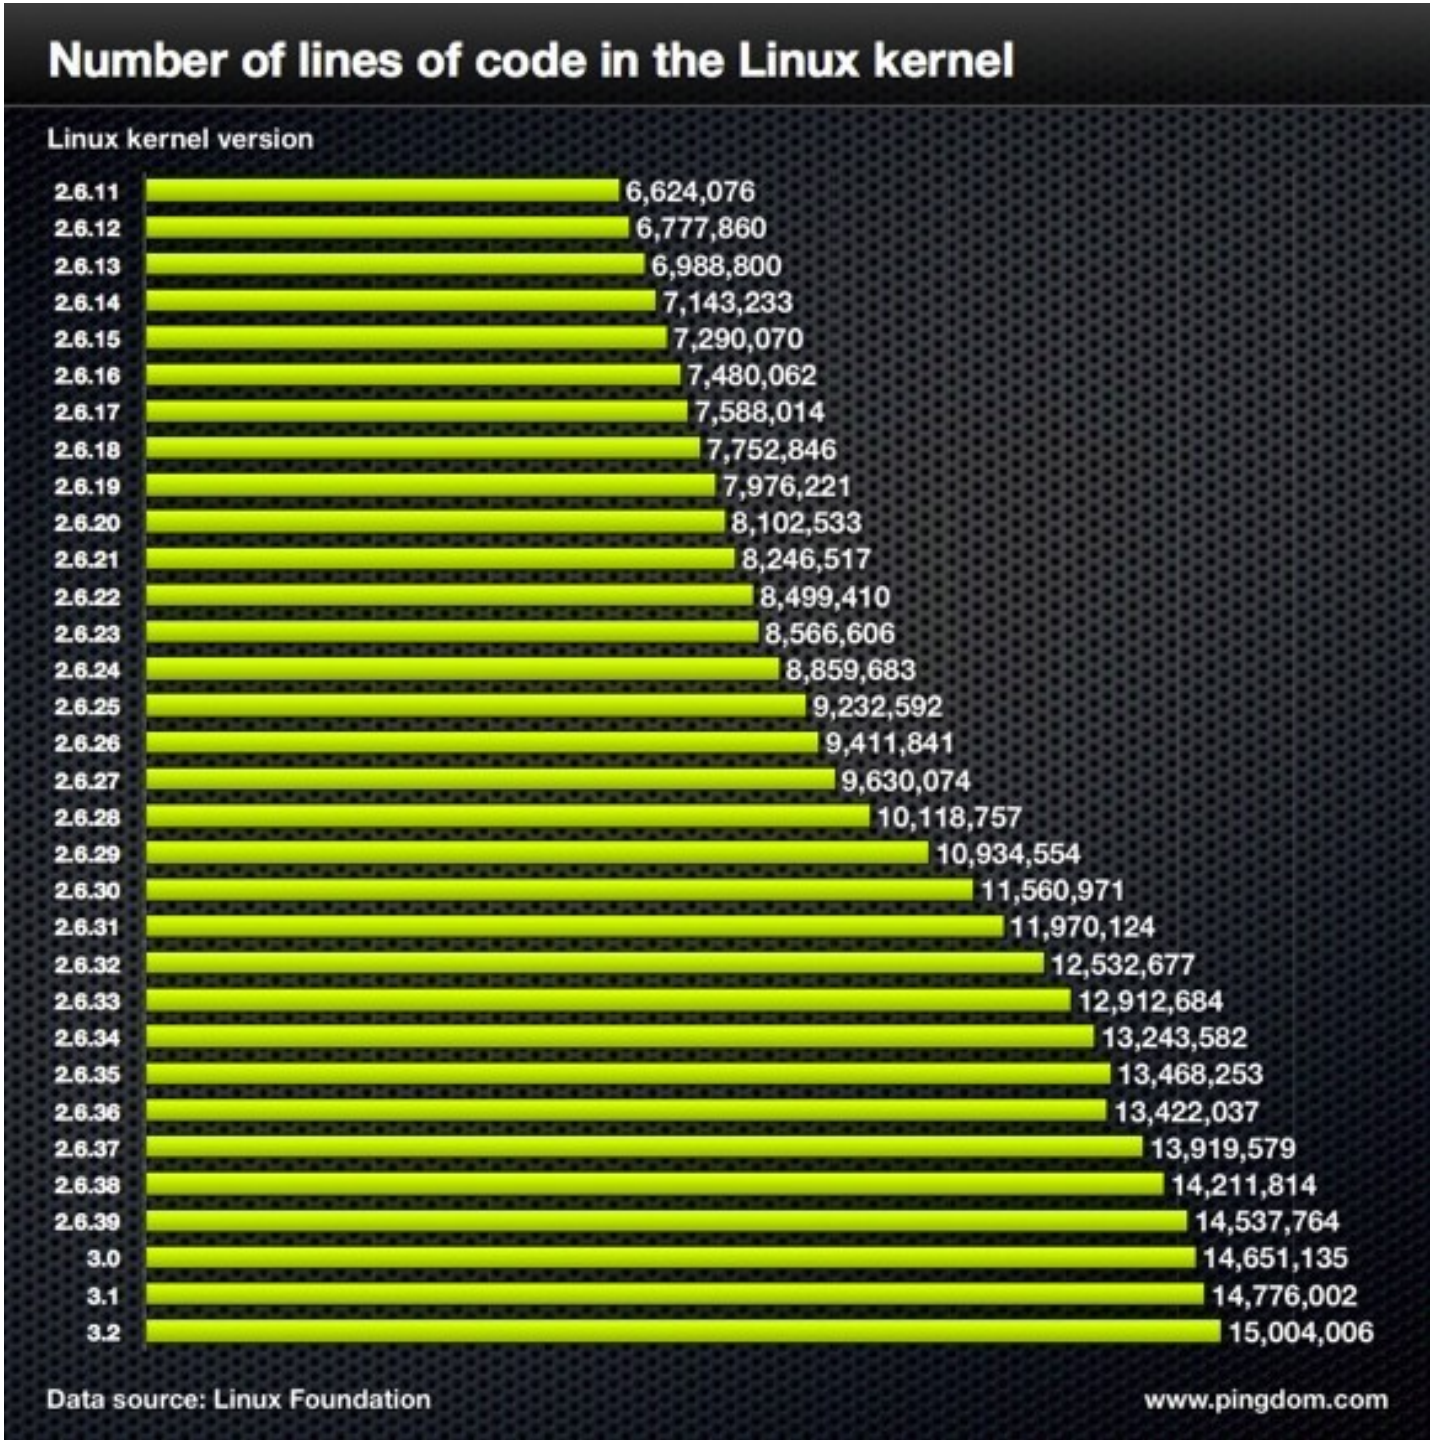
Number of lines of code in the Linux kernel
Linux kernel version
2.6.11 6,624,076
2.6.12 6,777,860
2.6.13 6,988,800
2.6.14 7,143,233
2.6.15 7,290,070
2.6.16 7,480,062
2.6.17 7,588,014
2.6.18 7,752,846
2.6.19 7,976,221
2.6.20 8,102,533
2.6.21 8,246,517
2.6.22 8,499,410
2.6.23 8,566,606
2.6.24 8,859,683
2.6.25 9,232,592
2.6.26 9,411,841
2.6.27 9,630,074
2.6.28 10,118,757
2.6.29 10,934,554
2.6.30 11,560,971
2.6.31 11,970,124
2.6.32 12,532,677
2.6.33 12,912,684
2.6.34 13,243,582
2.6.35 13,468,253
2.6.36 13,422,037
2.6.37 13,919,579
2.6.38 14,211,814
2.6.39 14,537,764
3.0 14,651,135
3.1 14,776,002
3.2 15,004,006
Data source: Linux Foundation
www.pingdom.com

Top 10 contributors to the Linux kernel since version 2.6.36

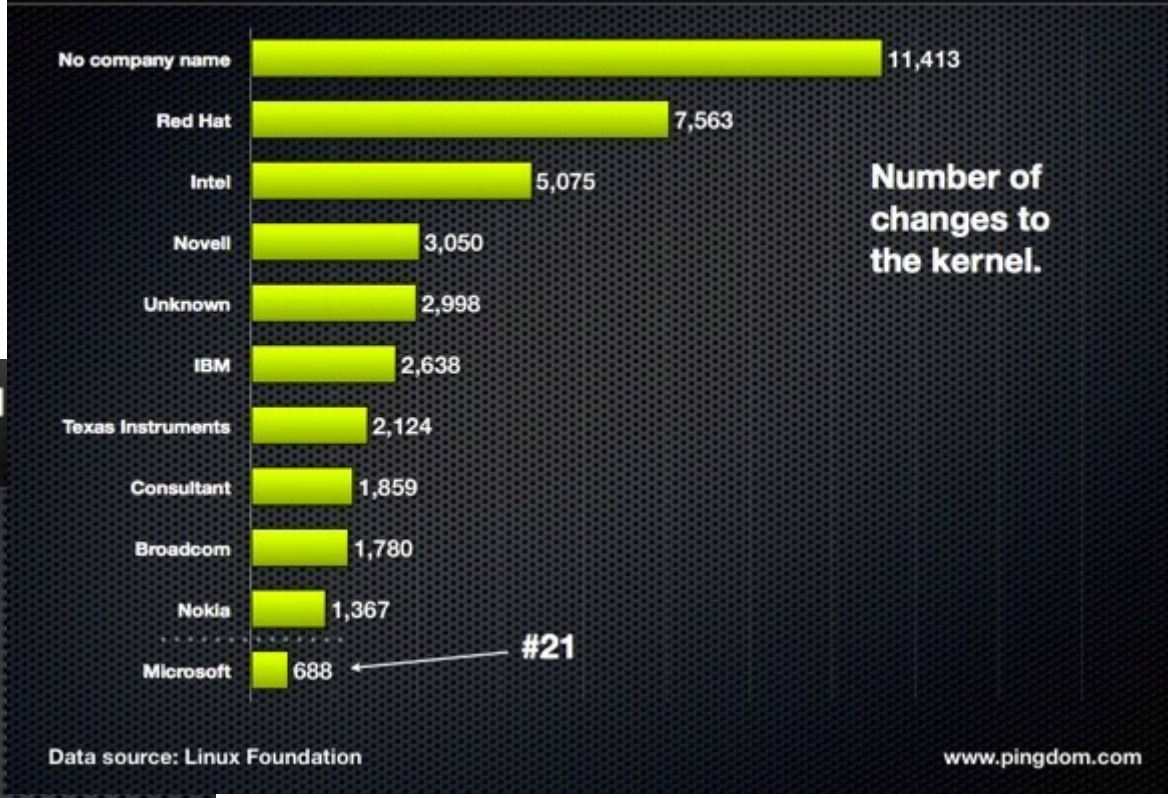

Number of lines of code added to the Linux kernel per each day of development

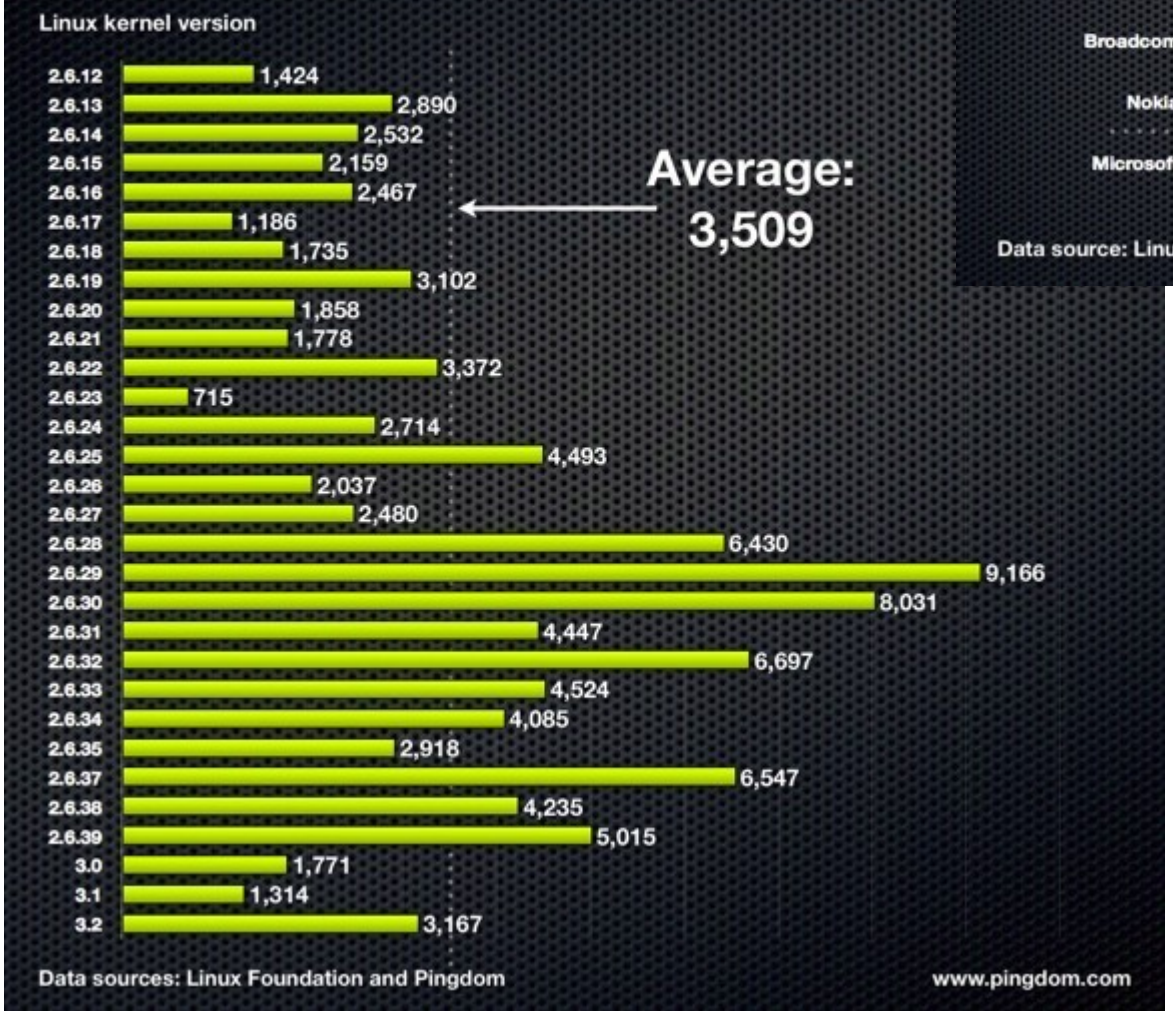

System Layout

- Devices connected by buses (USB, PCI, PCIx)
- Host chipset (DMA logic, IRQ controller) connects buses and CPU

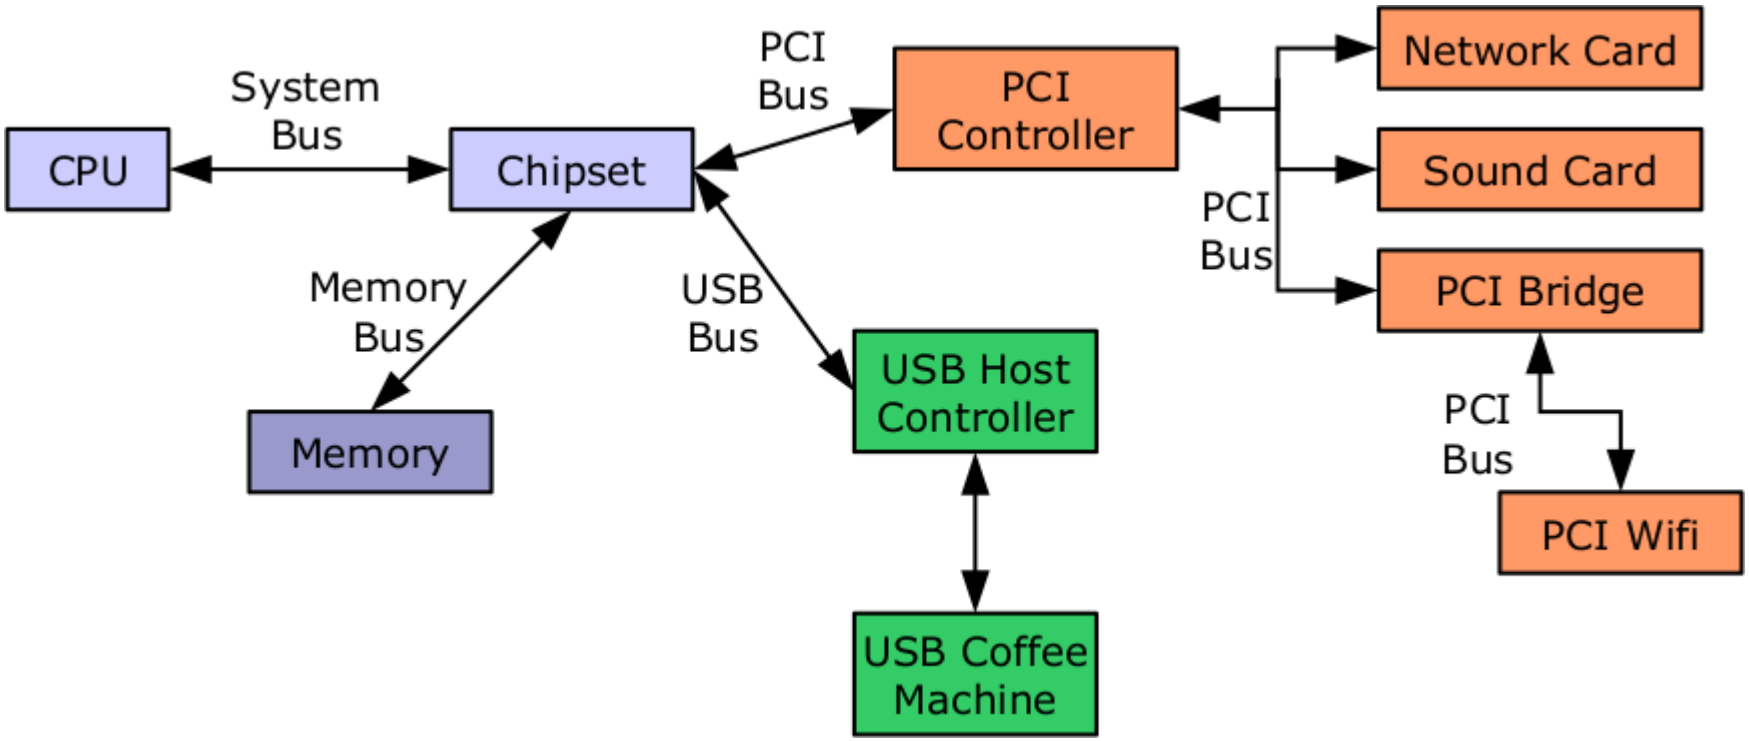

Example: Intel 925x chipset

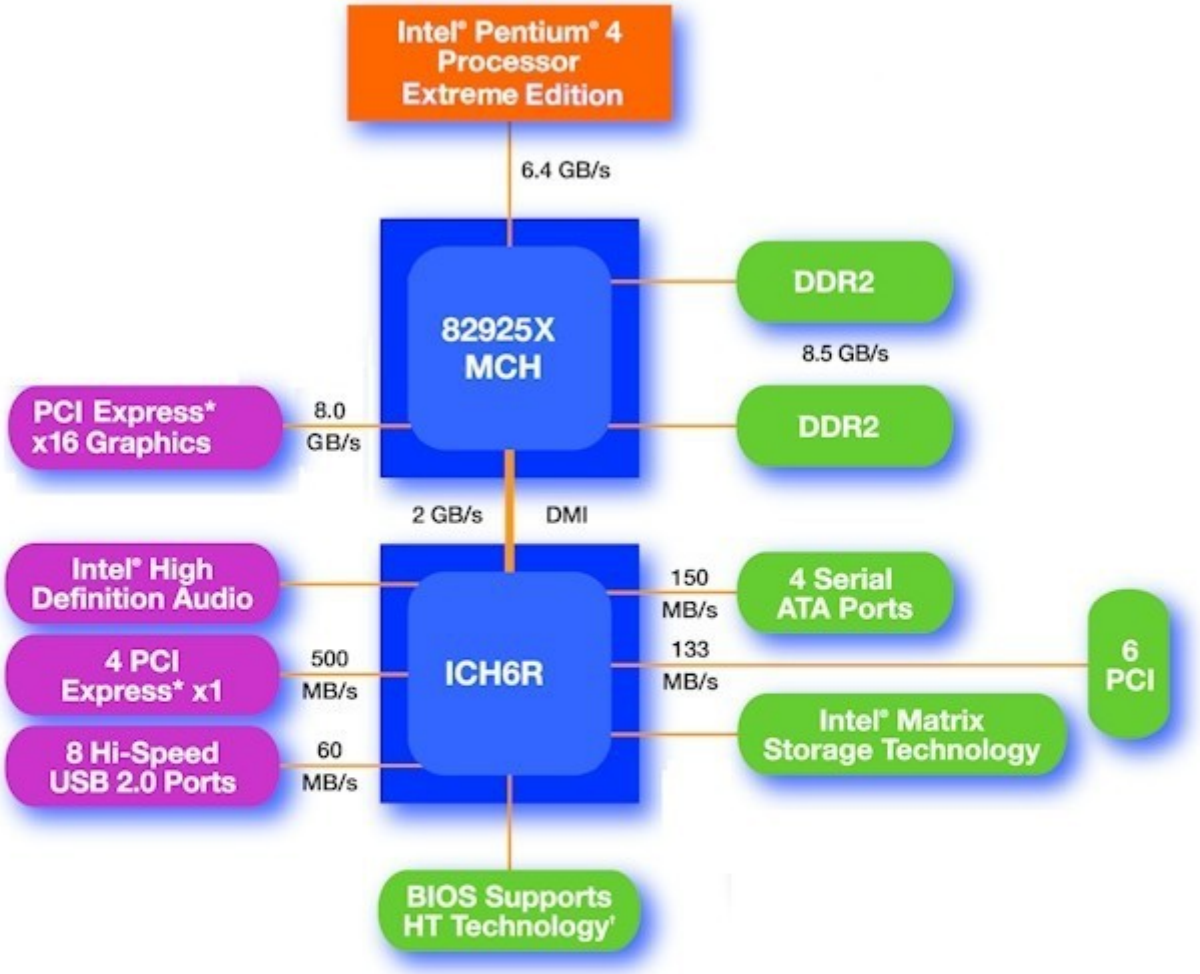

† Hyper-Threading (HT) Technology requires a computer system with an Intel® Pentium® 4 processor supporting HT Technology and a HT Technology enabled chipset, BIOS and operating system. Performance will vary depending on the specific hardware and software you use. See www.intel.com/info/hyperthreading for more information including details on which processors support HT Technology.

Bus & Devices

- Problem: more and more devices
  - need means of dynamic device discovery
- Probing
  - try out every driver to see if it works
- Plug-n-Play
  - first try of dynamic system description
  - device manufacturers provide unique IDs
- PCI: dedicated config space
- ACPI: system description without relying on underlying bus/chipset

Bus: USB

- Intel, 1996
- Tree of devices
  - root = Host Controller (UHCI, OHCI, EHCI)
  - Device drivers use Host Controller (HC) to communicate with their device via USB Request Blocks (URBs)
  - USB is a serial bus
    - HC serializes URBs
- Wide range of device classes (input, storage, peripherals, ...)
  - classes allow generic drivers

Attack iOS through USB charger!

- BlackHat 2013
  - MACTANS: INJECTING MALWARE INTO IOS DEVICES VIA MALICIOUS CHARGERS
  - http://www.blackhat.com/us-13/briefings.html#Lau
- "we demonstrate how an iOS device can be compromised within one minute of being plugged into a malicious charger. We first examine Apple's existing security mechanisms to protect against arbitrary software installation, then describe how USB capabilities can be leveraged to bypass these defense mechanisms."

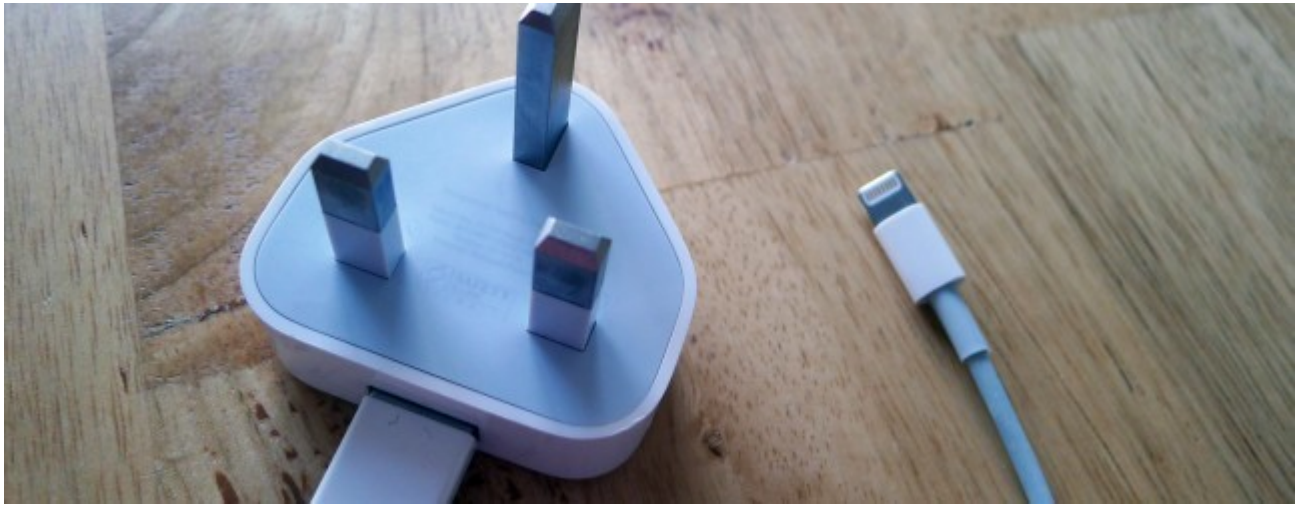

Device Driver Model

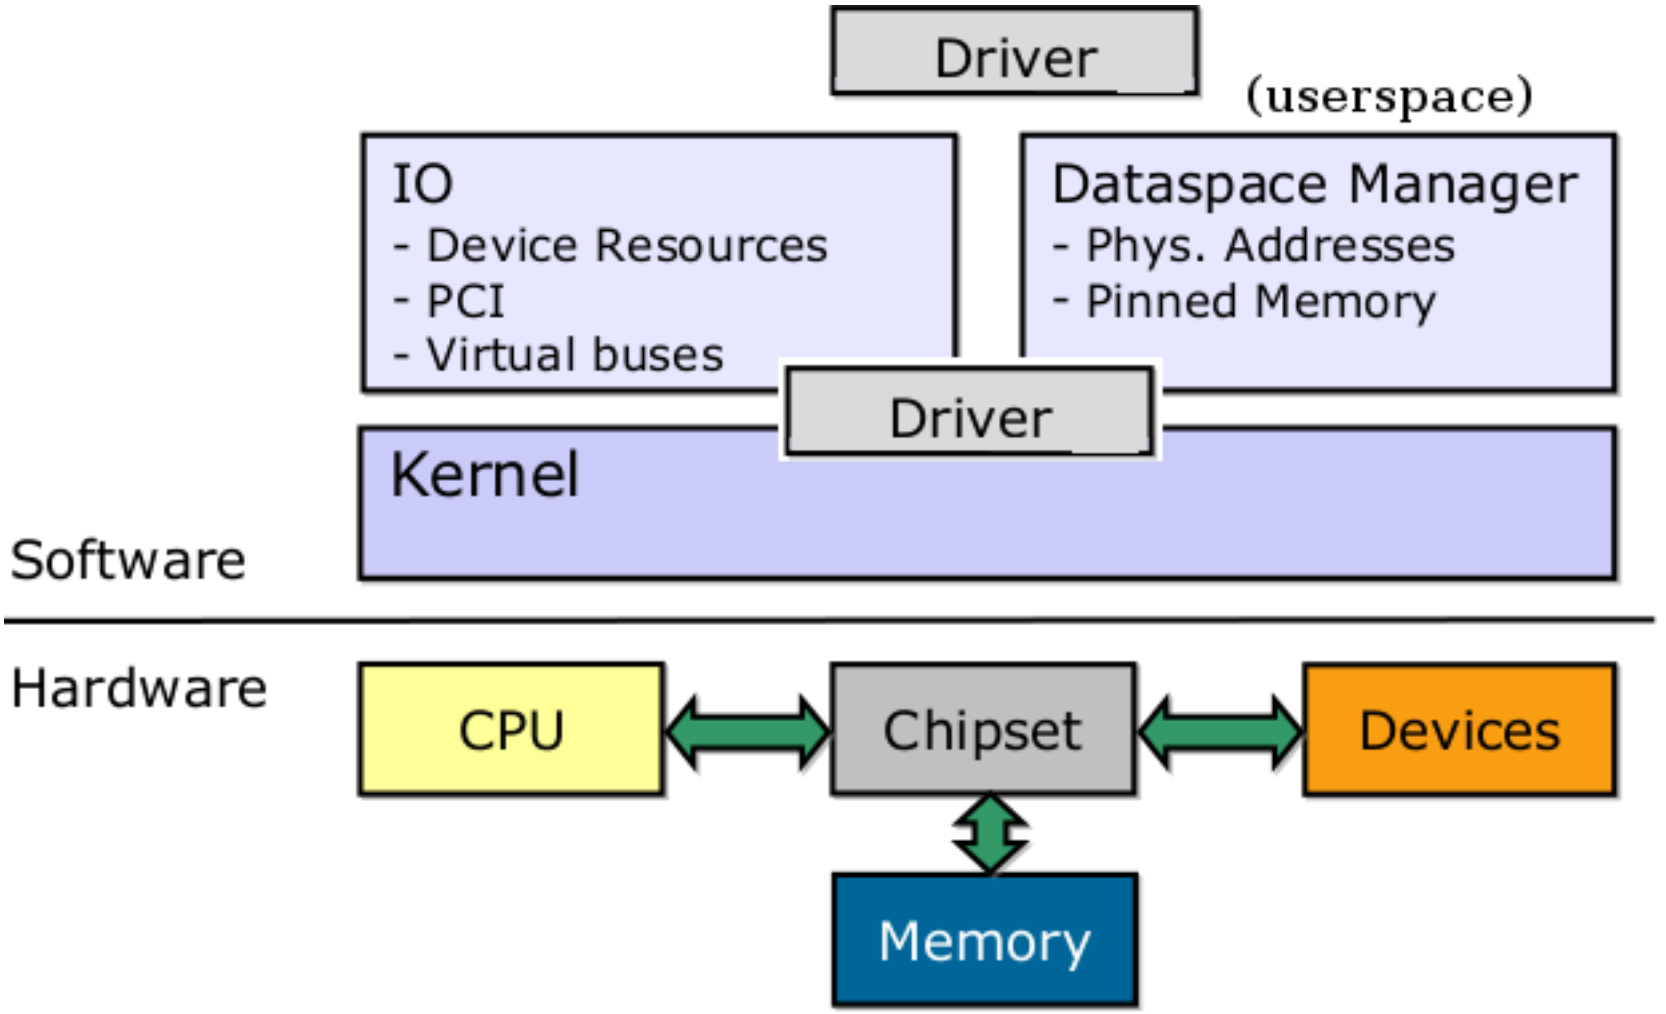

Bugs in Linux Device Driver

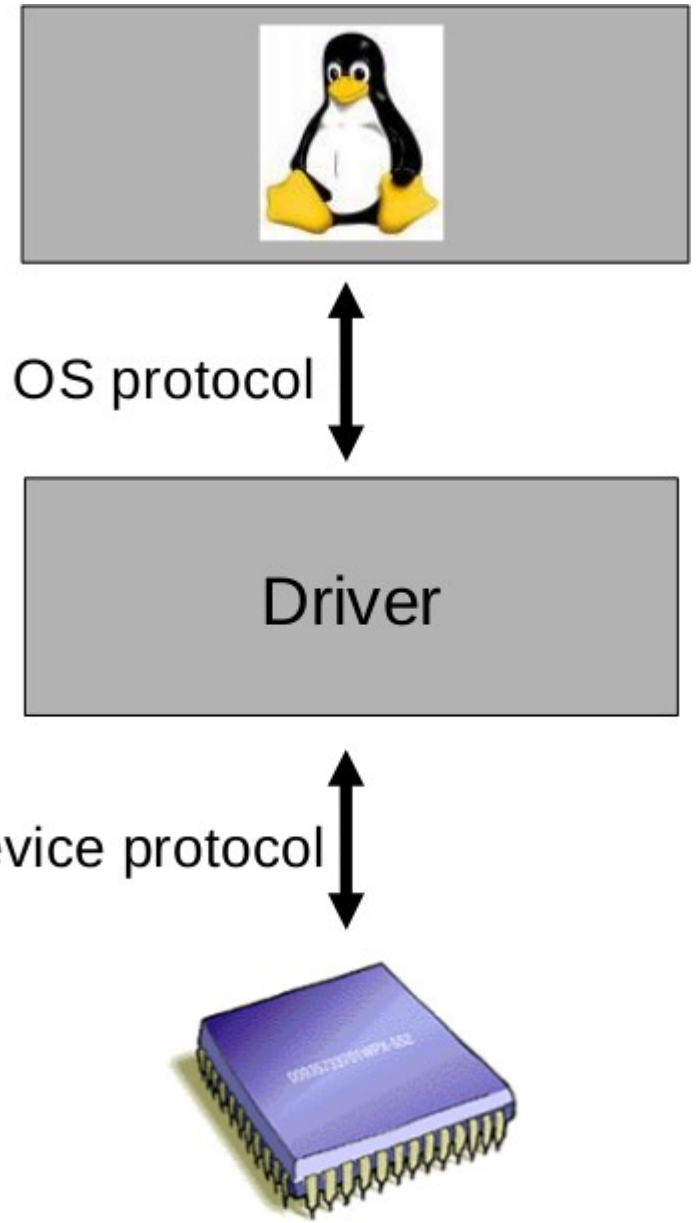

Bugs in Linux Device Driver

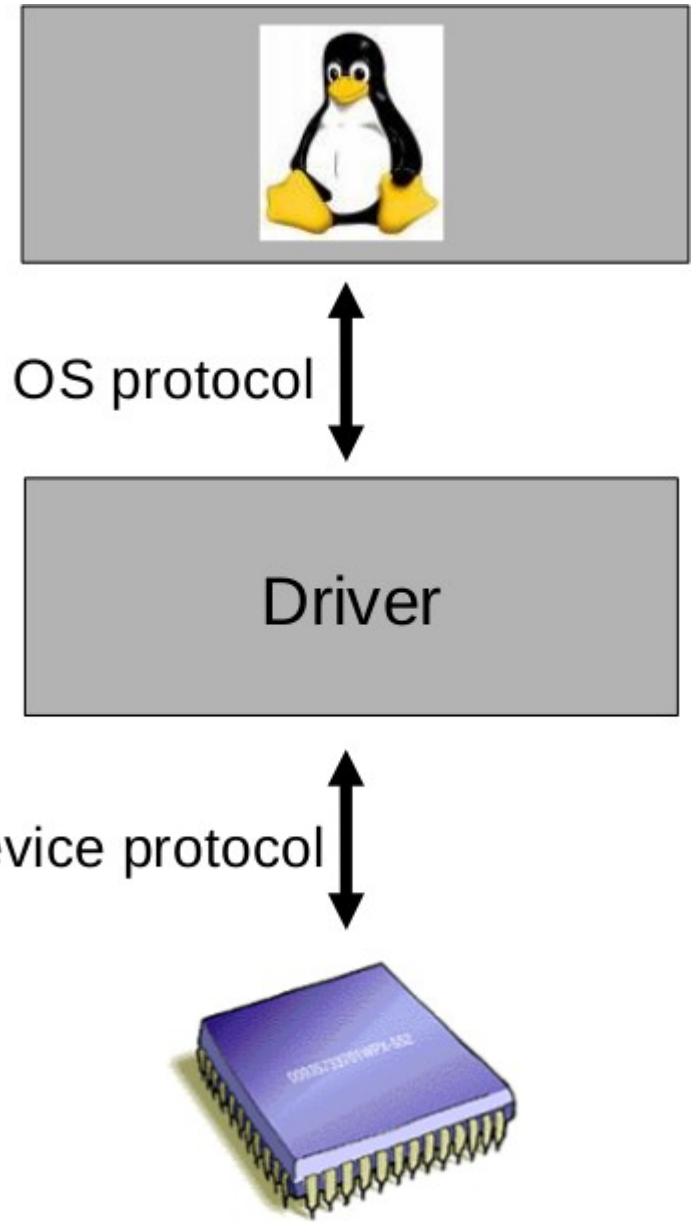

Device protocol violation examples:

- ✗Issuing a command to uninitialized device
- ✗Writing an invalid register value
- ✗Incorrectly managing DMA descriptors

Linux Device Driver Bug Portion

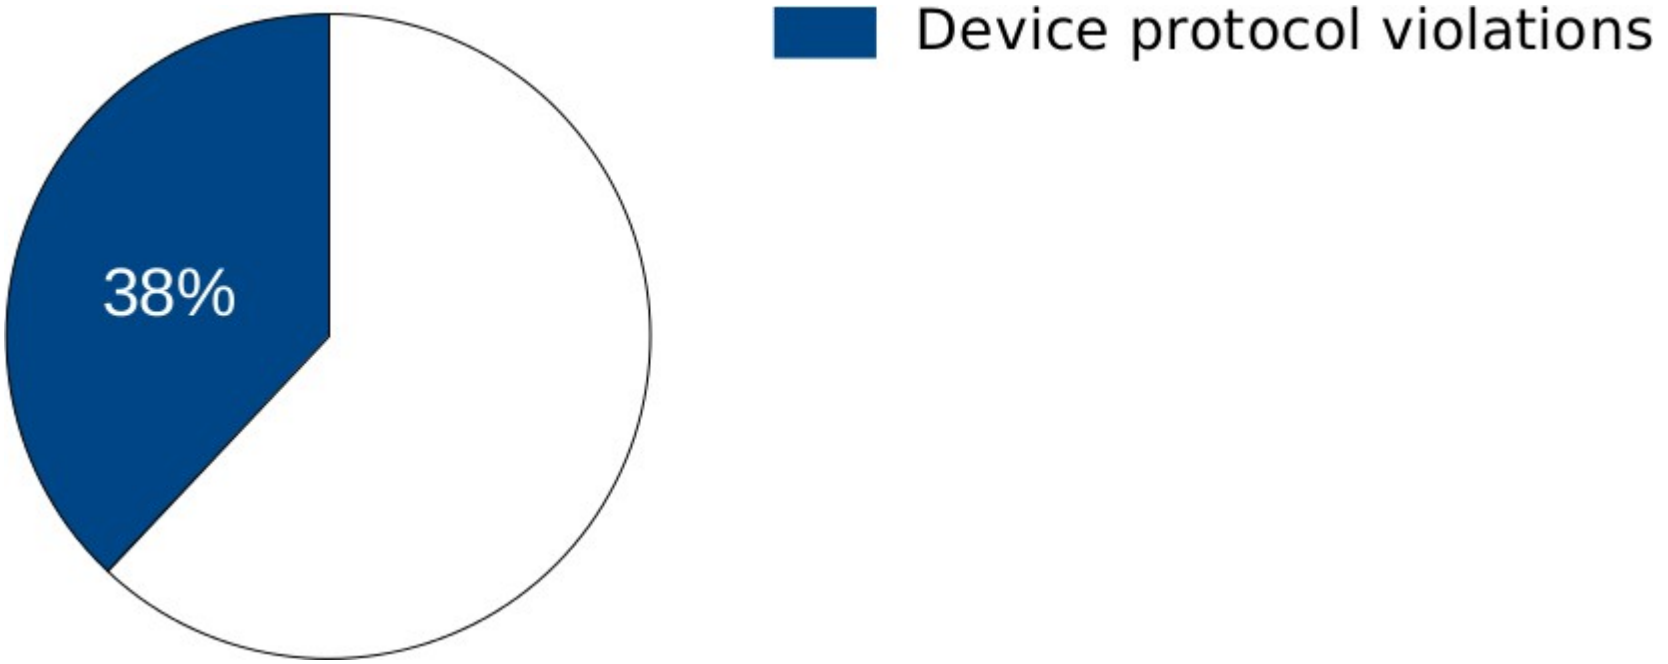

Bugs in Linux Device Driver

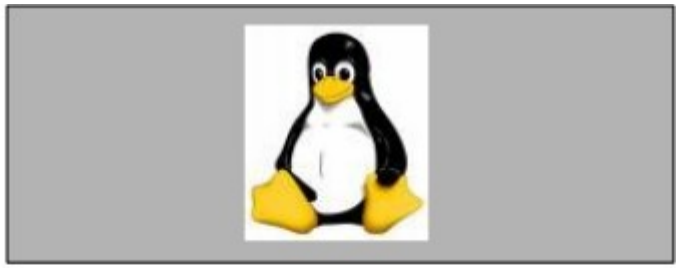

OS protocol

Driver

device protocol

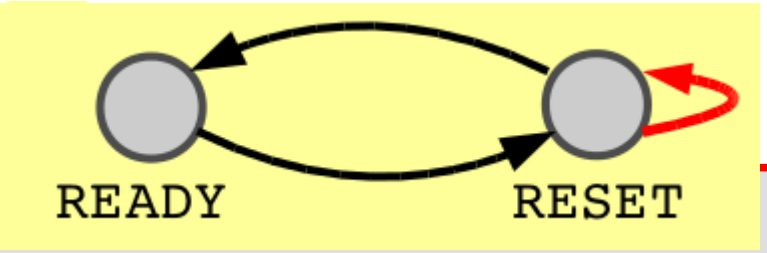

Mellanox Infinihost controller Driver

```
if(cur_state==IB_RESET &&
    new_state==IB_RESET){
   return 0;
}
```

Linux Device Driver Bug Portion

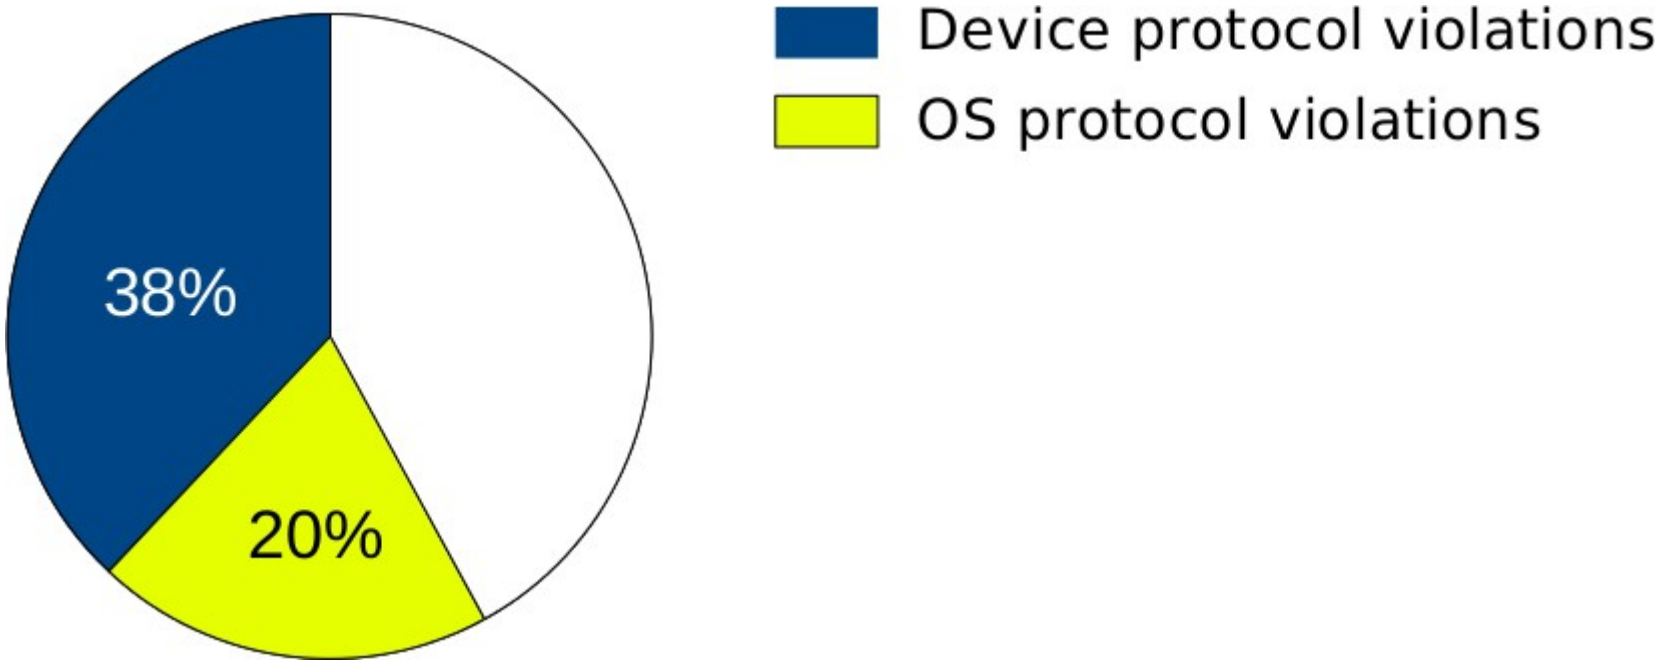

Concurrency errors

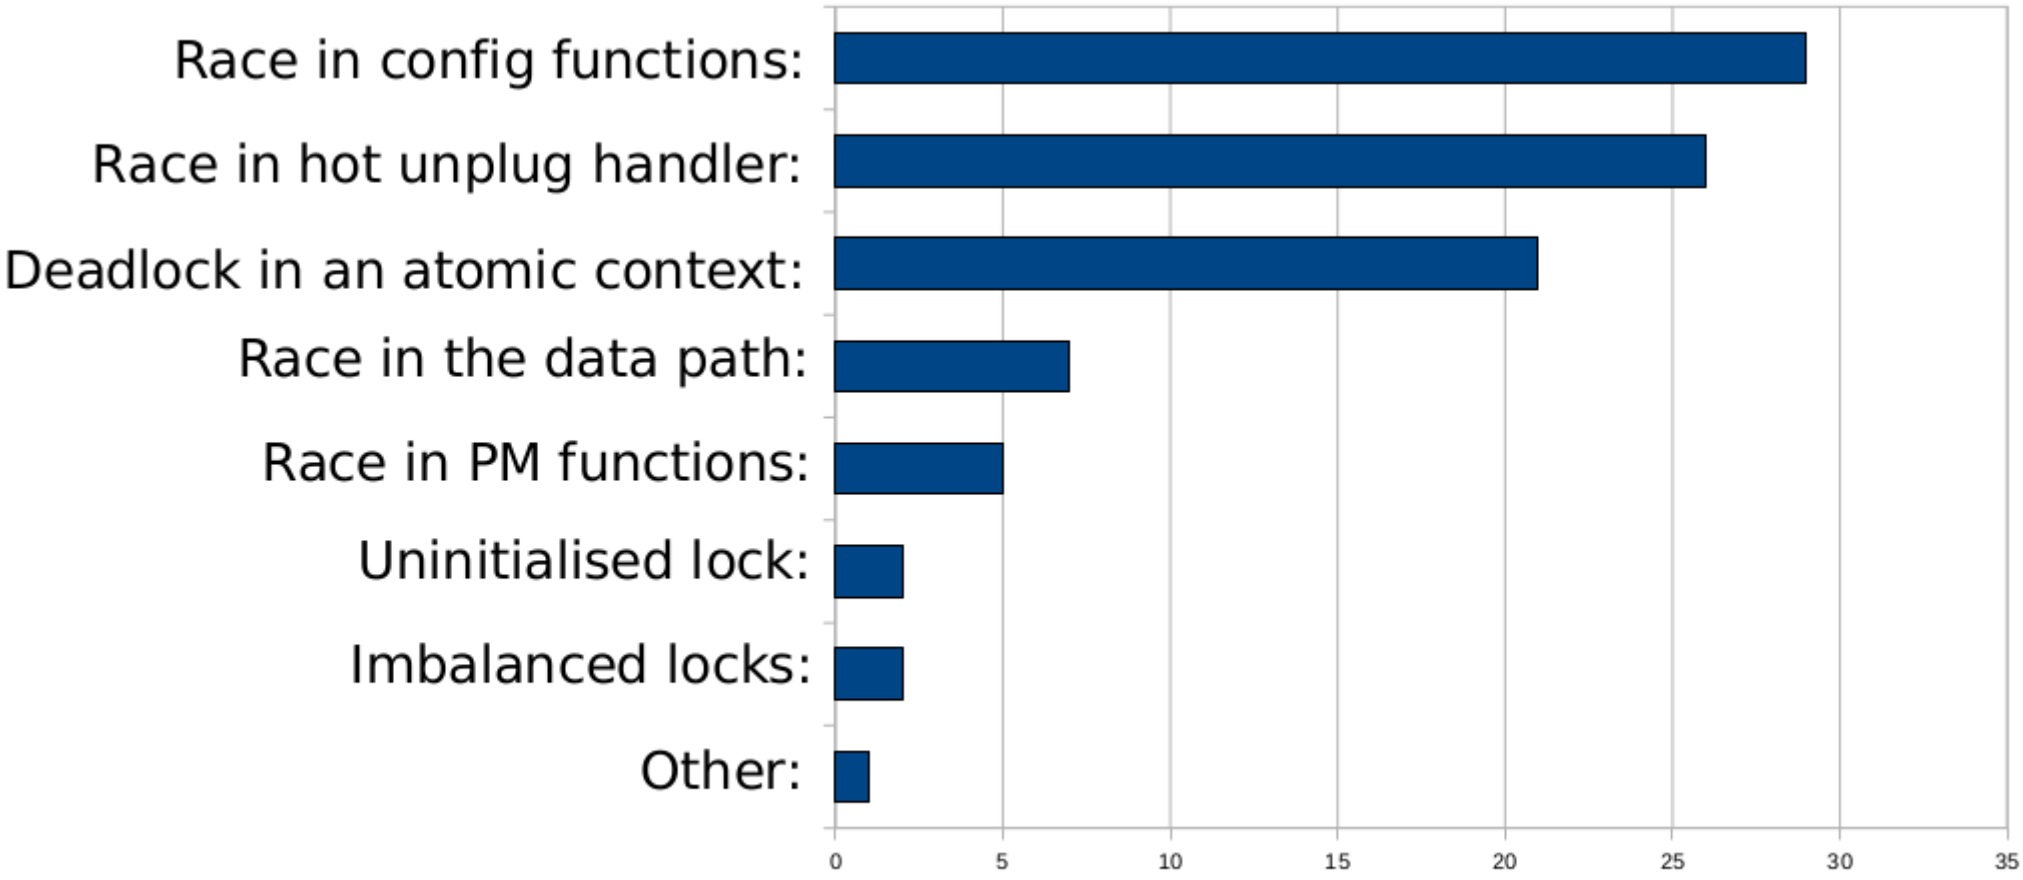

Further study about concurrency bugs

- Markus Peloquin, Lena Olson, Andrew Coonce, University of Wisconsin–Madison, *"Simultaneity Safari: A Study of Concurrency Bugs in Device Drivers"* (2009)
- Types of Device Driver Bugs

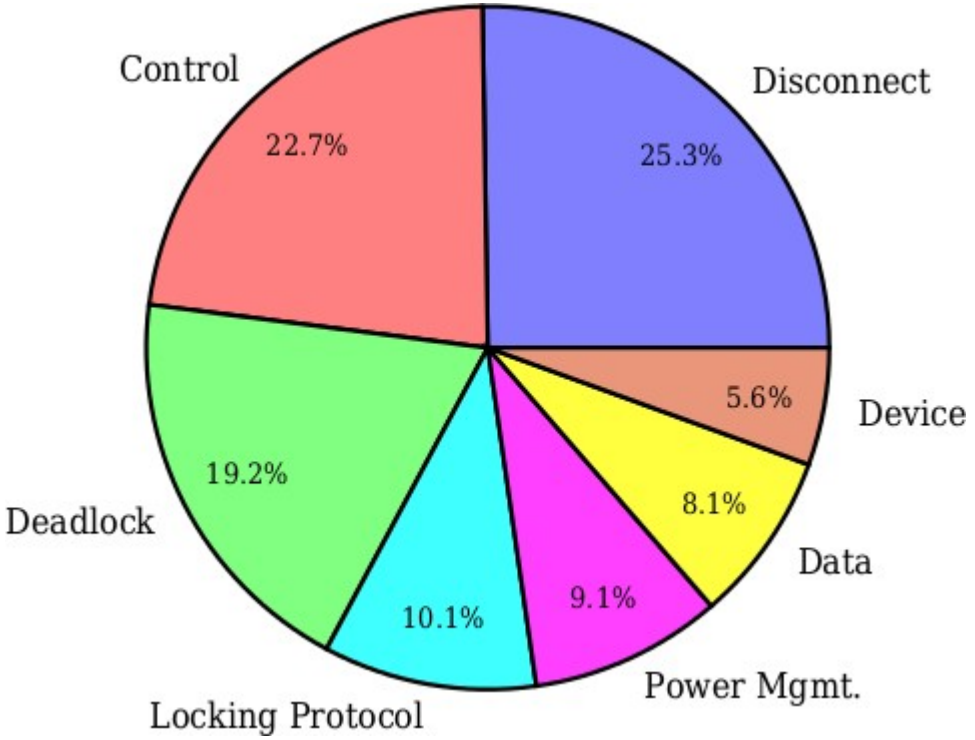

Linux Device Driver Bug Portion

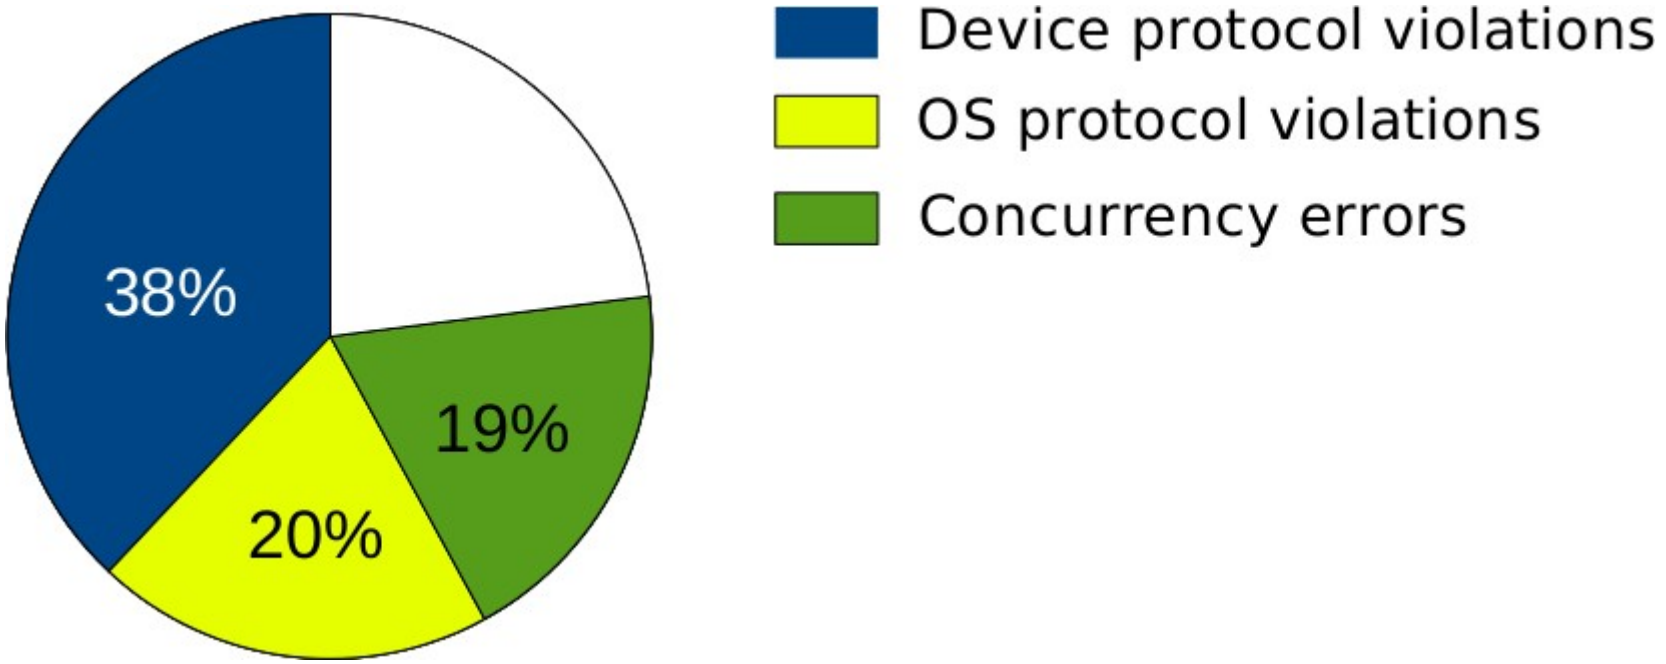

Approaches

General methods

Dealing with faulty drivers

Runtime isolation

Mach, L4, Nooks, MINIX, XFI, SafeDrive, etc.

- Performance overhead
- Transparent recovery is hard

Static analysis

SLAM, MC, Singularity, etc.

- Detects a limited subset of bugs

Approaches: Runtime Isolation

SUD-UML

[Tolerating Malicious Device Drivers in Linux, MIT CSAIL]

- In user-space, there is an unmodified Ethernet device driver running on top of SUD -UML.
- A separate driver process runs for each device driver. Shown in kernel-space are two SUD kernel modules, an Ethernet proxy driver (used by all Ethernet device drivers in SUD), and a safe PCI device access module (used by all PCI card drivers in SUD).

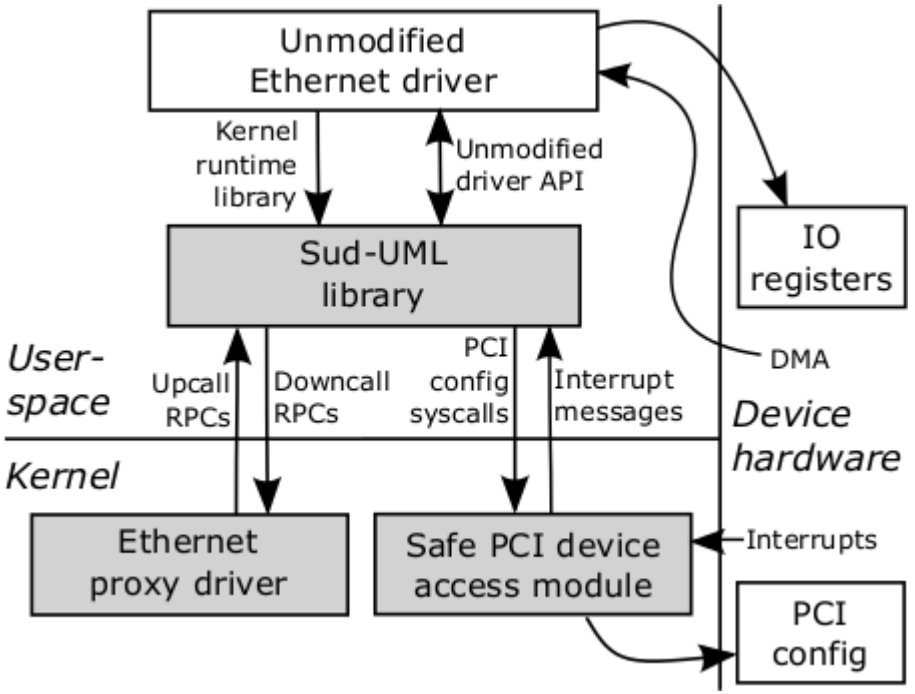

User-level Drivers

- Microkernel (MINIX/L4) / Hybrid kernel (XNU/DragonFly BSD) style
- Isolate components
  - device drivers (disk, network, graphic, ...)
  - stacks (TCP/IP, file systems, ...)
- Separate address spaces each
  - More robust components
- Problems
  - Overhead
    - hardware multiplexing
    - context switches
  - Need to handle I/O privileges

Device Driver OS: Virtualization technique

- LeVasseur et. al.: *"Unmodified Device Driver Reuse and Improved System Dependability via Virtual Machines"*, OSDI 2004
- provide a Linux environment to run drivers on L4 microkernel
  - Device Driver Environment (DDE)

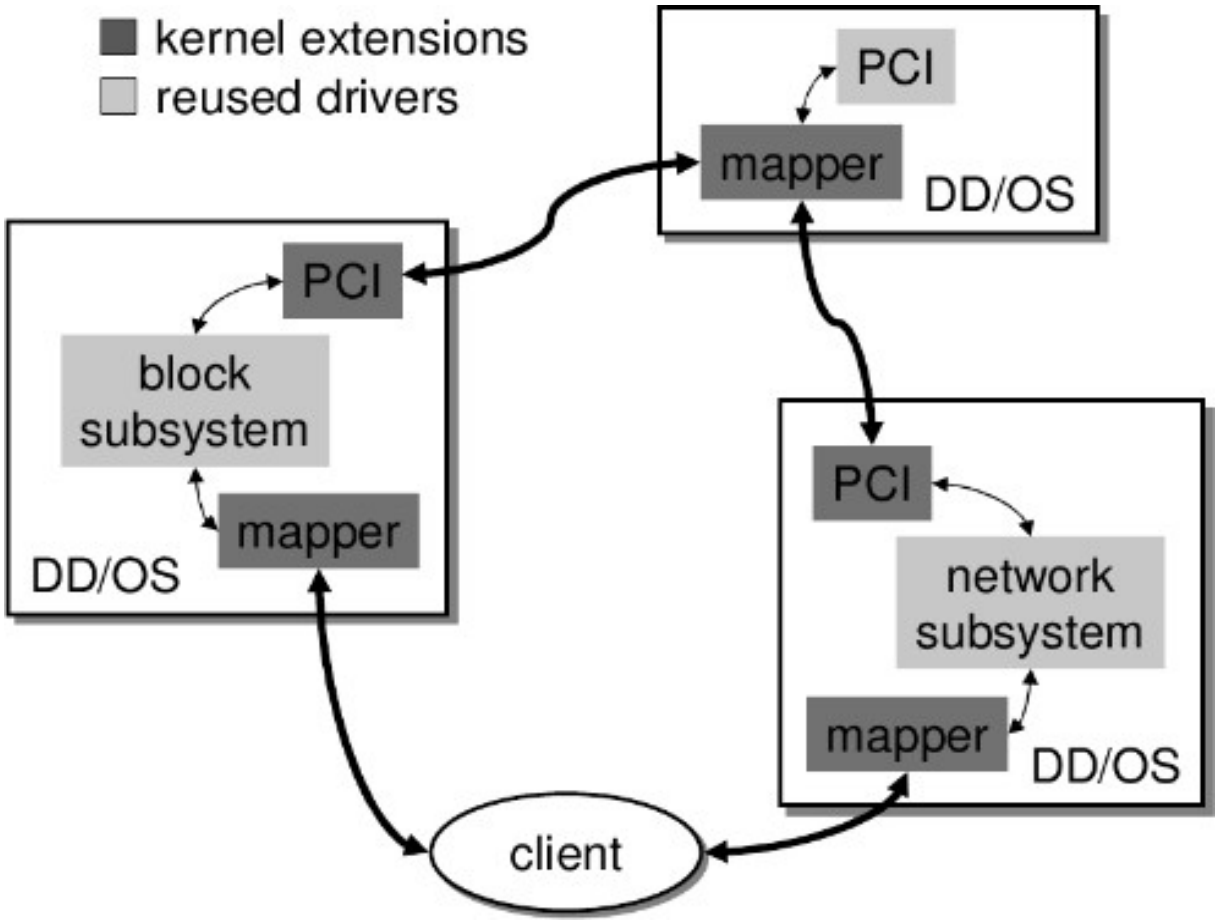

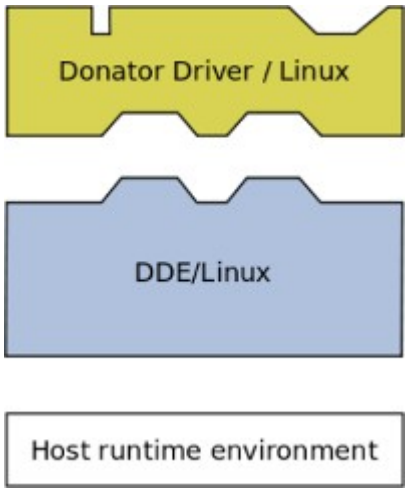

Approaches: Static Analysis

Static Analysis

- Coccinelle: Faults in Linux: Ten Years Later (ASPLOS 2011)
- Dingo: Taming Device Drivers (EuroSys 2009)
- KLEE: Automatic generation of high-coverage tests (EuroSys 2008)
- RWset: Attacking path explosion (TACAS 2008)
- EXE: Automatically generating inputs of death (CCS 2006)

Static Analysis: Instrumentation

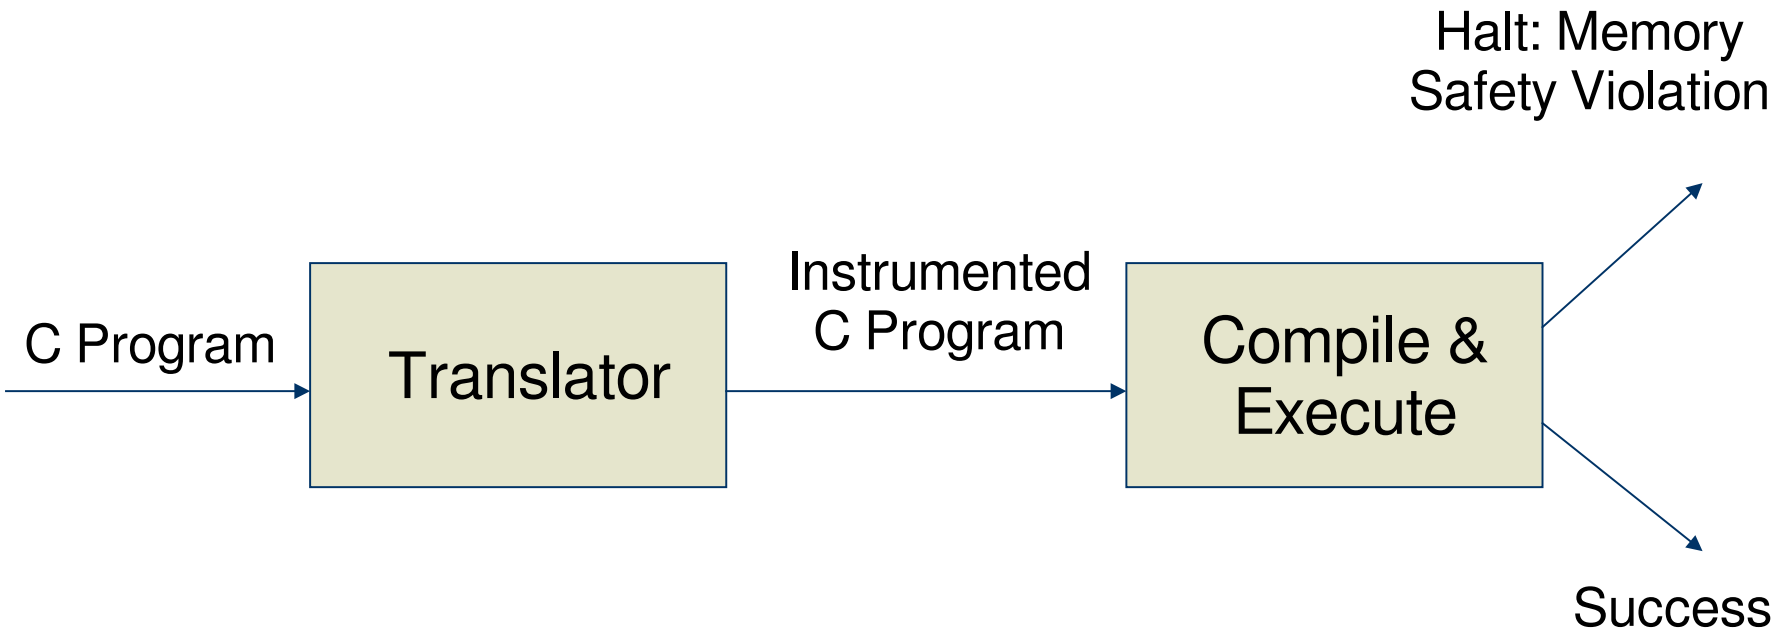

- Facts
  - 50% of software errors are due to pointers
  - 50% of security errors due to buffer overruns
- Run-time bookkeeping for memory safety
  - Array bounds information
  - Some run-time type information

Static Analysis: Instrumentation

- C statement "p++", infer p is not SAFE

```
struct { int a; int b; } *p1, *p2;
int *q = (int *)p1;    // this cast is fine
int **r = (int **)p2; // this one is not:
                       // p2 and r must be DYN
```

- DYNamic Pointer:

On use:
- null check
- bounds check
- tag check/update

Can do:
- dereference
- pointer arithmetic
- arbitrary typecasts

DYN pointer

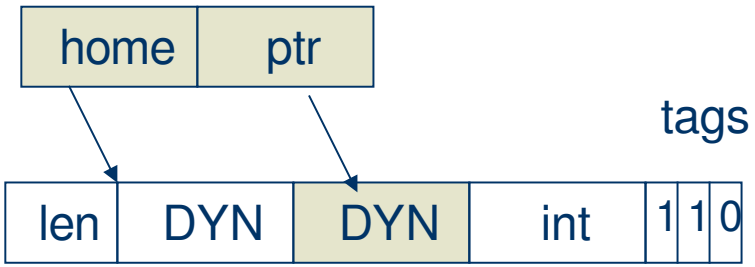

Static Analyzer for Detecting Buffer Overrun Errors in C

- “static”: no test runs
- “C”: full ANSI C + (GNU C)
- Examples

```
int *c = (int *)malloc(sizeof(int)*10);
c[i] = 1;  c[i + f()] = 1;  c[*k + (*g)()] = 1;
x = c+5;  x[1] = 1;
z->a = c; (z->a)[i] = 1;
foo(c+2);  int foo(int *d) {… d[i] = 1; …}
```

Static Analyzer: Internals

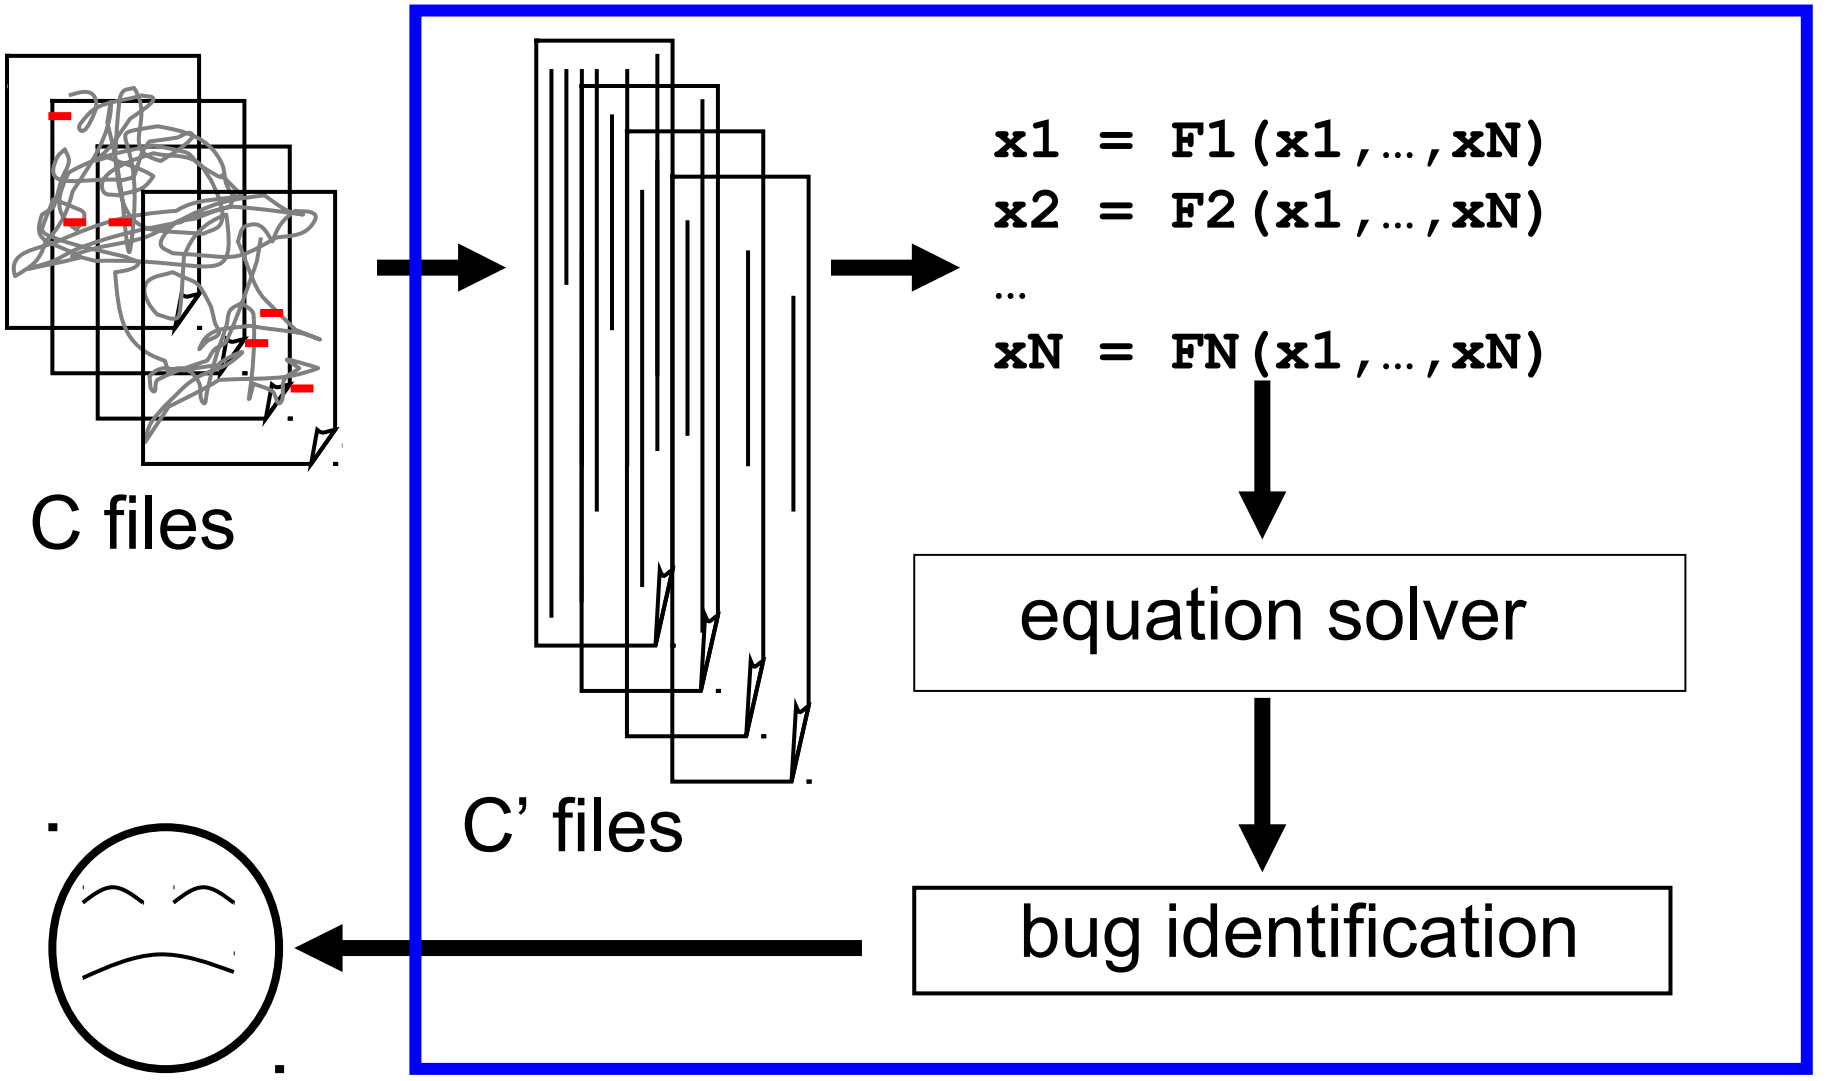

Static Analyzer - Example: cdc_acm.c (Linux device driver)

cdc-acm.c = (~/project/airac/demo-20050121) - VIM

```
601                             data = cfacm->interface[j + 1];
602                     } else if (intf == cfacm->interface[j + 1]) {
603                             ifdata = cfacm->interface[j]->altsetting + 0;
604                             data = cfacm->interface[j];
605                     } else
606                             continue;
607
608                     if (ifdata->desc.bInterfaceClass != 10 || ifdata->desc.b
    NumEndpoints < 2)
609                             continue;
610
611                     epctrl = &ifcom->endpoint[0].desc;
612                     epread = &ifdata->endpoint[0].desc;
613                     epwrite = &ifdata->endpoint[1].desc;
614
615                     if ((epctrl->bEndpointAddress & 0x80) != 0x80 || (epctrl
    ->bmAttributes & 3) != 3 ||
616                        (epread->bmAttributes & 3) != 2 || (epwrite->bmAttrib
    utes & 3) != 2 ||
617                        ((epread->bEndpointAddress & 0x80) ^ (epwrite->bEndpo
    intAddress & 0x80)) != 0x80)
618                             continue;
619
620                     if ((epread->bEndpointAddress & 0x80) != 0x80) {
621                             epread = &ifdata->endpoint[1].desc;
622                             epwrite = &ifdata->endpoint[0].desc;
623                     }
624
625                     for (minor = 0; minor < ACM_TTY_MINORS && acm_table[mino
    r]; minor++);
626                     if (acm_table[minor]) {
627                             err("no more free acm devices");
628                             return -ENODEV;
629                     }
630
631                     if (!(acm = kmalloc(sizeof(struct acm), GFP_KERNEL))) {
632                             err("out of memory");
633                             return -ENOMEM;
634                     }
635                     memset(acm, 0, sizeof(struct acm));
636
637                     ctrlsize = epctrl->wMaxPacketSize;
638                     readsize = epread->wMaxPacketSize;
639                     acm->writesize = epwrite->wMaxPacketSize;
640                     acm->control = intf;
641                     acm->data = data;
642                     acm->minor = minor;
643                     acm->dev = dev;
644
645                     acm->bh.func = acm_rx_tasklet;
646                     acm->bh.data = (unsigned long) acm;
647                     INIT_WORK(&acm->work, acm_softint, acm);
648
649                     if (!(buf = kmalloc(ctrlsize + readsize + acm->writesize
    , GFP_KERNEL))) {
650                             err("out of memory");
651                             kfree(acm);
652                             return -ENOMEM;
653                     }
-- VISUAL LINE --                                  2         626,4-25      75%
```

sab: ~/project/airac/demo-20050121

```
(0:26:0)$ ls
cdc-acm.c     osInputQueue.c     pp-rmt.c       rmt.c.out  src/
demo.sh*      pp-cdc-acm.c       preprocessed/  show.sh*
hardToSwat.c  pp-osInputQueue.c  rmt.c          show.vim
(netj@tty10@sab)-(2005-01-31 Mon 18:17:50)-------(~/project/airac/demo-20050121)
(0:27:0)$ ./show.sh cdc-acm.c
Airac: analyzing pp-cdc-acm.c main.cn
NOTE: Airac ignores all assembly codes during analysis.
      Airac ignores all compilation directives such as inline, attribute, etc.

Warning: "linux-2.6.4/include/asm/apic.h":31: assuming sizeof() has [-Inf,Inf] v
alue. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm/rwsem.h":175: assuming sizeof() has [-Inf,Inf]
 value. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm-generic/siginfo.h":275: assuming sizeof() has
[-Inf,Inf] value. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm-generic/siginfo.h":275: assuming sizeof() has
[-Inf,Inf] value. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm-generic/siginfo.h":275: assuming sizeof() has
[-Inf,Inf] value. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm-generic/siginfo.h":278: assuming sizeof() has
[-Inf,Inf] value. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm-generic/siginfo.h":278: assuming sizeof() has
[-Inf,Inf] value. (expression is neither a id nor type name.)
Warning: "linux-2.6.4/include/asm-generic/siginfo.h":278: assuming sizeof() has
[-Inf,Inf] value. (expression is neither a id nor type name.)

Parsing and transformation done:
     31034 ast nodes
       431 global variables
       410 function definitions
      2268 initial abstract memory entries

Airac options: the inlining/unrolling depth 1
               the unrolling bound 0

Analysis begins
  Fixpoint iterations with widening.................................
  The number of alarm candidates: 4
  Fixpoint iterations with narrowing.....................

  There are some overflows and underflows that may be bugs.
     Overflow (name: acm_table, size: [32, 32])
        Index: [0, +Inf]  (file: "cdc-acm.c", line: 626, column: 44)
        Index: [0, +Inf]  (file: "cdc-acm.c", line: 703, column: 47)
        Index: [1, +Inf]  (file: "cdc-acm.c", line: 625, column: 70)

  Number of alarms (buffers: 1, accesses: 3)
  User + system time: 32.998000 s
  Wall-clock time: 33 s
(netj@tty10@sab)-(2005-01-31 Mon 18:18:47)-------(~/project/airac/demo-20050121)
(0:28:0)$
```

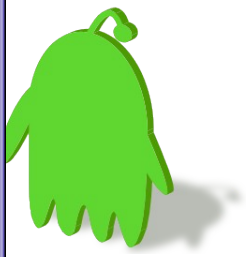

Static Analyzer - Coverity

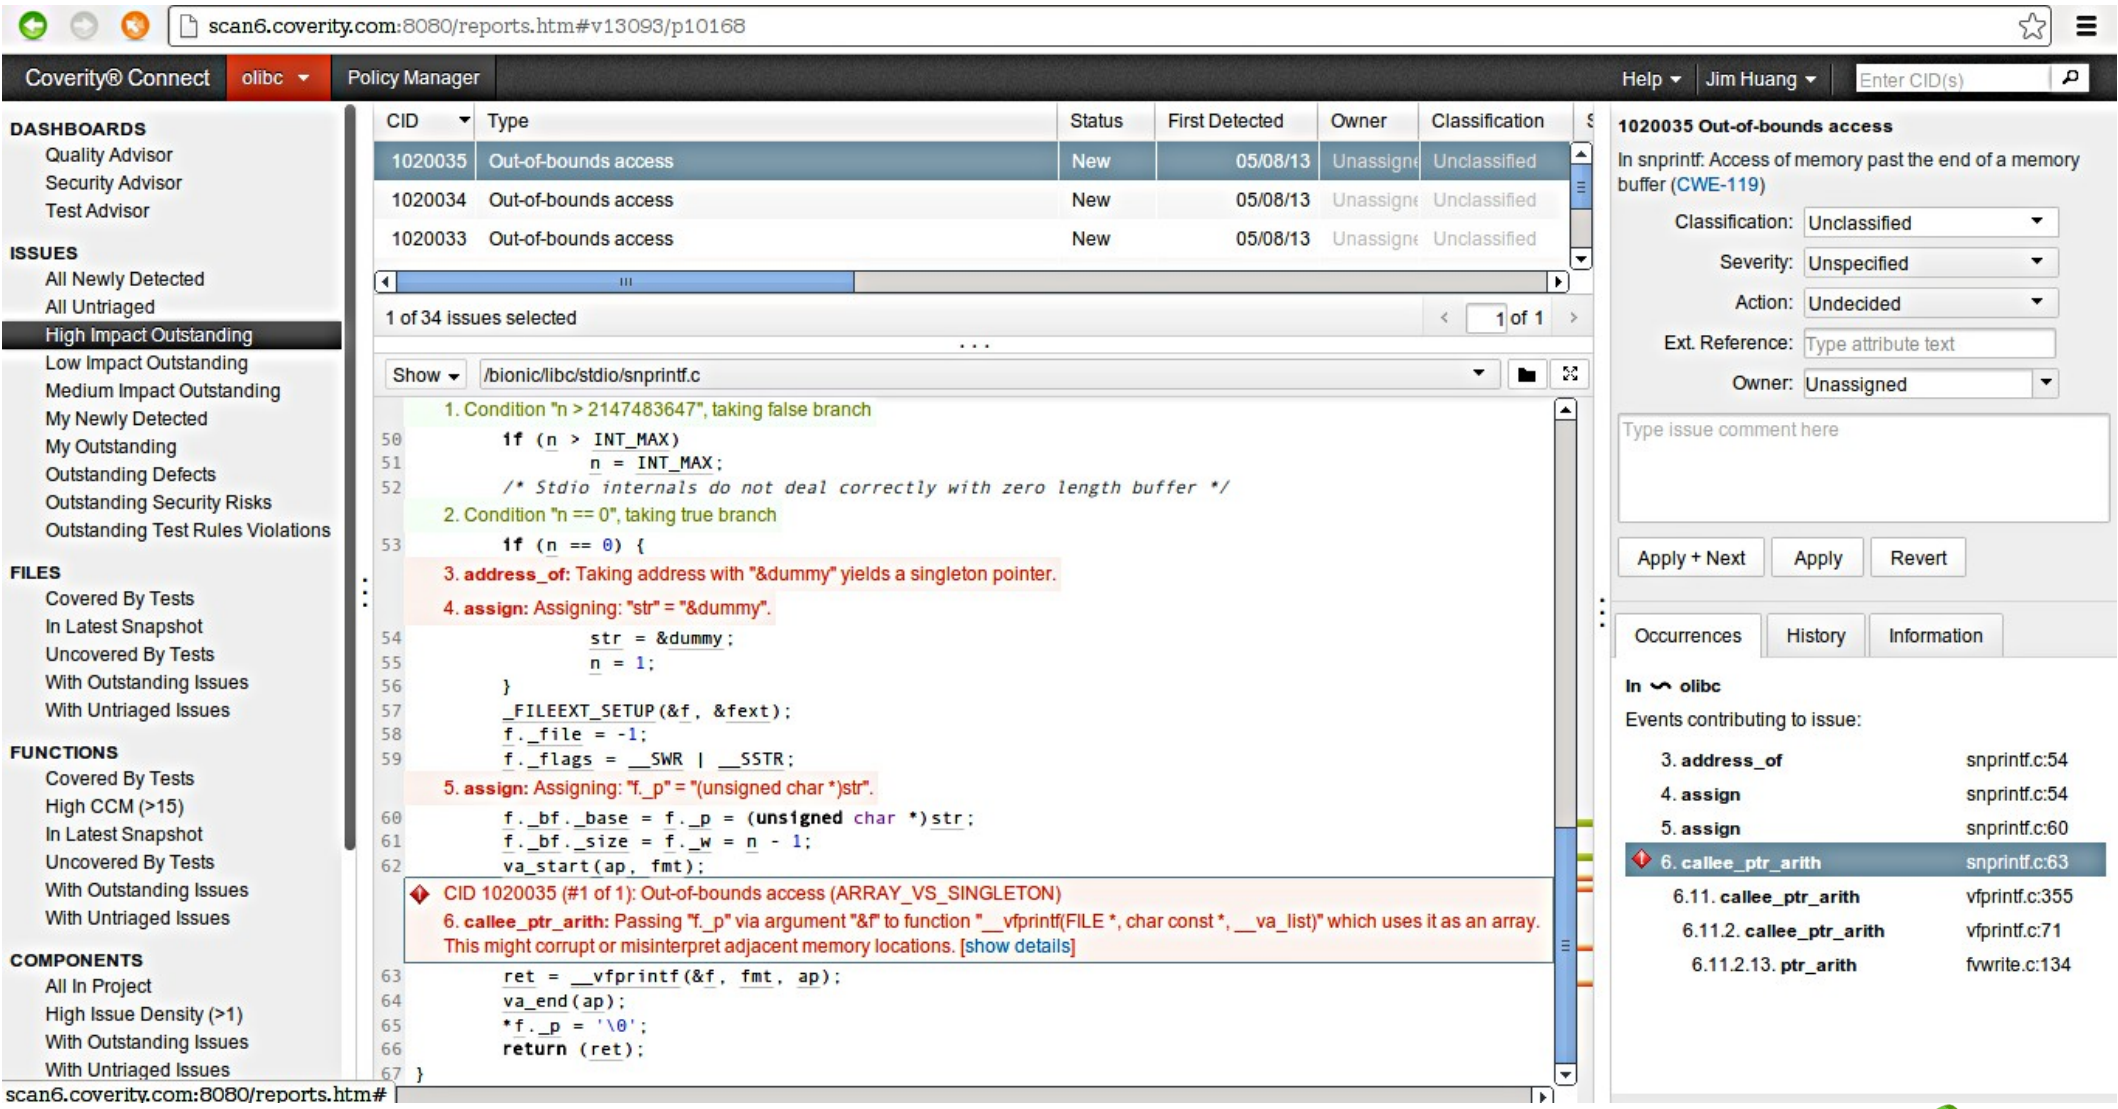

Securing Driver: Dingo

[Dingo: Taming device drivers, 2009]

- Observations:
  - drivers fail to obey device spec
  - developers misunderstand OS interface
  - multi-threading is bad
- Drivers run as part of the kernel
  - Need to deal with concurrent invocations
  - Shared state must be maintained
- Synchronization is hard to get right
  - Race conditions and deadlocks
  - 20% of bugs in device drivers

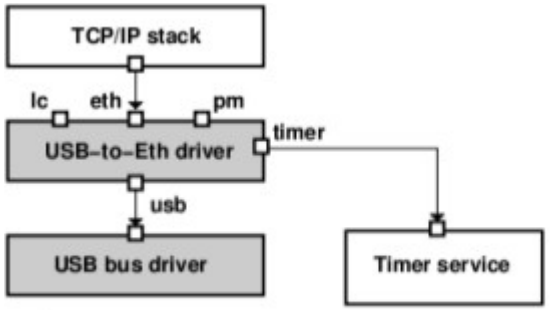

Ports of the USB-to-Ethernet adapter driver

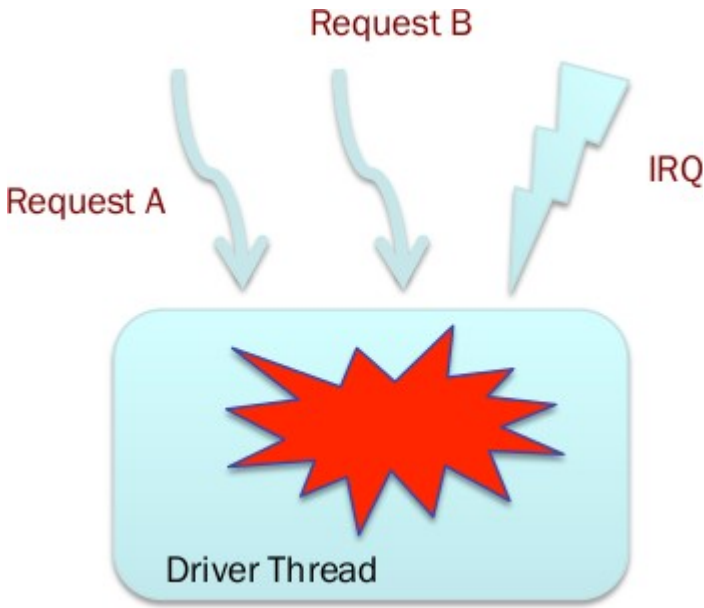

Securing Driver: Dingo

- Tingu: state-chart-based specification of device protocols
  - Event-based state transition
  - Timeouts
  - Variables

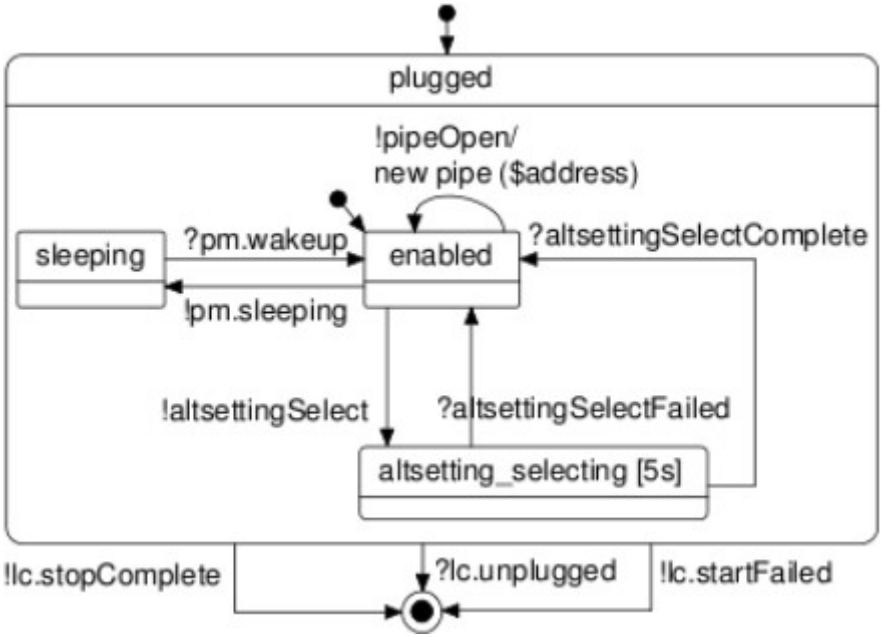

Securing Driver: Dingo

- Device driver architecture
- Single-threaded
  - Builtin atomicity
  - Not a performance problem for most drivers
- Event-based
  - Developers implement a Tingu specification
- Can use Tingu specs to generate runtime driver monitors

Deal with concurrency bugs

Threads

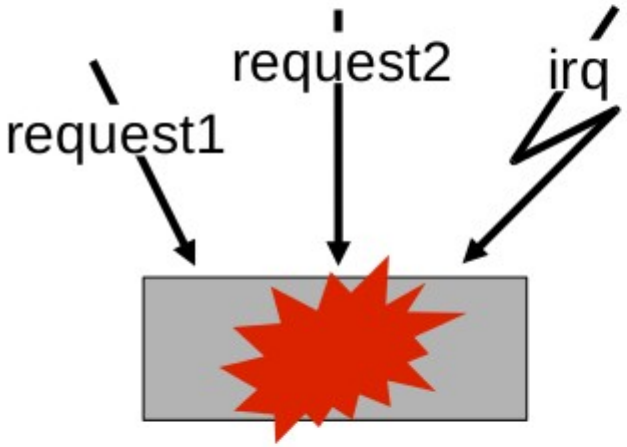

Events

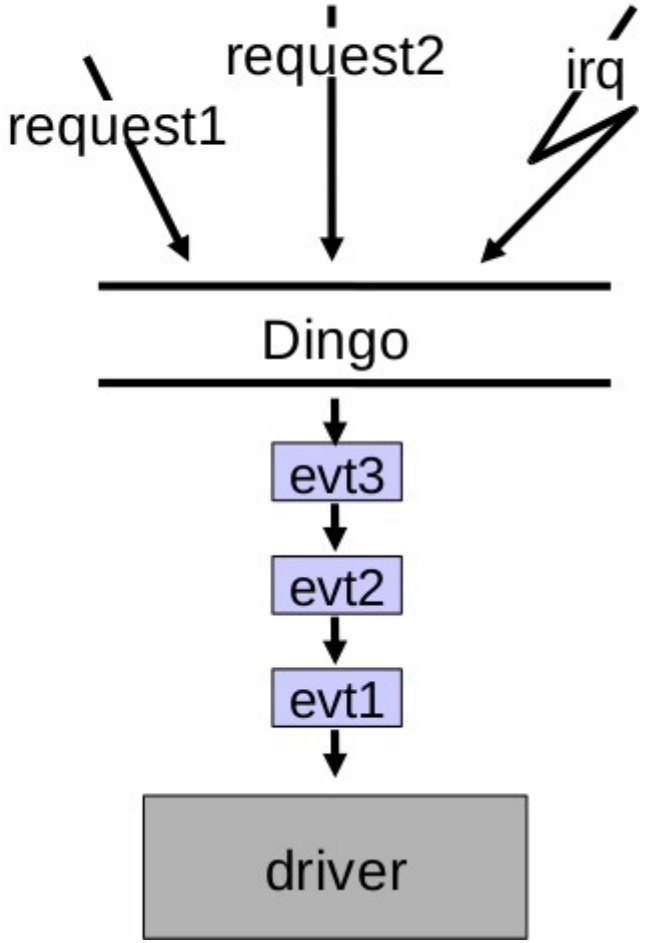

Event-based Device Driver

```
probe() {
  ...
  write();
  msleep(10);
  read();
  ...

}
```

'Rip' the stack

```
probe() {
  ...
  write();
  drv->state = 1;
  schedule_timeout(10, drv, timeout);
  return;
}

timeout(drv) {
  switch(drv->state) {
    case 1:
      read();
      break;
    ...
  }
}
```

Insightful Researches

- DevIL (OSDI 2000): generate driver from an IDL spec of the device interface

  "...our vision is that Devil specifications either should be written by device vendors or should be widely available aspublic domain libraries..."
- Termite (SOSP 2009): use device driver spec (VHDL) to generate
  - Lets vendors generate drivers on their own
- RevNIC (EuroSys 2010):
  - Obtain I/O trace from existing driver (Windows)
  - Analyze driver binary
  - Generate Linux driver

Conclusion

- Device drivers are hard than expected while quality and stability are considered.
- Security risks exist inside every area of system software. Device driver is the major.
- It is a common technique to introduce virtual buses for isolating device resources.
- Performing static analysis as early as possible when you design the device driver model and adapt legacy implementations upon the revised frameworks.

Reference

- *“Dingo: Taming Device Drivers”*, Leonid Ryzhyk, Peter Chubb, Ihor Kuz, Gernot Heiser, UNSW/NICTA/Open Kernel Labs (2009)
- *"Hardware and Device Drivers"*, Björn Döbel, TU Dresden (2012)
- *"Configuration Coverage in the Analysis of Large-Scale System Software"*, Reinhard Tartler, Daniel Lohmann, Christian Dietrich, Christoph Egger, Julio Sincero, Friedrich-Alexander University (2011)
- *“AIRAC: A Static Analyzer for Detecting All Buffer Overrun Errors in C Programs"*, Kwangkeun Yi, Seoul National University (2005)
- *“CCured: Taming C Pointers”*, George Necula, Scott McPeak, Wes Weimer, Berkeley (2002)

http://0xlab.org